

Managing Network Datalinks in Oracle® Solaris 11.3



Part No: E54788
November 2016

Part No: E54788

Copyright © 2011, 2016, Oracle and/or its affiliates. All rights reserved.

This software and related documentation are provided under a license agreement containing restrictions on use and disclosure and are protected by intellectual property laws. Except as expressly permitted in your license agreement or allowed by law, you may not use, copy, reproduce, translate, broadcast, modify, license, transmit, distribute, exhibit, perform, publish, or display any part, in any form, or by any means. Reverse engineering, disassembly, or decompilation of this software, unless required by law for interoperability, is prohibited.

The information contained herein is subject to change without notice and is not warranted to be error-free. If you find any errors, please report them to us in writing.

If this is software or related documentation that is delivered to the U.S. Government or anyone licensing it on behalf of the U.S. Government, then the following notice is applicable:

U.S. GOVERNMENT END USERS. Oracle programs, including any operating system, integrated software, any programs installed on the hardware, and/or documentation, delivered to U.S. Government end users are "commercial computer software" pursuant to the applicable Federal Acquisition Regulation and agency-specific supplemental regulations. As such, use, duplication, disclosure, modification, and adaptation of the programs, including any operating system, integrated software, any programs installed on the hardware, and/or documentation, shall be subject to license terms and license restrictions applicable to the programs. No other rights are granted to the U.S. Government.

This software or hardware is developed for general use in a variety of information management applications. It is not developed or intended for use in any inherently dangerous applications, including applications that may create a risk of personal injury. If you use this software or hardware in dangerous applications, then you shall be responsible to take all appropriate fail-safe, backup, redundancy, and other measures to ensure its safe use. Oracle Corporation and its affiliates disclaim any liability for any damages caused by use of this software or hardware in dangerous applications.

Oracle and Java are registered trademarks of Oracle and/or its affiliates. Other names may be trademarks of their respective owners.

Intel and Intel Xeon are trademarks or registered trademarks of Intel Corporation. All SPARC trademarks are used under license and are trademarks or registered trademarks of SPARC International, Inc. AMD, Opteron, the AMD logo, and the AMD Opteron logo are trademarks or registered trademarks of Advanced Micro Devices. UNIX is a registered trademark of The Open Group.

This software or hardware and documentation may provide access to or information about content, products, and services from third parties. Oracle Corporation and its affiliates are not responsible for and expressly disclaim all warranties of any kind with respect to third-party content, products, and services unless otherwise set forth in an applicable agreement between you and Oracle. Oracle Corporation and its affiliates will not be responsible for any loss, costs, or damages incurred due to your access to or use of third-party content, products, or services, except as set forth in an applicable agreement between you and Oracle.

Documentation Accessibility

For information about Oracle's commitment to accessibility, visit the Oracle Accessibility Program website at <http://www.oracle.com/pls/topic/lookup?ctx=acc&id=docacc>.

Access to Oracle Support

Oracle customers that have purchased support have access to electronic support through My Oracle Support. For information, visit <http://www.oracle.com/pls/topic/lookup?ctx=acc&id=info> or visit <http://www.oracle.com/pls/topic/lookup?ctx=acc&id=trs> if you are hearing impaired.

Référence: E54788

Copyright © 2011, 2016, Oracle et/ou ses affiliés. Tous droits réservés.

Ce logiciel et la documentation qui l'accompagne sont protégés par les lois sur la propriété intellectuelle. Ils sont concédés sous licence et soumis à des restrictions d'utilisation et de divulgation. Sauf stipulation expresse de votre contrat de licence ou de la loi, vous ne pouvez pas copier, reproduire, traduire, diffuser, modifier, accorder de licence, transmettre, distribuer, exposer, exécuter, publier ou afficher le logiciel, même partiellement, sous quelque forme et par quelque procédé que ce soit. Par ailleurs, il est interdit de procéder à toute ingénierie inverse du logiciel, de le désassembler ou de le décompiler, excepté à des fins d'interopérabilité avec des logiciels tiers ou tel que prescrit par la loi.

Les informations fournies dans ce document sont susceptibles de modification sans préavis. Par ailleurs, Oracle Corporation ne garantit pas qu'elles soient exemptes d'erreurs et vous invite, le cas échéant, à lui en faire part par écrit.

Si ce logiciel, ou la documentation qui l'accompagne, est livré sous licence au Gouvernement des Etats-Unis, ou à quiconque qui aurait souscrit la licence de ce logiciel pour le compte du Gouvernement des Etats-Unis, la notice suivante s'applique:

U.S. GOVERNMENT END USERS. Oracle programs, including any operating system, integrated software, any programs installed on the hardware, and/or documentation, delivered to U.S. Government end users are "commercial computer software" pursuant to the applicable Federal Acquisition Regulation and agency-specific supplemental regulations. As such, use, duplication, disclosure, modification, and adaptation of the programs, including any operating system, integrated software, any programs installed on the hardware, and/or documentation, shall be subject to license terms and license restrictions applicable to the programs. No other rights are granted to the U.S. Government.

Ce logiciel ou matériel a été développé pour un usage général dans le cadre d'applications de gestion des informations. Ce logiciel ou matériel n'est pas conçu ni n'est destiné à être utilisé dans des applications à risque, notamment dans des applications pouvant causer des dommages corporels. Si vous utilisez ce logiciel ou matériel dans le cadre d'applications dangereuses, il est de votre responsabilité de prendre toutes les mesures de secours, de sauvegarde, de redondance et autres mesures nécessaires à son utilisation dans des conditions optimales de sécurité. Oracle Corporation et ses affiliés déclinent toute responsabilité quant aux dommages causés par l'utilisation de ce logiciel ou matériel pour ce type d'applications.

Oracle et Java sont des marques déposées d'Oracle Corporation et/ou de ses affiliés. Tout autre nom mentionné peut correspondre à des marques appartenant à d'autres propriétaires qu'Oracle.

Intel et Intel Xeon sont des marques ou des marques déposées d'Intel Corporation. Toutes les marques SPARC sont utilisées sous licence et sont des marques ou des marques déposées de SPARC International, Inc. AMD, Opteron, le logo AMD et le logo AMD Opteron sont des marques ou des marques déposées d'Advanced Micro Devices. UNIX est une marque déposée d'The Open Group.

Ce logiciel ou matériel et la documentation qui l'accompagne peuvent fournir des informations ou des liens donnant accès à des contenus, des produits et des services émanant de tiers. Oracle Corporation et ses affiliés déclinent toute responsabilité ou garantie expresse quant aux contenus, produits ou services émanant de tiers, sauf mention contraire stipulée dans un contrat entre vous et Oracle. En aucun cas, Oracle Corporation et ses affiliés ne sauraient être tenus pour responsables des pertes subies, des coûts occasionnés ou des dommages causés par l'accès à des contenus, produits ou services tiers, ou à leur utilisation, sauf mention contraire stipulée dans un contrat entre vous et Oracle.

Accessibilité de la documentation

Pour plus d'informations sur l'engagement d'Oracle pour l'accessibilité à la documentation, visitez le site Web Oracle Accessibility Program, à l'adresse <http://www.oracle.com/pls/topic/lookup?ctx=acc&id=docacc>.

Accès aux services de support Oracle

Les clients Oracle qui ont souscrit un contrat de support ont accès au support électronique via My Oracle Support. Pour plus d'informations, visitez le site <http://www.oracle.com/pls/topic/lookup?ctx=acc&id=info> ou le site <http://www.oracle.com/pls/topic/lookup?ctx=acc&id=trs> si vous êtes malentendant.

Contents

Using This Documentation	11
 1 Introduction to Managing Network Datalinks	 13
What's New in Managing Network Datalinks in Oracle Solaris 11.3	13
Features and Components That Are Used to Manage Network Datalinks	14
Link Aggregations	14
Virtual Local Area Networks	15
Private Virtual Local Area Networks	15
Bridged Networks	15
Link Layer Discovery Protocol	15
Data Center Bridging	16
 2 Configuring High Availability by Using Link Aggregations	 17
Overview of Link Aggregations	17
Benefits of Link Aggregation	18
Trunk Aggregations	19
Using a Switch	20
Back-to-Back Trunk Aggregation Configuration	21
Using a Switch With the Link Aggregation Control Protocol	22
Defining Aggregation Policies for Load Balancing	22
Datalink Multipathing Aggregations	23
Advantages of DLMP Aggregation	23
How DLMP Aggregation Works	24
Failure Detection in DLMP Aggregation	26
Requirements for Link Aggregations	27
Creating a Link Aggregation	28
▼ How to Create a Link Aggregation	28
Adding a Link to an Aggregation	31
▼ How to Add a Link to an Aggregation	32
Removing a Link From an Aggregation	33

Modifying a Trunk Aggregation	33
Configuring Probe-Based Failure Detection for DLMP Aggregation	34
▼ How to Configure Probe-Based Failure Detection for DLMP	35
Monitoring Probe-Based Failure Detection	38
Deleting a Link Aggregation	40
▼ How to Delete a Link Aggregation	41
Switching Between Trunk and DLMP Aggregations	41
▼ How to Switch Between Link Aggregation Types	42
Use Case: Configuring a Link Aggregation	43
Comparing Trunk and DLMP Aggregation	45
 3 Configuring Virtual Networks by Using Virtual Local Area Networks	47
Overview of Deploying VLANs	47
When to Use VLANs	48
Assigning VLAN Names	48
VLAN Topology	48
Using VLANs With Virtualization	51
Planning a VLAN Configuration	52
Configuring a VLAN	52
▼ How to Configure a VLAN	53
Configuring VLANs Over a Link Aggregation	57
▼ How to Configure VLANs Over a Link Aggregation	57
Configuring VLANs on a Legacy Device	59
▼ How to Configure VLANs on a Legacy Device	59
Displaying VLAN Information	60
Modifying VLANs	61
Modifying the VLAN ID of a VLAN	61
Migrating a VLAN to Another Underlying Link	61
Deleting a VLAN	63
Using VLAN Tagging for SR-IOV-Enabled Ports	64
Use Case: Combining Link Aggregations and VLAN Configurations	64
Use Case: Configuring Flows for a VLAN	67
Use Case: Configuring Bandwidth for VLAN VNICs	68
 4 Configuring Private Virtual Local Area Networks	71
Overview of Private VLANs	71
Advantages of Using PVLANS	73
PVLAN Ports	73
Tagging the Outgoing Traffic	73

PVLAN Configuration Requirements	74
PVLANs With Zones	74
Configuring a Private VLAN	75
Modifying Private VLANs	76
Deleting a Private VLAN	77
Assigning a PVLAN to a Zone	77
▼ How to Assign a PVLAN to a Zone	77
5 Administering Bridging Features	79
Overview of Bridged Networks	79
Simple Bridged Network	80
Bridged Network Ring	81
How a Bridged Network Works	82
Bridging Protocols	83
STP Daemon	84
TRILL Daemon	85
Creating a Bridge	85
Modifying the Protection Type for a Bridge	87
Adding Links to an Existing Bridge	87
Removing Links From a Bridge	88
Setting Link Properties for a Bridge	88
Displaying Bridge Configuration Information	89
Displaying Information About Configured Bridges	89
Displaying Configuration Information About Bridge Links	91
Deleting a Bridge From the System	91
▼ How to Delete a Bridge From the System	91
Administering VLANs on Bridged Networks	92
▼ How to Configure VLANs Over a Datalink That Is Part of a Bridge	92
VLANs and the STP and TRILL Protocols	93
Debugging Bridges	93
6 Exchanging Network Connectivity Information With Link Layer Discovery Protocol	95
Overview of LLDP	95
Components of an LLDP Implementation	96
Information Sources of the LLDP Agent	97
LLDP Agent Modes	97
Information the LLDP Agent Advertises	98
Mandatory TLV Units	98

Optional TLV Units	98
TLV Unit Properties	99
Enabling LLDP on the System	101
▼ How to Install the LLDP Package	101
▼ How to Enable LLDP Globally	102
▼ How to Enable LLDP for Specific Ports	103
Specifying TLV Units and Values for the LLDP Packet of an Agent	105
▼ How to Specify TLV Units for the LLDP Packet of an Agent	105
▼ How to Define TLV Units	107
Disabling LLDP	109
▼ How to Disable LLDP	109
Monitoring LLDP Agents	110
Displaying the Advertised Information	110
Displaying LLDP Statistics	113
 7 Managing Converged Networks by Using Data Center Bridging	 115
Overview of Data Center Bridging	115
Considerations When Using DCB	117
Priority-Based Flow Control	117
Enhanced Transmission Selection	118
Enabling DCBX	119
▼ How to Enable the Data Center Bridging Exchange Feature Manually	119
Setting the Mode of Operation for DCB	120
▼ How to Set the Mode of Operation for DCB	121
Customizing Priority-Based Flow Control for DCB	121
Setting the PFC-Related Datalink Properties	122
Setting the PFC TLV Units	123
Displaying PFC Configuration Information	124
Displaying Datalink Properties	124
Displaying the Capability of the Local Host to Synchronize PFC Information	125
Displaying PFC Mapping Information Between Host and Peer	125
Displaying Priority Definitions	126
Application Priority Configurations	126
Customizing Enhanced Transmission Selection for DCB	127
Setting the ETS-Related Datalink Properties	128
Setting ETS TLV Units	129
Recommending ETS Configuration to the Peer	130
Displaying ETS Configuration Information	132

A Link Aggregations and IPMP: Feature Comparison	135
B Packet Format of Transitive Probes	137
Index	139

Using This Documentation

- **Overview** – Provides an overview of the advanced features that are used to manage network datalinks to improve network performance. Describes how to combine links into aggregations by using trunk or DLMP aggregation, divide your network into subnetworks by using virtual local area networks, connect separate network segments by using bridges, exchange network connectivity information by using Link Layer Discovery Protocol, and manage converged networks by using data center bridging.
- **Audience** – System administrators.
- **Required knowledge** – Basic and some advanced network administration skills.

Product Documentation Library

Documentation and resources for this product and related products are available at <http://www.oracle.com/pls/topic/lookup?ctx=E53394>.

Feedback

Provide feedback about this documentation at <http://www.oracle.com/goto/docfeedback>.

Introduction to Managing Network Datalinks

This chapter introduces the advanced features that are used to manage network datalinks, which are described in the remaining chapters of this book. Use of the different technologies described in this book depends on specific circumstances. Also, some hardware configurations might require you to use a specific type of feature. Therefore, you do not need to complete all the configuration procedures in this book. Instead, select and deploy the technologies that address your network requirements.

Before you perform any of the configurations that are described in this book, you must have completed basic network configuration and have an understanding about the fundamental datalink configuration. For information about basic network configuration, see [Configuring and Managing Network Components in Oracle Solaris 11.3](#). For information about elements of managing datalink configuration, see [Chapter 2, “Administering Datalink Configuration in Oracle Solaris” in Configuring and Managing Network Components in Oracle Solaris 11.3](#).

For a summary of network configuration features in Oracle Solaris, see [Chapter 1, “Summary of Oracle Solaris Network Administration” in Strategies for Network Administration in Oracle Solaris 11.3](#).

This chapter contains the following topics:

- [“What's New in Managing Network Datalinks in Oracle Solaris 11.3” on page 13](#)
- [“Features and Components That Are Used to Manage Network Datalinks” on page 14](#)

What's New in Managing Network Datalinks in Oracle Solaris 11.3

For existing customers, this section highlights the following key changes in this release:

- **Private virtual local area networks** – Oracle Solaris 11.3 includes support for private virtual local area networks (PVLANS), as defined in the RFC 5517 standard. You can use a PVLAN to divide a regular VLAN (primary) into sub-VLANs (secondary). For more information, see [“Overview of Private VLANs” on page 71](#).

- **Converged Enhanced Ethernet (CEE) data center bridging exchange (DCBX)** – Oracle Solaris 11.3 includes support for Converged Enhanced Ethernet (CEE) data center bridging exchange (DCBX) in addition to IEEE so that Oracle Solaris Fibre Channel over Ethernet (FcoE) can be deployed in environments with a diverse set of switches when using data center bridging (DCB). Oracle Solaris CEE DCBX supports priority-based flow control (PFC) and Application type-length-value (TLV). For more information, see [“Setting the Mode of Operation for DCB” on page 120](#).

Features and Components That Are Used to Manage Network Datalinks

Managing network datalinks refers to the use of features and technologies to fine tune the way your systems process the network traffic. Systems that are configured with these technologies can manage network traffic better, which contributes to the improvement of the network's total performance. Although these features address different areas of network operations, they provide common benefits such as network connectivity, network administration, and efficiency.

You use the `dladm` command to manage network datalinks by using the features such as link aggregation, virtual local area networks, and bridged networks. In addition to managing datalinks, you can use the `dladm` command to manage capture links. The Firewall uses capture links to log packets that match the rule with the `log` action. Packets intercepted by a capture link can be read by any `libpcap` tool such as `tcpdump` or `wireshark`. For more information, see [Securing the Network in Oracle Solaris 11.3](#).

Link Aggregations

Link aggregations enable you to pool multiple datalink resources that you administer as a single unit. You can improve the bandwidth and provide high availability for applications by combining multiple physical NICs together. Link aggregation of network datalinks ensures that a system has continuous access to the network. Trunk aggregation and DLMP aggregation are the two types of link aggregation.

Trunk aggregation provides consolidated bandwidth of the underlying datalinks for the clients configured over the aggregation. DLMP aggregation provides high availability across multiple switches for the clients configured over the aggregation. DLMP aggregation also supports link-based failure detection and probe-based failure detection to ensure continuous availability of the network to send and receive traffic. For more information about different types of link aggregation and procedures for configuring and administering link aggregations, see [Chapter 2, “Configuring High Availability by Using Link Aggregations”](#).

Virtual Local Area Networks

Virtual local area networks (VLANs) enable you to divide your network into subnetworks without having to add resources to the physical network environment. Therefore, the subnetworks are virtual and you use the same physical network resources. VLANs provide applications with isolated subnetworks so that only the applications in the same VLAN can communicate with each other. You can configure multiple virtual networks within a single network unit, for example, a switch by combining VLANs and Oracle Solaris zones. For more information about configuring and administering VLANs, see [Chapter 3, “Configuring Virtual Networks by Using Virtual Local Area Networks”](#).

Private Virtual Local Area Networks

The private virtual local area networks (PVLANS) standard defined in RFC 5517 enables you to divide a regular VLAN (primary) into sub-VLANs (secondary). A regular VLAN is a single broadcast domain. PVLAN partitions one broadcast domain into multiple subdomains. For more information about configuring and administering PVLANS, see [Chapter 4, “Configuring Private Virtual Local Area Networks”](#).

Bridged Networks

Bridges connect separate network segments, which are paths between two nodes. When connected by a bridge, the attached network segments communicate as if they were a single network segment. Bridges use a packet-forwarding mechanism to connect subnetworks together and enable a system to transmit packets to their destinations by using the shortest connection routes. For more information about bridged networks and procedures to administer bridges, see [Chapter 5, “Administering Bridging Features”](#).

Link Layer Discovery Protocol

Link Layer Discovery Protocol (LLDP) enables exchanging of connectivity and management information between the systems on the network for the purpose of topology discovery. The information can include system capabilities, management addresses, and other information relevant to network operations. The network diagnostics service uses LLDP to detect problems that might lead to limited or degraded network connectivity. For more information about LLDP and procedures to configure LLDP, see [Chapter 6, “Exchanging Network Connectivity Information With Link Layer Discovery Protocol”](#).

Data Center Bridging

Data center bridging (DCB) is used to manage the bandwidth, relative priority, and flow control of multiple traffic types when sharing the same network link, for example, when sharing a datalink between networking and storage protocols. DCB enables information exchange with the peers about the features that support converged network by using LLDP. The information is related to the configurations affecting the integrity of network packets especially in heavy traffic environments, such as data centers. DCB enables efficient network infrastructure by consolidating storage area network (SAN) and local area network (LAN) and thereby reducing operational and management costs in large deployments.

You can configure DCB features such as priority-based flow control (PFC) for the prevention of packet loss and enhanced transmission selection (ETS) for bandwidth sharing among packets based on class of service (CoS) priorities. For more information, see [Chapter 7, “Managing Converged Networks by Using Data Center Bridging”](#).

Configuring High Availability by Using Link Aggregations

This chapter provides an overview of link aggregations and describes procedures to configure and administer link aggregations.

This chapter contains the following topics:

- [“Overview of Link Aggregations” on page 17](#)
- [“Trunk Aggregations” on page 19](#)
- [“Datalink Multipathing Aggregations” on page 23](#)
- [“Requirements for Link Aggregations” on page 27](#)
- [“Creating a Link Aggregation” on page 28](#)
- [“Adding a Link to an Aggregation” on page 31](#)
- [“Removing a Link From an Aggregation” on page 33](#)
- [“Modifying a Trunk Aggregation” on page 33](#)
- [“Configuring Probe-Based Failure Detection for DLMP Aggregation” on page 34](#)
- [“Monitoring Probe-Based Failure Detection” on page 38](#)
- [“Deleting a Link Aggregation” on page 40](#)
- [“Switching Between Trunk and DLMP Aggregations” on page 41](#)
- [“Use Case: Configuring a Link Aggregation” on page 43](#)
- [“Comparing Trunk and DLMP Aggregation” on page 45](#)

Overview of Link Aggregations


Link aggregation is a method to combine several physical datalinks on a system to improve the bandwidth for applications. These physical datalinks are configured together as a single logical unit to increase the throughput of network traffic and to achieve high availability at the datalink layer.

You can include any link aggregation that is listed in the available NICs as an individual interface by using Oracle Enterprise Manager Ops Center. You can also display the details of both trunk and DLMP aggregation by using Oracle Enterprise Manager Ops Center. For more

information about the Oracle Enterprise Manager Ops Center, see <http://www.oracle.com/pls/topic/lookup?ctx=oc122&id=OPCCM>.

The following figure shows an example of a simple link aggregation configured on a system.

FIGURE 1 Link Aggregation Configuration



The illustration shows an aggregation `aggr1` that consists of three underlying datalinks, `net0`, `net1`, and `net2`. These datalinks are dedicated to serve the traffic that traverses the system through the aggregation. The underlying links are hidden from external applications. Instead, the logical datalink `aggr1` is accessible.

Benefits of Link Aggregation

Link aggregation has the following benefits:

- **Increased bandwidth** – The capacity of multiple links is combined into one logical link.
- **Automatic failover and failback** – The traffic from a failed link is automatically switched over to other working links in the aggregation, thereby achieving high availability.
- **Improved administration** – All underlying links are administered as a single unit.
- **Less drain on the network address pool** – The entire aggregation can be assigned one IP address.

Because link aggregation groups multiple links into a single logical datalink, features of datalinks such as link protection and resource management work well with link aggregations. For information about link protection, see [Chapter 1, “Using Link Protection in Virtualized Environments”](#) in *Securing the Network in Oracle Solaris 11.3*. For information about

resource management, see [Chapter 7, “Managing Network Resources” in *Managing Network Virtualization and Network Resources in Oracle Solaris 11.3*](#).

Link aggregation overcomes some of the problems that are encountered when using high availability features with network virtualization, for example IPMP. With link aggregation, you create the aggregation in a global zone first, then specify it as the underlying link when you configure the non-global zone. When the zone boots, it is assigned a virtual network interface card (VNIC) on the link aggregation. Unlike IPMP, which needs to be configured in every non-global zone, link aggregations are configured in global zones only and provide highly available VNICs to non-global zones. The global zone can also configure properties such as bandwidth on the VNICs assigned to the zone.

Note - Link aggregations perform similar functions to IP multipathing (IPMP) to improve network performance and availability at the datalink layer. For a comparison of these two technologies, see [Appendix A, “Link Aggregations and IPMP: Feature Comparison”](#).

The following types of link aggregations are supported:

- Trunk aggregations
- Datalink multipathing (DLMP) aggregations

For the differences between these two types of link aggregations, see [“Comparing Trunk and DLMP Aggregation” on page 45](#).

Trunk Aggregations

Trunk aggregations are based on the IEEE 802.3ad standard and work by enabling multiple flows of traffic to be spread across a set of aggregated ports. IEEE 802.3ad requires switch configuration, as well as switch-vendor proprietary extensions in order to work across multiple switches. In trunk aggregations, the clients configured over the aggregations get a consolidated bandwidth of the underlying links, because each network port is associated with every configured datalink over the aggregation. When you create a link aggregation, the aggregation is by default created in the trunk mode. You might use a trunk aggregation in the following situations:

- For systems in the network that run applications with distributed heavy traffic, you can dedicate a trunk aggregation to that application's traffic to take advantage of the increased bandwidth.
- For sites with limited IP address space that require large amounts of bandwidth, you need only one IP address for the trunk aggregation of datalinks.
- For sites that need to hide any internal datalinks, the IP address of the trunk aggregation hides these datalinks from external applications.
- For applications that need reliable network connection, trunk aggregation protects network connections against link failure.

Trunk aggregation supports the following features:


- Using a switch
- Using a switch with the Link Aggregation Control Protocol (LACP)
- Back-to-back trunk aggregation configuration
- Aggregation policies and load balancing

The following sections describe the features of trunk aggregations.

Using a Switch

Systems that are configured with trunk aggregations might use an external switch to connect to other systems. The following figure depicts a local network with two systems where each system has a trunk aggregation configured.

FIGURE 2 Trunk Aggregation Using a Switch




The two systems are connected by a switch. System A has a trunk aggregation that consists of two datalinks, net1 and net2. These datalinks are connected to the switch through aggregated ports. System B has a trunk aggregation of four datalinks, net1 through net4. These datalinks are also connected to the aggregated ports on the switch. In this trunk aggregation topology, the switch must support the IEEE 802.3ad standard and the switch ports must be configured for aggregation. See the switch manufacturer's documentation to configure the switch.

Back-to-Back Trunk Aggregation Configuration

Trunk aggregations support back-to-back configuration. Instead of using a switch, two systems are directly connected to run parallel aggregations, as shown in the following figure.

FIGURE 3 Back-to-Back Trunk Aggregation Configuration



The illustration shows trunk aggregation trunk0 on System A directly connected to trunk aggregation trunk0 on System B by means of the corresponding links between their respective underlying datalinks. This setup enables Systems A and B to provide redundancy, high availability, and high-speed communication between both systems. Each system also has net0 configured for traffic flow within the local network.

The most common application of a back-to-back trunk aggregation is the configuration of mirrored database servers in large deployments such as data centers. Both servers must be updated together and therefore require significant bandwidth, high-speed traffic flow, and reliability.

Using a Switch With the Link Aggregation Control Protocol

If your setup of a trunk aggregation has a switch and the switch supports LACP, you can enable LACP for the switch and the system. Oracle Switch ES1-24 supports LACP and for more information about configuring LACP, see [Sun Ethernet Fabric Operating System, LA Administration Guide](#). If you are using any other switch, see the switch manufacturer's documentation to configure the switch.

LACP enables a more reliable detection method of datalink failures. Without LACP, a link aggregation relies only on the link state reported by the device driver to detect the failure of an aggregated datalink. With LACP, LACPDUs are exchanged at regular intervals to ensure that the aggregated datalinks can send and receive traffic. LACP also detects some misconfiguration cases, for example, when the grouping of datalinks does not match between the two peers.

LACP exchanges special frames called Link Aggregation Control Protocol Data Units (LACPDUs) between the aggregation and the switch if LACP is enabled on the system. LACP uses these LACPDUs to maintain the state of the aggregated datalinks.

Use the `dladm create-aggr` command to configure an aggregation's LACP to one of the following three modes:

- `off` – The default mode for aggregations. The system does not generate LACPDUs.
- `active` – The system generates LACPDUs at specified intervals.
- `passive` – The system generates an LACPDU only when it receives an LACPDU from the switch. When both the aggregation and the switch are configured in passive mode, they do not exchange LACPDUs.

For information about how to configure LACP, see [“How to Create a Link Aggregation” on page 28](#).

Defining Aggregation Policies for Load Balancing

You can define a policy for the outgoing traffic that specifies how to distribute load across the available links of an aggregation, thus establishing load balancing. You can use the following load specifiers to enforce various load balancing policies:

- L2 – Determines the outgoing link by using the MAC (L2) header of each packet
- L3 – Determines the outgoing link by using the IP (L3) header of each packet
- L4 – Determines the outgoing link by using the TCP, UDP, or other ULP (L4) header of each packet

Any combination of these policies is also valid. The default policy is L4.

Datalink Multipathing Aggregations

Datalink multipathing (DLMP) aggregation is a type of link aggregation that provides high availability across multiple switches without requiring switch configuration. DLMP aggregation supports link-based failure detection and probe-based failure detection to ensure continuous availability of the network to send and receive traffic.

Advantages of DLMP Aggregation

DLMP aggregation provides the following advantages:

- The IEEE 802.3ad standard implemented by trunk aggregations does not have provisions to span multiple switches. Allowing failover between multiple switches in trunk mode requires vendor-proprietary extensions on the switches that are not compatible between vendors. DLMP aggregations allow failover between multiple switches without requiring any vendor-proprietary extensions.
- Using IPMP for high availability in the context of network virtualization is very complex. An IPMP group cannot be assigned directly to a zone. When network interface cards (NICs) have to be shared between multiple zones, you have to configure VNICs so that each zone gets one VNIC from each of the physical NICs. Each zone must group its VNICs into an IPMP group to achieve high availability. The complexity increases as you scale configurations, for example, in a scenario that includes large numbers of systems, zones, NICs, virtual NICs (VNICs), and IPMP groups. With DLMP aggregations, you create a VNIC or configure a zone's anet resource on top of the aggregation, and the zone sees a highly available VNIC.
- DLMP aggregation enables you to use the features of link layer such as link protection, user-defined flows, and the ability to customize link properties, such as bandwidth.

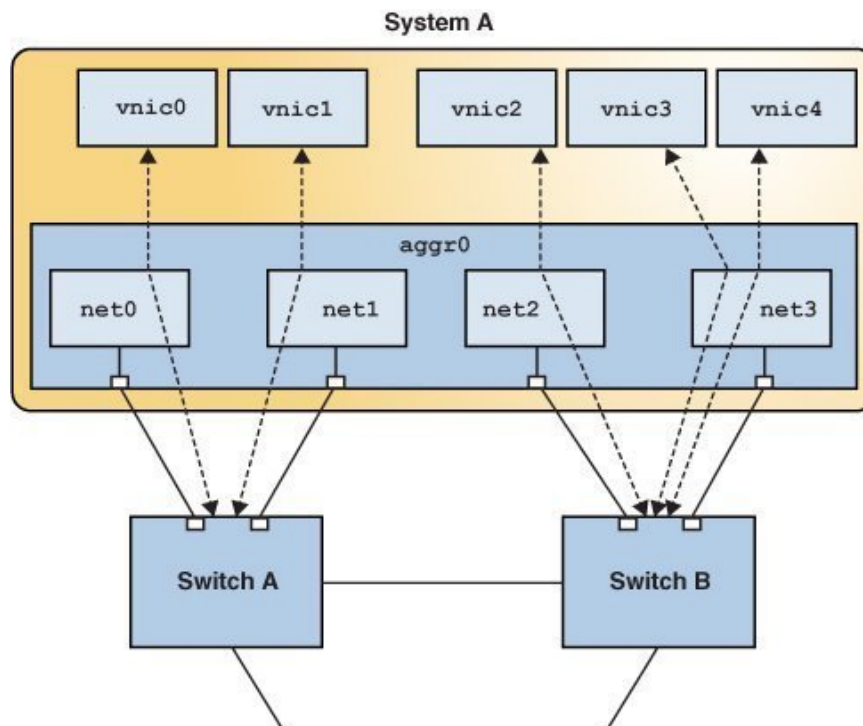
Note - You must not configure an IPMP group over a DLMP aggregation. However, you can configure an IPMP group over a trunk aggregation.

How DLMP Aggregation Works

In a trunk aggregation, each port is associated with every configured datalink over the aggregation. In a DLMP aggregation, a port is associated with any of the aggregation's configured datalinks.

The following figure shows how a DLMP aggregation works.

FIGURE 4 DLMP Aggregation



The figure shows System A with link aggregation `aggr0`. The aggregation consists of four underlying links, from `net0` to `net3`. VNICS `vnic0` through `vnic4` are also configured over the aggregation. The aggregation is connected to Switch A and Switch B, which in turn connect to other destination systems in the wider network.

VNICS are associated with aggregated ports through the underlying links. For example, in the figure, `vnic0` through `vnic3` are associated with the aggregated ports through the underlying

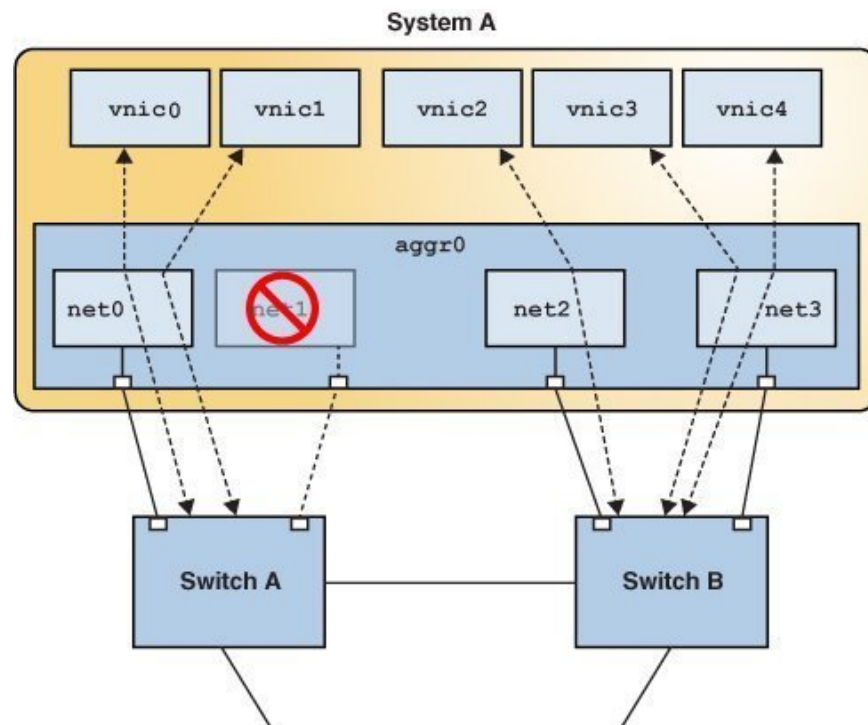
links net0 through net3. That is, if the number of VNICs and the number of underlying links are equal, then each port is associated with an underlying link.

If the number of VNICs exceeds the number of underlying links, then one port is associated with multiple datalinks. For example, in the figure the total number of VNICs exceeds the number of underlying links. Hence, vnic4 shares a port with vnic3.

When an aggregated port fails, all the datalinks that use that port are distributed among the other ports, thereby providing network high availability during failover. For example, if net0 fails, then DLMP aggregation shares the remaining port net1, between VNICs. The distribution among the aggregated ports occurs transparently to the user and independently of the external switches connected to the aggregation.

The following figure shows how DLMP aggregation works when a port fails. In the figure, net1 has failed and the link between switch and net1 is down. vnic1 shares a port with vnic0 through net0.

FIGURE 5 DLMP Aggregation When a Port Fails



Failure Detection in DLMP Aggregation

Failure detection in DLMP aggregation is a method to detect the failure of the aggregated ports. A port is considered to have failed when it cannot send or receive traffic. The port might fail because of the following reasons:

- Damage or cut in the cable
- Switch port goes down
- Failure in upstream network path

DLMP aggregation performs failure detection on the aggregated ports to ensure continuous availability of the network to send or receive traffic. When a port fails, the clients associated with that port are failed over to an active port. Failed aggregated ports remain unusable until they are repaired. The remaining active ports continue to function while any existing ports are deployed as needed. After the failed port recovers from the failure, clients from other active ports can be associated with it.

DLMP aggregation supports both link-based and probe-based failure detection.

Link-Based Failure Detection

Link-based failure detection detects failure when the cable is cut or when the switch port is down. It therefore can only detect failures caused by the loss of direct connection between the datalink and the first-hop switch. Link-based failure detection is enabled by default when a DLMP aggregation is created.

Probe-Based Failure Detection

Probe-based failure detection detects failures between an end system and the configured target IP addresses. This feature overcomes the known limitations of link-based failure detection. Probe-based failure detection is useful when a default router is down or when the network becomes unreachable. The DLMP aggregation detects failure by sending and receiving probe packets.

To enable probe-based failure detection in DLMP aggregation, you must configure the `probe-ip` property.

Note - In a DLMP aggregation, if no `probe-ip` is configured, then probe-based failure detection is disabled and only link-based failure detection is used.

Probe-based failure detection is performed by using the combination of two types of probes: Internet Control Message Protocol (ICMP (L3)) probes and transitive (L2) probes, which work together to determine the health of the aggregated physical datalinks.

- **ICMP Probing**

ICMP probing operates by sending ICMP packets to a target IP address. You can specify the source and target IP addresses manually or let the system choose them. If a source address is not specified, the source address is chosen from the IP addresses that are configured on the DLMP aggregation or on a VNIC configured on the DLMP aggregation. If a target address is not specified, the target address is chosen from one of the next hop routers on the same subnet as one of the specified source IP addresses. For more information, see [“How to Configure Probe-Based Failure Detection for DLMP” on page 35](#).

- **Transitive Probing**

Transitive probing is performed when the health state of all network ports cannot be determined by using only ICMP probing. This can happen if the source IP addresses chosen for ICMP probing are not on the list of addresses that are configured on a particular port. As a result, this port cannot receive ICMP probe replies. Transitive probing works by sending an L2 packet out of the ports without a configured source IP address to another port that has a configured source IP address. If the other port can reach the ICMP target and the ports sending the transitive probe receive the L2 reply from the other port, it means that the port sending the transitive probe can also reach the ICMP target.

In a DLMP aggregation, you can optionally specify a VLAN that can be used for both ICMP and transitive probing to support configurations where untagged traffic is restricted. You can specify a single VLAN ID by using the `probe-vlan-id` property and the specified VLAN will be used for both sending the ICMP probes to the target IP and to check the health of the member ports that depend on the transitive probes. The L3 probe target will be automatically selected and validated based on this property. You can provide a value between 0 and 4094 for the `probe-vlan-id` property. The default value is 0, which indicates that the probes are untagged.

Oracle Solaris includes proprietary protocol packets for transitive probes that are transmitted over the network. For more information, see [Appendix B, “Packet Format of Transitive Probes”](#).

Requirements for Link Aggregations

The link aggregation configuration has the following requirements:

- The datalinks to be configured into an aggregation should not have any IP interface configured over them.
- You can create link aggregations only from the global zone. You cannot use datalinks to create a link aggregation from a non-global zone even if you do not have any IP interface configured over the datalinks. Link aggregation only combines multiple physical NICs

but in a non-global zone all interfaces are virtual. Therefore, you cannot create a link aggregation from the non-global zone.

- All the datalinks in the aggregation must run at the same speed and in full-duplex mode.
- For DLMP aggregations, you must have at least one switch to connect the aggregation to the ports in the other systems. You cannot use a back-to-back setup when configuring DLMP aggregations.
- On SPARC based systems, each datalink must have its own unique MAC address. For information, see [“How to Ensure That the MAC Address of Each Interface Is Unique” in *Configuring and Managing Network Components in Oracle Solaris 11.3*](#).
- (Trunk aggregation only) Devices must support link state notification as defined in the IEEE 802.3ad Link Aggregation Standard in order for a port to attach to a trunk aggregation or to detach from a trunk aggregation. Devices that do not support link state notification can be aggregated only by using the -f option of the `dladm create-aggr` command. For such devices, the link state is always reported as UP. For information about the use of the -f option, see [“How to Create a Link Aggregation” on page 28](#).

Creating a Link Aggregation

Link aggregation groups the underlying ports into a single logical group. The aggregation uses these underlying ports exclusively and you cannot perform any other operation such as configuring VNICs or assigning IP addresses on these ports. However, you can configure VNICs on top of the aggregation and not on the individual ports.

You must remove any existing IP interface on these ports before creating the link aggregation.

You can also configure a VLAN and create VNICs over the link aggregation that you have created. For information about how to create VLANs over link aggregations, see [“How to Configure VLANs Over a Link Aggregation” on page 57](#).

Note - Link aggregation works only on full-duplex, point-to-point links that operate at identical speeds. Make sure that the datalinks in your aggregation conform to this requirement.

▼ How to Create a Link Aggregation

Before You Begin If you are creating a trunk aggregation and are using a switch in the aggregation, configure the ports to be used as an aggregation on the switch. If the switch supports LACP, configure LACP in either active or passive mode.

See the switch manufacturer's documentation to configure the switch.

1. Become an administrator.

For more information, see [“Using Your Assigned Administrative Rights” in *Securing Users and Processes in Oracle Solaris 11.3*](#).

2. Display the datalink information to identify the physical datalinks for the aggregation.

```
# dladm show-phys
```

3. Ensure that the datalink that you intend to aggregate are not in use by any application.

For example, if an IP interface is created over the datalink, remove the IP interface first.

a. Determine the link state.

```
# ipadm show-if
```

IFNAME	CLASS	STATE	ACTIVE	OVER
lo0	loopback	ok	yes	--
net0	ip	ok	no	--

The output indicates that an IP interface exists over the datalink net0.

b. Remove the IP interface.

```
# ipadm delete-ip interface
```

where *interface* specifies the IP interface over the link. For more information, see the [ipadm\(1M\)](#) man page.

4. Create a link aggregation.

```
# dladm create-aggr [-f] [-m mode] [-P policy] [-L LACP-mode] \
[-T time] [-u address] -l link1 -l link2 [...] aggr
```

-f Forces the creation of the aggregation. Use this option when you are attempting to aggregate devices that do not support link state notification.

-m mode Mode must be set to one of the following values. The default mode is trunk.

- trunk – IEEE 802.3ad compliant link aggregation mode
- dmp – Datalink multipathing mode

-P policy (Trunk aggregation only) Specifies the load balancing policy for the aggregation. Supported values are L2, L3, and L4. For more information, see [“Defining Aggregation Policies for Load Balancing” on page 22](#).

<code>-L LACP-mode</code>	(Trunk aggregation only) Specifies the mode of LACP if it is used. Supported values are <code>off</code> , <code>active</code> , or <code>passive</code> . For information about the modes, see “Using a Switch” on page 20 .
<code>-T time</code>	(Trunk aggregation only) Specifies the LACP timer value. The supported values are <code>short</code> or <code>long</code> .
<code>-u address</code>	Specifies the fixed unicast address for the aggregation.
<code>-l linkn</code>	Specifies the datalinks that you want to aggregate.
<code>aggr</code>	Specifies the name of the aggregation, which can be any customized name. For information about the rules to assign names, see “Rules for Valid Link Names” in <i>Configuring and Managing Network Components in Oracle Solaris 11.3</i> .

5. (Optional) Check the status of the aggregation that you created.

- Display the aggregations and links with the status information.

```
# dladm show-link
```

- Display the aggregations with the status and per port information.

```
# dladm show-aggr -x
```

The state of the aggregation should be up.

Example 1 Creating a Trunk Aggregation

This example shows the commands to create a link aggregation with two underlying datalinks, `net0` and `net1`. The aggregation is also configured to transmit LACP packets. The example begins with the removal of existing IP interfaces over the underlying datalinks.

```
# ipadm show-if
IFNAME      CLASS      STATE      ACTIVE      OVER
lo0         loopback   ok         yes         --
net0        ip         ok         no          --
# ipadm delete-ip net0
# dladm create-aggr -L active -l net0 -l net1 trunk0
# dladm show-aggr -x
LINK        PORT          SPEED DUPLEX  STATE      ADDRESS          PORTSTATE
trunk0      --            1000Mb full  up         8:0:27:49:10:b8  --
              net0          1000Mb full  up         8:0:27:49:10:b8  attached
              net1          1000Mb full  up         8:0:27:e4:d9:46  attached
```

Example 2 Creating a DLMP Aggregation and Configuring an IP Interface on Top of the Aggregation

This example shows how to create a DLMP aggregation. The aggregation has three underlying datalinks `net0`, `net1`, and `net2`. An IP interface is created on top of the aggregation `aggr0` and a VNIC `vnic1` is created on top of it.

```
# dladm create-aggr -m dlmp -l net0 -l net1 -l net2 aggr0
# dladm show-link
LINK      CLASS  MTU    STATE  OVER
net0      phys   1500   up     --
net1      phys   1500   up     --
net2      phys   1500   up     --
aggr0     aggr   1500   up     net0 net1 net2
# dladm show-aggr -x
LINK      PORT      SPEED DUPLEX  STATE  ADDRESS          PORTSTATE
aggr0     --        1000Mb full  up     8:0:27:49:10:b8  --
          net0      1000Mb full  up     8:0:27:49:10:b8  attached
          net1      1000Mb full  up     8:0:27:e4:d9:46  attached
          net2      1000Mb full  up     8:0:27:38:7a:97  attached
# ipadm create-ip aggr0
# ipadm create-addr -a local=10.10.10.1 aggr0/v4
# dladm create-vnic -l aggr0 vnic1
```

Next Steps You can perform further configuration of the aggregation such as creating IP interfaces and VNICs. You can use the created aggregation for configuring both non-global zones and kernel zones.

- For information about creating IP interfaces, see [Chapter 3, “Configuring and Administering IP Interfaces and Addresses in Oracle Solaris” in *Configuring and Managing Network Components in Oracle Solaris 11.3*](#).
- For information about configuring VLANs over a link aggregation, see [“Configuring VLANs Over a Link Aggregation” on page 57](#).
- For information about configuring VNICs, see [“How to Configure VNICs and Etherstubs” in *Managing Network Virtualization and Network Resources in Oracle Solaris 11.3*](#).
- For information about configuring zones, see [Creating and Using Oracle Solaris Zones](#).

Adding a Link to an Aggregation

You can include additional datalinks to an existing aggregation. If you are adding datalinks to a trunk aggregation, you might have to reconfigure the switch to accommodate the additional datalinks even though LACP is configured on the switch.

▼ How to Add a Link to an Aggregation

1. Become an administrator.

For more information, see [“Using Your Assigned Administrative Rights” in *Securing Users and Processes in Oracle Solaris 11.3*](#).

2. Ensure that no IP interfaces are configured over the link. If any IP interface is configured, delete that IP interface.

```
# ipadm show-if
# ipadm delete-ip interface
```

where *interface* is the IP interface configured over the datalink.

3. Add the link to the aggregation.

```
# dladm add-aggr -l link [-l link] [...] aggr
```

where *link* represents a datalink that you are adding to the aggregation and *aggr* is the name of the aggregation.

4. (Trunk aggregations only) If necessary, reconfigure the switch.

You might need to reconfigure the switch to accommodate the additional datalinks based on how the switch is configured, even if LACP is enabled on the switch.

See the switch manufacturer's documentation to perform any reconfiguration tasks on the switch.

Example 3 Adding a Link to an Aggregation

This example shows how to add a link to the aggregation `aggr0`.

```
# dladm show-link
LINK      CLASS    MTU      STATE    OVER
net0      phys     1500     up       --
net1      phys     1500     up       --
aggr0     aggr     1500     up       net0 net1
net3      phys     1500     up       --
# ipadm delete-ip net3
# dladm add-aggr -l net3 aggr0
# dladm show-link
LINK      CLASS    MTU      STATE    OVER
net0      phys     1500     up       --
net1      phys     1500     up       --
aggr0     aggr     1500     up       net0 net1 net3
net3      phys     1500     up       --
```


Removing a Link From an Aggregation

You can remove the individual datalinks associated with an aggregation by using the `dladm remove-aggr` command. You need to reconfigure the switch when you remove links from the aggregation.

Become an administrator and use the following command:

```
# dladm remove-aggr -l link aggr
```

EXAMPLE 4 Removing a Link From an Aggregation

This example shows how to remove a link from the aggregation `aggr0`.

```
# dladm show-link
LINK    CLASS    MTU    STATE    OVER
net0    phys      1500   up       --
net1    phys      1500   up       --
aggr0    aggr      1500   up       net0 net1 net3
net3    phys      1500   up       --
# dladm remove-aggr -l net3 aggr0
# dladm show-link
LINK    CLASS    MTU    STATE    OVER
net0    phys      1500   up       --
net1    phys      1500   up       --
aggr0    aggr      1500   up       net0 net1
net3    phys      1500   unknown  --
```

Modifying a Trunk Aggregation

You can modify selected attributes such as policy, lacpmode, and time for the trunk aggregation. These attributes are not supported in DLMP aggregations.

- To modify the load balancing policy of the aggregation, become an administrator and issue the following command:

```
# dladm modify-aggr -P policy aggr
```

policy Represents one or more of the load balancing policies L2, L3, and L4, as explained in [“Defining Aggregation Policies for Load Balancing”](#) on page 22.

aggr Specifies the aggregation whose policy you want to modify.

- To modify the LACP mode of the aggregation, become an administrator and use the following command:

```
# dladm modify-aggr -L LACP-mode -T time aggr
```

-L *LACP-mode* Indicates the LACP mode in which the aggregation must operate. Possible values are active, passive, and off.

-T *time* Indicates the LACP timer value, either short or long.

aggr Specifies the aggregation whose LACP mode you want to modify.

Note - You can modify the load balancing policy, the LACP mode, and the LACP timer value of an aggregation even when the aggregation is active or in use. You can also modify these attributes even if an IP interface and a VNIC is configured over the aggregation.

EXAMPLE 5 Modifying a Trunk Aggregation

This example shows how to modify the load balancing policy of link aggregation `aggr0` to L2 and to change the LACP mode to active.

```
# dladm modify-aggr -P L2 aggr0
# dladm modify-aggr -L active -T short aggr0
# dladm show-aggr
LINK      MODE      POLICY  ADDRPOLICY      LACPACTIVITY  LACPTIMER
aggr0     trunk      L2      auto             active         short
```

Configuring Probe-Based Failure Detection for DLMP Aggregation

You must configure the `probe-ip` property for a DLMP aggregation to enable probing. Otherwise, probing is disabled by default and only link-based failure detection is used. For more information, see [“Probe-Based Failure Detection” on page 26](#).

You can use the following datalink properties with the `dladm set-linkprop` command to configure probe-based failure detection:

- `probe-ip` – Specifies a comma separated list of source IP address, target IP addresses, or source and target IP addresses. These addresses will be used for ICMP probing.

Note - Specifying the source and the target IP addresses is optional. If you do not specify the source and the IP target addresses, the source and the probing IP addresses are automatically selected.

You can use + to separate the source from the target. You can specify the targets as an IP address or the host name of the target. For information about how to specify the source address and the target address, see [“How to Configure Probe-Based Failure Detection for DLMP” on page 35](#).

- `probe-vlan-id` – Specifies the VLAN ID to be used for both ICMP and transitive probing. Valid values are from 0 to 4094. The value 0 indicates that the transitive probes are untagged. The default value is 0.
- `probe-fdt` – Specifies the failure detection time. You can configure the expected failure detection time value in seconds. The default value is 10 seconds.

▼ How to Configure Probe-Based Failure Detection for DLMP

Before You Begin Create a DLMP aggregation. For more information, see [“How to Create a Link Aggregation” on page 28](#).

1. Become an administrator.

For more information, see [“Using Your Assigned Administrative Rights” in *Securing Users and Processes in Oracle Solaris 11.3*](#).

2. (Optional) Display all the existing aggregations to identify the aggregation for configuring probe-based failure detection.

```
# dladm show-aggr
```

3. Configure probing for the aggregation by either using automatically selected source and target IP address or by manual selection.

■ **To use automatic selection, issue the following command:**

If you do not specify target addresses, the DLMP aggregation chooses a target IP automatically from one of the next hop routers that are on the same subnet as the source IP address.

```
# dladm set-linkprop -p probe-ip=+ aggr
```

If the source IP address is not specified, all the IP addresses configured on the aggregation and the VNICs are the potential source IP addresses of the ICMP probes. Only IP

interfaces configured in the global zone are used as the source addresses. For example, IP addresses in any non-global zone that is using the DLMP aggregation are not used as a source address.

■ **To use manual selection, issue the following command:**

```
# dladm set-linkprop -p probe-ip=[source[,...]]+[target[,...]] aggr
```

<i>source</i>	<p>Specifies the source IP address for probing, which is optional. You can format the source IP address in any one of the following ways:</p> <ul style="list-style-type: none">■ <code>probe-ip=IP-address[/prefix-length]+</code> IP address and its prefix length. For example, <code>10.130.10.1/24+</code>.■ <code>probe-ip=addr-obj-name+</code> Address object name. For example, <code>vnic1/addr1+</code>.■ <code>probe-ip=interface-name+</code> Interface name. The name of the aggregated interface or any VNIC that is configured over the aggregation. For example, <code>[aggr1]</code>. The interface must be an IP interface with a configured IP address.■ <code>probe-ip=hostname+</code> Host name. For example, <code>sdg1+</code>
<i>target</i>	<p>Specifies the target IP address for probing, which is optional. When you specify the target IP address, it must be in any one of the following ways:</p> <ul style="list-style-type: none">■ <code>probe-ip=+IP-address</code> For example, <code>+10.130.10.1</code>.■ <code>probe-ip=+hostname</code> For example, <code>+sdg1</code>
<i>aggr</i>	<p>Name of the aggregation.</p>

Note - It is recommended that you do not specify target IP addresses explicitly unless you have a particular address that you want to use as a target. If you do not specify target addresses, the DLMP aggregation chooses a target IP automatically from one of the next hop routers that are on the same subnet as the source IP address. For more information about how to determine or configure next hop router, see [Chapter 1, “Administering TCP/IP Networks” in *Administering TCP/IP Networks, IPMP, and IP Tunnels in Oracle Solaris 11.3*](#).

4. **(Optional) Set the VLAN ID for the probes.**

```
# dladm set-linkprop -p probe-vlan-id=probe-VID aggr
```

where *probe-VID* is the VLAN ID to be used for both ICMP and transitive probes. The valid values of *probe-VID* ranges from 0 to 4094. The *probe-VID* value 0 indicates that the transitive probes are untagged. The default value is 0.

5. (Optional) Set the failure detection time.

```
# dladm set-linkprop -p probe-fdt=fdt aggr
```

where *fdt* is the failure detection time specified in seconds. The default value is 10 seconds.

6. (Optional) Display the probe statistics.

```
# dlstat show-aggr -n -P [[t],[i],[all]]
```

where *-P* specifies the type of the probe to display. You can specify any of the following in a comma-separated list as the probe type:

- *t* – Displays the transitive probes.
- *i* – Displays the ICMP probes.
- *all* – Displays both the transitive and ICMP probes.

Example 6 Configuring Probe-Based Failure Detection

This example shows how you can configure probe-based failure detection by using the default automatic selection of source and probing IP addresses. To use the default automatic selection, you need to set the *ip-probe* property of the DLMP aggregation to +.

1. Create a DLMP aggregation.

```
# dladm create-aggr -m dlmp -l net0 -l net1 -l net2 aggr1
# dladm show-aggr -x
```

LINK	PORT	SPEED	DUPLEX	STATE	ADDRESS	PORTSTATE
<i>aggr0</i>	--	1000Mb	full	up	8:0:27:49:10:b8	--
	<i>net0</i>	1000Mb	full	up	8:0:27:49:10:b8	attached
	<i>net1</i>	1000Mb	full	up	8:0:27:e4:d9:46	attached
	<i>net2</i>	1000Mb	full	up	8:0:27:38:7a:97	attached

2. Create an IP interface over the aggregation.

```
# ipadm create-ip aggr1
```

3. Assign an IP address to the interface.

```
# ipadm create-addr -a 192.168.85.137/24 aggr1
```

4. Configure probe-based failure detection.

```
# dladm set-linkprop -p probe-ip=+ aggr1
```

Since the source and the target IP addresses are not specified, the IP address 192.168.85.137 configured on the aggregation, `aggr1`, becomes the source IP address of the ICMP probes.

5. Set the VLAN ID for the probes.

```
# dladm set-linkprop -p probe-vlan-id=2100 aggr1
```

6. Set the failure detection time.

```
# dladm set-linkprop -p probe-fdt=15 aggr1
```

7. Display the properties that are set.

```
# dladm show-linkprop -p probe-ip,probe-fdt aggr1
```

LINK	PROPERTY	PERM	VALUE	EFFECTIVE	DEFAULT	POSSIBLE
aggr1	probe-ip	rw	+	+	--	--
aggr1	probe-fdt	rw	15	15	10	1-600

8. Display the probe statistics of the aggregation.

```
# dlstat show-aggr -n -P t,i aggr1
```

TIME	AGGR	PORT	LOCAL	TARGET	PROBE	NETRTT	RTT
0.45s	aggr1	net0	net0	net1	t16148	--	--
0.45s	aggr1	net0	net0	net1	t16148	0.63ms	0.81ms
1.08s	aggr1	net1	net1	net0	t16148	--	--
1.08s	aggr1	net1	net1	net0	t16148	0.72ms	0.99ms
2.07s	aggr1	net1	192.168.85.137	192.168.85.137	i15535	--	--
2.07s	aggr1	net1	192.168.85.137	192.168.85.137	i15535	0.18ms	0.54ms

For information about each of the probe statistics field displayed in the output, see [Example 7, “Displaying Probe-Related Information,” on page 39](#).

Monitoring Probe-Based Failure Detection

You can monitor probe-based failure detection by using the `dladm show-aggr`, `dlstat show-aggr`, and `ipadm show-addr` commands.

You can use the `dlstat show-aggr` command with the following syntax to display the probe-related information:

```
# dlstat show-aggr -n -P [[t],[i],[all]]
```

where `-P` specifies the probe type to display. You can specify any of the following in a comma-separated list as the probe type:

- `t` – Displays the transitive probes.

- `i` – Displays ICMP probes.
- `all` – Displays both the transitive and ICMP probes.

EXAMPLE 7 Displaying Probe-Related Information

The following example displays the statistics of the probes for the DLMP aggregation, `aggr1`.

```
# dlstat show-aggr -n -P t,i aggr1
TIME      AGGR      PORT  LOCAL      TARGET      PROBE  NETRTT      RTT
0.53s     aggr1     net0   net0        net1        t16148  --          --
0.53s     aggr1     net0   net0        net1        t16148  0.62ms     0.87ms
1.17s     aggr1     net1   net1        net0        t16148  --          --
1.17s     aggr1     net1   net1        net0        t16148  0.72ms     0.99ms
2.24s     aggr1     net1   192.168.0.1 192.168.0.2 i15535   --          --
2.24s     aggr1     net1   192.168.0.1 192.168.0.2 i15535   0.11ms     0.55ms
```

TIME Time at which the probe is sent in seconds. This time is relative to the time when you issue the `dlstat` command. If the probe is sent before you issue the `dlstat` command, the time is negative.

AGGR Aggregation name for which the probe is sent.

LOCAL ICMP probes: Source IP address of the probes.
Transitive probes: Port name from which the transitive probe originated.

TARGET ICMP probes: Destination IP address of the probes.
Transitive probes: Port name of the probe that is targeted.

PROBE Identifier number representing the probe. The prefix `t` is for transitive probes and the prefix `i` is for ICMP probes.

NETRTT Network round-trip-time for the probe. This value is the time period between sending the probe and receiving the acknowledgment by the DLMP aggregation.

RTT Total round-trip-time for the probe. This value is the time period between sending the probe and completing the process of acknowledgment by the DLMP aggregation.

For more information, see the [dlstat\(1M\)](#) man page.

EXAMPLE 8 Displaying Detailed Information About the Aggregated Port

The following example displays the detailed aggregation information on each underlying port.

```
# dladm show-aggr -x
LINK      PORT      SPEED  DUPLEX  STATE  ADDRESS          PORTSTATE
aggr1     --        100Mb  full    up     1e:34:db:fa:50:a2 --
          net0    100Mb  full    up     1e:34:db:fa:50:a2 attached
          net1    100Mb  full    up     b2:c0:6a:3e:c5:b5 attached
```

For more information, see the [dladm\(1M\)](#) man page.

EXAMPLE 9 Displaying the State of the Aggregated Ports

The following example shows the state of the ports of the aggregation and the target IP addresses of the ports.

```
# dladm show-aggr -S -n
LINK      PORT      FLAGS  STATE  TARGETS  XTARGETS
aggr1     net1      u--3   active  192.168.0.2  net0
--        net0      u-2-   active  --         net1
```

EXAMPLE 10 Displaying probe-ip Property Values

The following example displays the details about the link property probe-ip for the specified DLMP aggregation.

```
# dladm show-linkprop -p probe-ip aggr1
LINK      PROPERTY  PERM  VALUE          EFFECTIVE  DEFAULT  POSSIBLE
aggr1     probe-ip  rw    192.168.0.2    192.168.0.2  --       --
```

EXAMPLE 11 Displaying the IP Address and the State of the Aggregation

The following example displays the IP address and the state of the aggregation.

```
# ipadm show-addr aggr1
ADDROBJ      TYPE  STATE  ADDR
aggr1/local1 static ok     192.168.0.1/24
```

Deleting a Link Aggregation

You can delete a link aggregation by using the `dladm delete-aggr` command. You must remove the IP interface and VNICs configured over the link aggregation before deleting the aggregation.

Note - When you remove a datalink by using the `dladm delete-phys` command, any aggregation (a layer 2 component) that is configured over the datalink is also deleted. This `dladm delete-phys` command also deletes the other layer 2 components such as flows, VNICs, and VLANs configured over the datalink. However, you must manually delete the layer 3 components configured over the datalink such as IP interfaces by using the `ipadm` command.

▼ How to Delete a Link Aggregation

1. **Become an administrator.**

For more information, see [“Using Your Assigned Administrative Rights” in *Securing Users and Processes in Oracle Solaris 11.3*](#).

2. **Delete the IP interface that is configured over the link aggregation.**

```
# ipadm delete-ip IP-aggr
```

where *IP-aggr* is the IP interface over the link aggregation.

3. **Delete the link aggregation.**

```
# dladm delete-aggr aggr
```

Example 12 Deleting a Link Aggregation

This example shows how to delete the aggregation `aggr0`. The deletion is persistent.

```
# ipadm delete-ip aggr0
# dladm delete-aggr aggr0
```

Switching Between Trunk and DLMP Aggregations

Switching between a trunk aggregation and a DLMP aggregation changes the entire configuration, so it affects the aggregation in a more comprehensive way than simply modifying the other link aggregation properties. To switch from one mode of the aggregation to the other when the aggregation is in use, you must remove all the existing configurations of the aggregation and then reconfigure the aggregation in the desired mode.

▼ How to Switch Between Link Aggregation Types

- Before You Begin**
- If you switch from a trunk aggregation to a DLMP aggregation, you must remove the switch configuration that was previously created for the trunk aggregation.
 - If you switch from a DLMP aggregation, you must ensure that all the links are on the same switch and the switch is configured for the aggregation.

1. Become an administrator.

For more information, see [“Using Your Assigned Administrative Rights” in *Securing Users and Processes in Oracle Solaris 11.3*](#).

2. Determine the current type of link aggregation.

```
# dladm show-aggr
```

The output `MODE` field indicates the current type of the aggregation. The value of `MODE` is `trunk` for a trunk aggregation and `dlmp` for a DLMP aggregation.

3. Switch the aggregation.

```
# dladm modify-aggr -m mode aggr
```

where *mode* is `trunk` if you are switching to a trunk aggregation or `dlmp` if you are switching to a DLMP aggregation, and *aggr* is the name of the aggregation.

4. Configure the switch according to the requirements of the new type of link aggregation.

See the switch manufacturer's documentation for configuring the switch.

5. (Optional) Verify the current link aggregation configuration.

```
# dladm show-aggr
```

Example 13 Switching From a Trunk Aggregation to a DLMP Aggregation

This example shows how to change an aggregation from a trunk aggregation to a DLMP aggregation.

```
# dladm show-aggr
LINK    MODE    POLICY  ADDRPOLICY    LACPACTIVITY  LACPTIMER
aggr0   trunk    L2      auto          active        short

# dladm modify-aggr -m dlmp aggr0
# dladm show-aggr
LINK    MODE    POLICY  ADDRPOLICY    LACPACTIVITY  LACPTIMER
aggr0   dlmp    --      --            --            --
```

Once the aggregation is switched, you must remove the previous switch configuration that was applied to the trunk aggregation.

Use Case: Configuring a Link Aggregation

The following end-to-end use case shows how to accomplish the following actions:

- Create a DLMP aggregation.
- Add links to the aggregation.
- Configure an IP interface over the aggregation.
- Configure a VNIC over the aggregation.
- Configure probe-based failure detection for the aggregation.
- Configure the target IP address in the routing table.
- Monitor the ICMP and transitive probes.

1. Become an administrator.

For more information, see [“Using Your Assigned Administrative Rights” in *Securing Users and Processes in Oracle Solaris 11.3*](#).

2. Display datalink information to identify the datalinks for aggregation.

```
# dladm show-link
LINK      CLASS      MTU      STATE    OVER
net0      phys       1500     up       --
net1      phys       1500     up       --
net2      phys       1500     up       --
```

3. Ensure that the datalinks that you want aggregate do not have IP interfaces configured over the link. Delete the interface if any interface is configured on any of the links.

```
# ipadm show-if
IFNAME     CLASS      STATE     ACTIVE    OVER
lo0        loopback   ok        yes       --
net0       ip         ok        no        --
# ipadm delete-ip net0
```

4. Create a DLMP aggregation with the links net0 and net1.

```
# dladm create-aggr -m dlmp -l net0 -l net1 aggr1
```

5. Add another link, net2, to the aggregation.

```
# dladm add-aggr -l net2 aggr1
```

Reconfigure the switch to accommodate the new links if the existing switch configuration requires it. See the switch manufacturer's documentation.

6. Configure an IP interface on top of the aggregation aggr1.

```
# ipadm create-ip aggr1
# ipadm create-addr -a local=10.10.10.1 aggr1/v4
```

7. Create a VNIC on top of the aggregation.

```
# dladm create-vnic -l aggr1 vnic1
```

8. Configure probe-based failure detection for the aggregation.

```
# dladm set-linkprop -p probe-ip+= aggr1
```

Since the source and the target addresses are not specified, they will be automatically chosen.

9. Display the state of the aggregated ports and the targets.

```
# dladm show-aggr -S
```

LINK	PORT	FLAGS	STATE	TARGETS	XTARGETS
aggr1	net0	u--3	active	10.10.10.2	net2 net1
--	net1	u-2-	active	--	net2 net0
--	net2	u-2-	active	--	net0 net1

10. Monitor the ICMP probe statistics.

```
# dlstat show-aggr -n -P i
```

TIME	AGGR	PORT	LOCAL	TARGET	PROBE	NETRTT	RTT
1.16s	aggr1	net0	10.10.10.1	10.10.10.2	i33	--	--
1.16s	aggr1	net0	10.10.10.1	10.10.10.2	i33	0.08ms	0.33ms
2.05s	aggr1	net0	10.10.10.1	10.10.10.2	i34	--	--
2.05s	aggr1	net0	10.10.10.1	10.10.10.2	i34	0.01ms	0.64ms
4.05s	aggr1	net0	10.10.10.1	10.10.10.2	i35	--	--
4.05s	aggr1	net0	10.10.10.1	10.10.10.2	i35	0.10ms	0.35ms
5.54s	aggr1	net0	10.10.10.1	10.10.10.2	i36	--	--
5.54s	aggr1	net0	10.10.10.1	10.10.10.2	i36	0.08ms	0.34ms

11. Monitor the transitive probe statistics between the ports.

```
# dlstat show-aggr -n -P t
```

TIME	AGGR	PORT	LOCAL	TARGET	PROBE	NETRTT	RTT
0.30s	aggr1	net2	net2	net0	t38	--	--
0.30s	aggr1	net2	net2	net0	t38	0.46ms	0.59ms
0.46s	aggr1	net0	net0	net1	t39	--	--
0.46s	aggr1	net0	net0	net1	t39	0.46ms	0.50ms
0.48s	aggr1	net1	net1	net0	t39	--	--
0.48s	aggr1	net1	net1	net0	t39	0.34ms	0.38ms
0.72s	aggr1	net2	net2	net1	t38	--	--
0.72s	aggr1	net2	net2	net1	t38	0.38ms	0.42ms
0.76s	aggr1	net0	net0	net2	t39	--	--
0.76s	aggr1	net0	net0	net2	t39	0.33ms	0.38ms
0.87s	aggr1	net1	net1	net2	t39	--	--
0.87s	aggr1	net1	net1	net2	t39	0.32ms	0.38ms

1.95s	aggr1	net2	net2	net0	t39	--	--
1.95s	aggr1	net2	net2	net0	t39	0.36ms	0.42ms
1.97s	aggr1	net2	net2	net1	t39	--	--
1.97s	aggr1	net2	net2	net1	t39	0.32ms	0.38ms
1.99s	aggr1	net0	net0	net1	t40	--	--
1.99s	aggr1	net0	net0	net1	t40	0.31ms	0.36ms
2.12s	aggr1	net1	net1	net0	t40	--	--
2.12s	aggr1	net1	net1	net0	t40	0.34ms	0.40ms
2.14s	aggr1	net0	net0	net2	t40	--	--

The aggregation `aggr0` with an IP interface configured over it is created. The VNIC `vnic1` is configured on top of the aggregation `aggr0`. Probe-based failure detection is configured without specifying either the source IP address or the target IP address of the probes. To enable probing, the target in the routing table is configured with an IP address, `10.10.10.2`, on the same subnet as the specified IP address, `10.10.10.1`. The ICMP and transitive probe statistics are monitored.

Comparing Trunk and DLMP Aggregation

This section presents a general comparison of the two types of link aggregation.

TABLE 1 Feature Comparison of Trunk and DLMP Aggregation

Feature	Trunk Aggregations	DLMP Aggregations
Link-based failure detection	Supported	Supported
LACP	Supported	Not supported
Use of standby interfaces	Not supported	Not supported ¹
Span multiple switches	Not supported unless using vendor proprietary solution	Supported
Switch configuration	Required	Not required
Policies for load balancing	Supported	Not applicable
Load spreading across all of the aggregation's ports	Supported	Limited ²
User-defined flows for resource management	Supported	Supported
Link protection	Supported	Supported
Back-to-back configuration	Supported	Not supported ³

¹ Each DLMP client is associated with exactly one DLMP port. The remaining ports act as available ports for the DLMP clients, but you cannot configure these available ports.

² The aggregation spreads its VNICs across all ports. However, individual VNICs cannot spread the load on multiple ports.

³ DLMP aggregations must always use an intermediary switch to send packets to other destination systems. However, no switch configuration is required for DLMP.

Configuring Virtual Networks by Using Virtual Local Area Networks

This chapter discusses about virtual local area networks (VLANs) and also describes the procedures to configure and modify VLANs.

This chapter contains the following topics:

- “Overview of Deploying VLANs” on page 47
- “Using VLANs With Virtualization” on page 51
- “Planning a VLAN Configuration” on page 52
- “Configuring a VLAN” on page 52
- “Configuring VLANs Over a Link Aggregation” on page 57
- “Configuring VLANs on a Legacy Device” on page 59
- “Displaying VLAN Information” on page 60
- “Modifying VLANs” on page 61
- “Deleting a VLAN” on page 63
- “Use Case: Combining Link Aggregations and VLAN Configurations” on page 64
- “Use Case: Configuring Flows for a VLAN” on page 67
- “Use Case: Configuring Bandwidth for VLAN VNICs” on page 68

Overview of Deploying VLANs

A virtual local area network (VLAN) is a subdivision of a local area network at the datalink layer of the protocol stack. You can create VLANs for local area networks that use switch technology. You can assign interfaces on the same system to different VLANs.

In Oracle Solaris, the elastic virtual switch (EVS) forms the back end for OpenStack networking. EVS facilitates communication between VM instances that are either on VLANs or VXLANs. For more information about OpenStack in Oracle Solaris, see [Installing and Configuring OpenStack \(Havana\) in Oracle Solaris](#).

When to Use VLANs

You can deploy VLANs if you need to do the following:

- Create a logical division of workgroups.
For example, if all systems on a floor of a building are connected on one switch-based local network, you could create a separate VLAN for each workgroup on the floor.
- Enforce differing security policies for the workgroups.
For example, the security requirements of a finance department and an information technology department are quite different. You can create a separate VLAN for each department and enforce the appropriate security policy on a per-VLAN basis.
- Reduce the size of broadcast domain and improve network efficiency. You can split workgroups into manageable broadcast domains.
For example, in a broadcast domain consisting of 25 users, if the broadcast traffic is intended only for 12 users, then setting up a separate VLAN for those 12 users can reduce traffic and improve network efficiency.

Assigning VLAN Names

VLANs demonstrate the advantage of using generic or customized names. In previous releases, the VLAN was identified by the physical point of attachment (PPA) that required combining the hardware-based name of the datalink and the VLAN ID. However, now in Oracle Solaris, you can select a more meaningful name to identify the VLAN. The name must conform to the rules for naming datalinks that are provided in [“Rules for Valid Link Names” in *Configuring and Managing Network Components in Oracle Solaris 11.3*](#). For example, you can assign a custom VLAN name such as `sales0` or `marketing1`.

VLAN names work in conjunction with VLAN IDs. Each VLAN in a local area network is identified by a VLAN ID, which is a part of the VLAN tag. The VLAN ID is assigned during VLAN configuration. When you configure switches to support VLANs, you need to assign a VLAN ID to each port. The VLAN ID on the port must be the same as the VLAN ID assigned to the interface that connects to the port.

By default, each port has a VLAN ID called Port VLAN ID. The packets that belong to this VLAN ID are not tagged with a VLAN tag. In Oracle Solaris, you can use the datalink property `default_tag` to display and change Port VLAN ID on an interface.

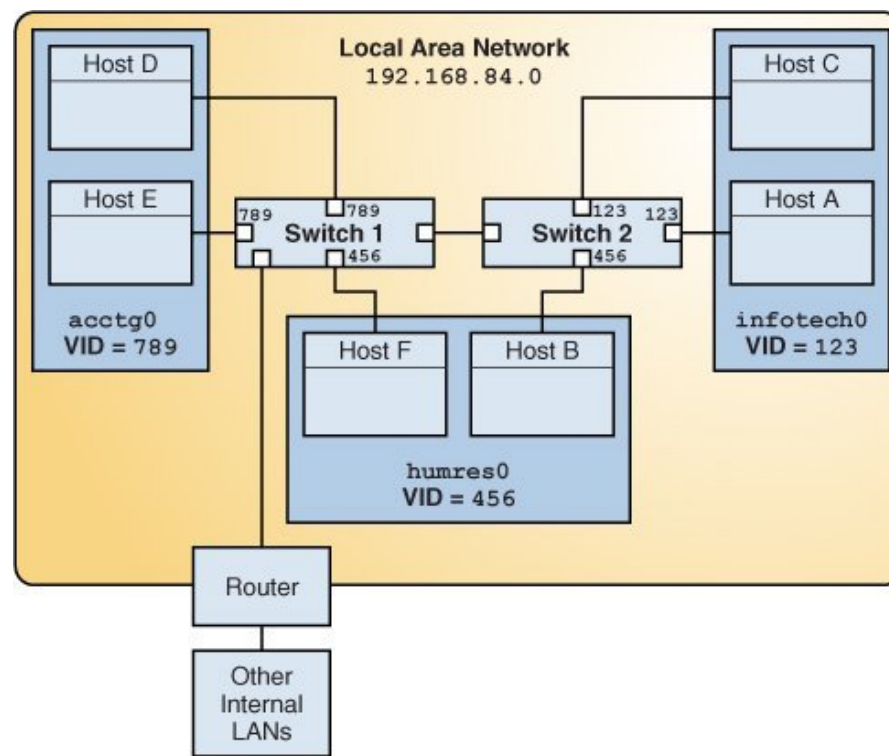
VLAN Topology

Switched LAN technology enables you to organize systems on a local network into VLANs. Before you divide a local network into VLANs, you must obtain switches that support the

VLAN technology. You can configure all ports on a switch to serve a single VLAN or multiple VLANs, depending on the VLAN topology. Each switch manufacturer has different procedures for configuring ports on a switch. Oracle Switch ES1-24 supports VLAN technology. For more information about configuration of VLAN on Oracle Switch ES1-24, see [Sun Ethernet Fabric Operating System, VLAN Administration Guide](#).

The following figure shows a local area network that has been divided into three VLANs.

FIGURE 6 Local Area Network With Three VLANs




In the illustration, the LAN has the subnet address 192.168.84.0.

This LAN is subdivided into three VLANs to correspond with three workgroups:

- acctg0 with VLAN ID 789 – Accounting group. This group owns Host D and Host E.
- humres0 with VLAN ID 456 – Human Resources group. This group owns Host B and Host F.
- infotech0 with VLAN ID 123 – Information Technology group. This group owns Host A and Host C.

A variation of previous figure is shown in the following figure, where only one switch is used and multiple hosts belonging to different VLANs connect to that single switch.

FIGURE 7 A Switch Connecting Multiple Hosts of Different VLANs



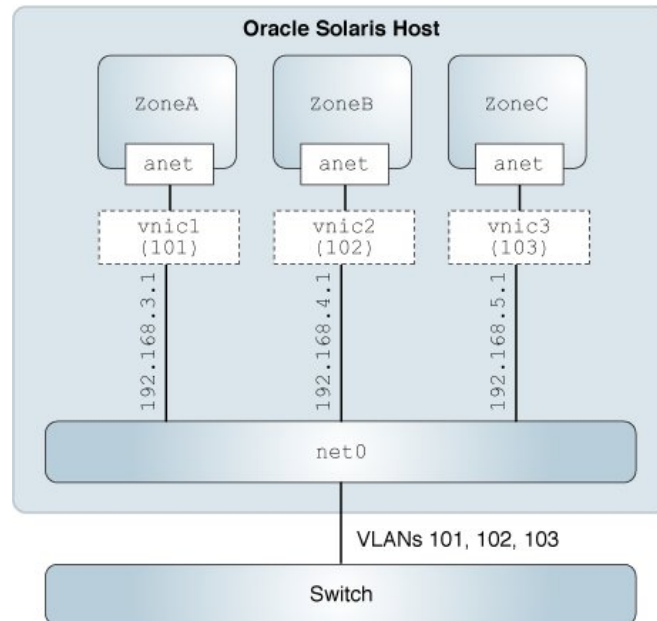
In the figure, Host A and Host C belong to the Information Technology VLAN with the VLAN ID 123. One of Host A's interface is configured with the VLAN ID 123. This interface connects to Port 1 on Switch 1, which is also configured with the VLAN ID 123. Host B is a member of the Human Resources VLAN with the VLAN ID 456. Host B's interface connects to Port 5 on Switch 1, which is configured with the VLAN ID 456. Finally, the interface of Host C is configured with the VLAN ID 123. The interface connects to Port 9 on Switch 1. Port 9 is also configured with the VLAN ID 123.

The illustration also shows that a single host can belong to multiple VLANs. For example, Host A has two VLANs configured over the interface of the host. The second VLAN is configured with the VLAN ID 456 and is connected to Port 3 that is configured with the VLAN ID 456. Therefore, Host A is a member of both the `infotech0` and the `humres0` VLANs.

Using VLANs With Virtualization

You can configure multiple virtual networks within a single network unit such as a switch by combining VLANs and Oracle Solaris Zones. The following figure shows a system with `net0` as the physical interface. ZoneA, ZoneB, and ZoneC are configured in the system with VLAN IDs, 101, 102, and 103 assigned to each of the zones respectively.

FIGURE 8 VLANs With Zones



You create a VNIC datalink by using the `dladm create-vnic` command. You create a VLAN datalink by using the `dladm create-vlan` command.

To see how to configure VLANs as shown in the figure, refer to [Example 16, “Configuring VLANs With Zones,”](#) on page 55. For more information about VLAN VNICs, see “How to Configure VNICs as VLANs” in *Managing Network Virtualization and Network Resources in Oracle Solaris 11.3* and “Modifying the VLAN IDs of VNICs” in *Managing Network Virtualization and Network Resources in Oracle Solaris 11.3*.

Planning a VLAN Configuration

Planning a VLAN configuration involves the following steps:

1. Examine the LAN topology and determine where subdivision into VLANs is appropriate.
For a basic example of such a topology, refer to [Figure 6, “Local Area Network With Three VLANs,” on page 49](#).
2. Create a numbering scheme for the VLAN IDs, and assign a VLAN ID to each VLAN.

Note - If VLAN numbering scheme already exists on the network, you must create VLAN IDs within the existing VLAN numbering scheme.

3. On each system, determine which interfaces are the components of a particular VLAN.
 - a. Determine which links are configured on the system by using the `dladm show-link` command.
 - b. Determine which VLAN ID will be associated with each datalink on the system.
 - c. Create the VLAN.
4. Check the connections of the datalinks to the switches of the network.
Note the VLAN ID of each datalink and the switch port where each interface is connected.
5. Configure each port on the switch with the same VLAN ID as the interface to which it is connected.

Refer to the switch manufacturer's documentation for configuration instructions.

Configuring a VLAN

The following procedure describes how to create a VLAN over a datalink by using the `dladm` command. You can create an IP interface over a VLAN and configure the interface with an IP address by using the `ipadm` command. For information about the `dladm` and `ipadm` commands, see the [`dladm\(1M\)`](#) and [`ipadm\(1M\)`](#) man pages.

You can also configure a VNIC with a VLAN ID to host the VLAN traffic. For more information, see [“How to Configure VNICs as VLANs” in *Managing Network Virtualization and Network Resources in Oracle Solaris 11.3*](#).

▼ How to Configure a VLAN

Before You Begin This procedure assumes that the zones are already created on the system. For information about zone configuration, refer to [Chapter 1, “How to Plan and Configure Non-Global Zones” in *Creating and Using Oracle Solaris Zones*](#).

- 1. Become an administrator.**

For more information, see [“Using Your Assigned Administrative Rights” in *Securing Users and Processes in Oracle Solaris 11.3*](#).

- 2. Determine the types of links that are in use on the system.**

```
# dladm show-link
```

- 3. Create a VLAN link over a datalink.**

```
# dladm create-vlan -l link -v vid VLAN-link
```

link Specifies the link on which the VLAN interface is being created.

vid Indicates the VLAN ID number.

VLAN-link Specifies the name of the VLAN, which can also be a meaningful custom name. For information about VLAN names, see [“Assigning VLAN Names” on page 48](#).

- 4. Verify the VLAN configuration.**

```
# dladm show-vlan
```

- 5. Create an IP interface over the VLAN.**

```
# ipadm create-ip interface
```

where *interface* provides the VLAN name.

- 6. Configure the IP interface with an IP address.**

```
# ipadm create-addr -a address interface
```

Example 14 Creating a VLAN

This example shows how to create the VLAN configuration that is illustrated in [Figure 6, “Local Area Network With Three VLANs,” on page 49](#).

1. Check the available links and create the VLANs over the specific links.

```
# dladm show-link
```

LINK	CLASS	MTU	STATE	OVER
net0	phys	1500	up	--
net1	phys	1500	up	--
net2	phys	1500	up	--

2. Host A:

```
# dladm create-vlan -l net0 -v 123 infotech0
```

Host C:

```
# dladm create-vlan -l net0 -v 123 infotech0
```

Host F:

```
# dladm create-vlan -l net0 -v 456 humres0
```

Host B:

```
# dladm create-vlan -l net0 -v 456 humres0
```

Host D:

```
# dladm create-vlan -l net0 -v 789 acctg0
```

Host E:

```
# dladm create-vlan -l net0 -v 789 acctg0
```

3. Display the VLANs created.

```
# dladm show-vlan
```

LINK	VID	SVID	PVLAN-TYPE	FLAGS	OVER
infotech0	123	--	--	----	net0
infotech0	123	--	--	----	net0
humres0	456	--	--	----	net0
humres0	456	--	--	----	net0
acctg0	789	--	--	----	net0
acctg0	789	--	--	----	net0

Example 15 Configuring VLANs With Zone's anet Resource

This example shows configuring VLANs with zone's anet resource. This example assumes that you have already configured different zones in the system. For more information about configuring zones, see [Creating and Using Oracle Solaris Zones](#).

1. Assign the VLAN IDs 11 and 12 to ZoneA and ZoneB respectively.

```
solaris#zonecfg -z ZoneA
```

```

zonecfg:ZoneA> add anet
zonecfg:ZoneA:anet> set linkname=net0
zonecfg:ZoneA:anet> set lower-link=net0
zonecfg:ZoneA:anet> set vlan-id=11
zonecfg:ZoneA:anet> set allowed-address=192.168.10.10/24
zonecfg:ZoneA:anet> set defrouter=192.168.6.10
zonecfg:ZoneA:net> end

solaris#zonecfg -z ZoneB
zonecfg:ZoneA> add anet
zonecfg:ZoneA:anet> set linkname=net1
zonecfg:ZoneA:anet> set lower-link=net1
zonecfg:ZoneA:anet> set vlan-id=12
zonecfg:ZoneA:anet> set allowed-address=192.168.11.10/24
zonecfg:ZoneA:anet> set defrouter=192.168.7.11
zonecfg:ZoneA:anet> end

```

2. Display the VLAN IDs assigned to the zones.

```

solaris:~# dladm show-vnic
LINK          OVER          SPEED  MACADDRESS      MACADDRTYPE  IDS
ZoneA/anet    net0            1000    2:8:20:47:8c:85  random       VID:11
ZoneB/anet    net1            1000    2:8:20:47:8c:85  random       VID:12

```

Note - Make sure that the VLANs are also configured on the switch.

Example 16 Configuring VLANs With Zones

This example shows how to create the VLAN configuration that is illustrated in [Figure 8, “VLANs With Zones,” on page 51](#). This example assumes that you have already configured different zones in the system. For more information about configuring zones, see [Creating and Using Oracle Solaris Zones](#).

1. Check the available links that can be used for configuring VLANs and then create the VLAN VNICs over the specific links.

```

root@solaris# dladm show-link
LINK    CLASS    MTU    STATE    OVER
net0    phys     1500    up       --
net1    phys     1500    up       --
net2    phys     1500    up       --

root@solaris# dladm create-vnic -v 101 -l net0 vnic1
root@solaris# dladm create-vnic -v 102 -l net0 vnic2
root@solaris# dladm create-vnic -v 103 -l net0 vnic3
root@solaris# dladm show-vnic

```

LINK	OVER	SPEED	MACADDRESS	MACADDRTYPE	IDS
vnic1	net0	1000	2:8:20:35:b:9a	random	VID:101
vnic2	net0	1000	2:8:20:fa:94:57	random	VID:102
vnic3	net0	1000	2:8:20:51:1c:4a	random	VID:103

When link information is displayed, the VLANs are included in the list.

```
root@solaris# dladm show-link
LINK      CLASS  MTU    STATE  OVER
net0      phys   1500   up     -
vnic1     vnic   1500   up     net0
vnic2     vnic   1500   up     net0
vnic3     vnic   1500   up     net0
```

- Specify the VLAN IDs 101, 102, and 103 to ZoneA, ZoneB, and ZoneC respectively.

```
root@solaris# zonecfg -z ZoneA
zonecfg:ZoneA> add net
zonecfg:ZoneA:net> set physical=vnic1
zonecfg:ZoneA:net> end
zonecfg:ZoneA> verify
zonecfg:ZoneA> commit
zonecfg:ZoneA> exit
root@solaris# zoneadm -z ZoneA reboot
```

```
root@solaris# zonecfg -z ZoneB
zonecfg:ZoneB> add net
zonecfg:ZoneB:net> set physical=vnic2
zonecfg:ZoneB:net> end
zonecfg:ZoneB> verify
zonecfg:ZoneB> commit
zonecfg:ZoneB> exit
root@solaris# zoneadm -z ZoneB reboot
```

```
root@solaris# zonecfg -z ZoneC
zonecfg:ZoneC> add net
zonecfg:ZoneC:net> set physical=vnic3
zonecfg:ZoneC:net> end
zonecfg:ZoneC> verify
zonecfg:ZoneC> commit
zonecfg:ZoneC> exit
root@solaris# zoneadm -z ZoneC reboot
```

- Display the VLAN VNICs assigned to the zones.

```
root@solaris:~# dladm show-vnic
LINK      OVER      SPEED  MACADDRESS      MACADDRTYPE  IDS
```


ZoneA/vnic1	net0	1000	2:8:20:47:8c:85	random	VID:101
ZoneB/vnic2	net0	1000	2:8:20:47:8c:85	random	VID:102
ZoneC/vnic3	net0	1000	2:8:20:47:8c:85	random	VID:103

4. Log in to each non-global zone to configure the VLAN with an IP address.

```

root@solaris:~# zlogin ZoneA
root@ZoneA:~# ipadm create-ip vnic1
root@ZoneA:~# ipadm create-addr -a 192.168.3.1 vnic1
vnic1/v4

root@solaris:~# zlogin ZoneB
root@ZoneB:~# ipadm create-ip vnic2
root@ZoneB:~# ipadm create-addr -a 192.168.4.1 vnic2
vnic2/v4

root@solaris:~# zlogin ZoneC
root@ZoneC:~# ipadm create-ip vnic3
root@ZoneC:~# ipadm create-addr -a 192.168.5.1 vnic3
vnic3/v4

```

Note - Make sure that the VLANs are also configured on the switch.

After all the VLANs have been configured with IP addresses, the configuration is complete. The three VLANs are operative and can host traffic for their respective zones.

Configuring VLANs Over a Link Aggregation

You can create VLANs on a link aggregation in a manner that is similar to configuring VLANs over an interface. Link aggregations are described in [Chapter 2, “Configuring High Availability by Using Link Aggregations”](#). This section combines configuring VLANs and link aggregations.

▼ How to Configure VLANs Over a Link Aggregation

Before You Begin Create the link aggregation. For information about how to create link aggregations, refer to [“How to Create a Link Aggregation”](#) on page 28.

1. **Become an administrator.**

For more information, see [“Using Your Assigned Administrative Rights”](#) in *Securing Users and Processes in Oracle Solaris 11.3*.

2. List the link aggregations that are configured on the system.

```
# dladm show-aggr
```

3. For every VLAN that you want to create over the link aggregation you selected, you can use the following command:

```
# dladm create-vlan -l link -v vid VLAN-link
```

link Specifies the link on which the VLAN interface is being created.

Note - In this procedure, the link refers to the link aggregation.

vid Indicates the VLAN ID number.

VLAN-link Specifies the name of the VLAN.

4. For every VLAN that you created in the previous step, create an IP interface over the VLAN.

```
# ipadm create-ip interface
```

where *interface* uses the VLAN name.

5. For each IP interface on a VLAN, configure a valid IP address.

```
# ipadm create-addr -a address interface
```

Example 17 Configuring Multiple VLANs Over a Link Aggregation

In this example, two VLANs are configured on a link aggregation. The VLANs are assigned VLAN IDs 193 and 194, respectively.

```
# dladm show-link
LINK    CLASS    MTU    STATE    OVER
net0    phys     1500   up       -
net1    phys     1500   up       -
aggr0    aggr     1500   up       net0 net1

# dladm create-vlan -l aggr0 -v 193 acctg0
# dladm create-vlan -l aggr0 -v 194 humres0

# ipadm create-ip acctg0
# ipadm create-ip humres0

# ipadm create-addr -a 192.168.10.0/24 acctg0
ipadm: acctg0/v4
# ipadm create-addr -a 192.168.20.0/24 humres0
```

```
ipadm: humres0/v4
```

Configuring VLANs on a Legacy Device

Certain legacy devices handle only packets whose maximum transmission unit (MTU) size, also known as frame size, is a maximum of 1514 bytes. Packets whose frame sizes exceed this maximum limit are dropped. For such cases, follow the same procedure listed in [“How to Configure a VLAN” on page 53](#). However, when creating the VLAN, use the `-f` option to force the creation of the VLAN.

▼ How to Configure VLANs on a Legacy Device

1. **Become an administrator.**

For more information, see [“Using Your Assigned Administrative Rights” in *Securing Users and Processes in Oracle Solaris 11.3*](#).

2. **Create the VLAN with the `-f` option.**

```
# dladm create-vlan -f -l link -v vid VLAN-link
```

`-f` Forces the creation of VLAN. Use this option when you are creating a VLAN on devices that do not allow frame sizes large enough to include a VLAN header.

`-l link` Specifies the link on which the VLAN interface is created. In this procedure, the link refers to the legacy device.

`-v vid` Indicates the VLAN ID number.

`VLAN-link` Specifies the name of the VLAN, which can also be an administratively-chosen name.

3. **Set a lower size for the maximum transmission unit (MTU).**

In the following example, mtu is set to 1496.

```
# dladm set-linkprop -p mtu=1496 VLAN-link
```

The lower MTU value gives space for the link layer to insert the VLAN header prior to transmission.

4. **Repeat Step 2 to set a MTU value for each node in the VLAN.**

For more information about changing link property values, refer to [“Administering Datalink Properties” in *Configuring and Managing Network Components in Oracle Solaris 11.3*](#).

Displaying VLAN Information

You can use the `dladm show-link` command to display information about VLANs because VLANs are datalinks. Use the `dladm show-vlan` command to display specific information about the VLANs.

The following example compares the type of information that you can obtain by using either the `dladm show-link` or `dladm show-vlan` command. The first example uses the `dladm show-link` command to display all of the datalinks on the system, including those that are not VLANs. The second example uses the `dladm show-vlan` command to display a subset of the datalink information that is relevant only to VLANs.

dladm show-link

LINK	CLASS	MTU	STATE	OVER
net0	phys	1500	up	--
net1	phys	1500	up	--
net2	phys	1500	up	--
web1	vlan	1500	up	net0
auth1	vlan	1500	up	net0
app1	vlan	1500	up	net0
web2	vlan	1500	up	net1
auth2	vlan	1500	up	net1
app2	vlan	1500	up	net1
web3	vlan	1500	up	net2
auth3	vlan	1500	up	net2

dladm show-vlan

LINK	VID	SVID	PVLAN-TYPE	FLAGS	OVER
web1	111	--	--	----	net0
auth1	112	--	--	----	net0
app1	113	--	--	----	net0
web2	111	--	--	----	net1
auth2	112	--	--	----	net1
app2	113	--	--	----	net1
web3	111	--	--	----	net2
auth3	113	--	--	----	net2

Modifying VLANs

You can modify a VLAN by using the `dladm modify-vlan` command in the following ways:

- Change the VLAN ID of a VLAN
- Migrate a VLAN to another underlying link

Modifying the VLAN ID of a VLAN

To change the VLAN ID of a VLAN, use one of the following commands:

- `dladm modify-vlan -v vid -L datalink`

where *vid* specifies the new VLAN ID that you are assigning to the VLAN and *datalink* refers to the underlying link over which the VLAN is configured.

Note - You can use the `dladm modify-vlan -v vid -L datalink` command syntax only if a single VLAN exists on the datalink. The command fails if you use it on a datalink that has multiple configured VLANs because each VLAN on a datalink must have unique VLAN IDs.

If you modify the VLAN ID on the link, you must also configure the switch port for the new VLAN ID.

- `dladm modify-vlan -v vid vlan`

Use this command to change the unique VLAN IDs of multiple VLANs over a single datalink. Each VLAN on the datalink has a unique VLAN ID so you must change the VLAN IDs one at a time. In the setup shown in [Figure 8, “VLANs With Zones,” on page 51](#), you would change the VLAN IDs of `web1`, `auth1`, and `app1` that are configured over `net0` as follows:

```
# dladm modify-vlan -v 123 web1
# dladm modify-vlan -v 456 app1
# dladm modify-vlan -v 789 auth1
```

Migrating a VLAN to Another Underlying Link

You can migrate a VLAN from one underlying datalink to another underlying datalink without deleting and reconfiguring the VLAN. The underlying link can be a physical link, a link aggregation, or an etherstub. For more information about etherstubs, see [“Network](#)

[Virtualization Components](#)” in *Managing Network Virtualization and Network Resources in Oracle Solaris 11.3*.

To successfully migrate a VLAN, the underlying datalink to which the VLAN is moved must be able to accommodate the datalink properties of the VLAN. If those properties are not supported, then migration fails and the user is notified. After a successful migration, all the applications that use the VLAN continue to operate normally provided that the VLAN remains connected to the network.

Certain hardware-dependent properties might change after a VLAN migration. For example, a VLAN always shares the same MAC address as its underlying datalink. Therefore, when you migrate a VLAN, the VLAN's MAC address changes to the primary MAC address of the target datalink. Other properties that might be affected are the datalink state, link speed, and MTU size. However, applications continue to operate without interruption.

Note - A migrated VLAN does not retain any of its hardware lane statistics from the original datalink. Available hardware lanes for the VLAN on the target datalink become the new source of statistics information. However, software statistics that are displayed by default by the `dlstat` command are preserved.

You can perform a VLAN migration either globally or selectively.

Global Migration

Global migration is used to migrate all of the VLANs that are configured over one datalink to another datalink. To perform a global migration, you need to specify only the source and target datalinks. The following example shows how to move all of the VLANs on `ether0` to `net1`.

```
# dladm modify-vlan -l net1 -L ether0
```

-l Refers to the target datalink to which the VLANs are migrated.

-L Refers to the original datalink over which the VLANs are configured.

Note - You must specify the target datalink before the source datalink.

Selective Migration

Selective migration is used to migrate only selected VLANs. To perform selective VLAN migration, you specify the VLANs that you want to move. In the following example, which is based on [Figure 8, “VLANs With Zones,” on page 51](#), VLANs are moved from `net0` to `net3`.

```
# dladm modify-vlan -l net3 web1,auth1,app1
```

Note - When migrating VLANs selectively, do not include the -L option, which applies only to global migration.

You can change the VLAN IDs of VLANs while performing a migration. Using [Figure 8, “VLANs With Zones,” on page 51](#) as the basis, the following example shows how you would migrate multiple VLANs and change their VLAN IDs at the same time.

```
# dladm show-vlan
LINK   VID   SVID   PVLAN-TYPE   FLAGS   OVER
web1    111   --     --           ----- net0
auth1   112   --     --           ----- net0
app1    113   --     --           ----- net0

# dladm modify-vlan -l net3 -v 123 web1
# dladm modify-vlan -l net3 -v 456 auth1
# dladm modify-vlan -l net3 -v 789 app1
# dladm show-vlan
LINK   VID   SVID   PVLAN-TYPE   FLAGS   OVER
web1    123   --     --           ----- net3
auth1   456   --     --           ----- net3
app1    789   --     --           ----- net3
```

Note - A parallel command, `dladm modify-vnic`, migrates VNICs that are configured as VLANs. You must use the correct subcommand depending on whether you are migrating VLANs or VNICs that are configured as VLANs. Use the `modify-vlan` subcommand on VLANs that are displayed by the `dladm show-vlan` command. Use the `modify-vnic` subcommand on VNICs, including those with VLAN IDs, that are displayed in the output of the `dladm show-vnic` command. For information about how to modify VNICs, see [“Modifying the VLAN IDs of VNICs” in *Managing Network Virtualization and Network Resources in Oracle Solaris 11.3*](#).

Deleting a VLAN

Use the `dladm delete-vlan` command to delete VLAN configurations on your system.

Note - You must first delete any existing IP configurations on the VLAN that you intend to delete before you can delete the VLAN. Deleting a VLAN fails if IP interfaces exist over the VLAN.

Note - When you remove a datalink by using the `dladm delete-phys` command, any VLAN that is configured over the datalink is also deleted. This `dladm delete-phys` command also deletes the other layer 2 components such as flows, VNICs, and aggregations configured over the datalink. However, you must manually delete the layer 3 components configured over the datalink such as IP interfaces by using the `ipadm` command.

EXAMPLE 18 Deleting a VLAN Configuration

This example shows how to delete a VLAN configuration.

```
# dladm show-vlan
LINK      VID    SVID    PVLAN-TYPE  FLAGS    OVER
web1      111    --      --          ----    net0
auth1     112    --      --          ----    net0
app1      113    --      --          ----    net0
web2      111    --      --          ----    net1
auth2     112    --      --          ----    net1
app2      113    --      --          ----    net1
web3      111    --      --          ----    net2
auth3     113    --      --          ----    net2

# ipadm delete-ip web1
# dladm delete-vlan web1
```

Using VLAN Tagging for SR-IOV-Enabled Ports

When using some Intel network adapters that support SR-IOV, a virtual function might be the target of malicious behavior. Unexpected software-generated frames can throttle traffic between the host and the virtual switch, which might negatively affect performance. To drop unexpected and potentially malicious frames, configure all SR-IOV-enabled ports to use VLAN tagging. For more information, see [Oracle VM Server for SPARC 3.3 Administration Guide](#).

Use Case: Combining Link Aggregations and VLAN Configurations

This section provides an example that shows how to create a combination of network configurations that uses link aggregations and VLANs.

In the following example, a system that uses four NICs must be configured to be a router for eight separate subnets. Therefore, eight links are configured, one for each subnet. First, a trunk aggregation is created on all four NICs. This untagged link that does not include a VLAN tag

in the outgoing frame becomes the default untagged subnet for the network to which the default route points.

VLAN interfaces are then configured over the link aggregation for the other subnets. The subnets are named based on a color-coded scheme. Accordingly, the VLAN names are likewise named to correspond to their respective subnets. The final configuration consists of eight links for the eight subnets: one untagged link and seven tagged VLAN links. The example begins with verifying whether IP interfaces already exist on the datalinks. These interfaces must be deleted before the datalinks can be combined into an aggregation.

1. Remove any IP interfaces that have been configured over the datalinks.

```
# ipadm show-if
IFNAME    CLASS    STATE    ACTIVE    OVER
lo0        loopback ok        yes       --
net0       ip        ok        yes       --
net1       ip        ok        yes       --
net2       ip        ok        yes       --
net3       ip        ok        yes       --
```

```
# ipadm delete-ip net0
# ipadm delete-ip net1
# ipadm delete-ip net2
# ipadm delete-ip net3
```

2. Create the trunk aggregation default0.

```
# dladm create-aggr -P L2,L3 -l net0 -l net1 -l net2 -l net3 default0
```

```
# dladm show-link
LINK      CLASS    MTU    STATE    OVER
net0      phys     1500   up       --
net1      phys     1500   up       --
net2      phys     1500   up       --
net3      phys     1500   up       --
default0  aggr     1500   up       net0 net1 net2 net3
```

3. Configure an IP interface over the aggregation.

```
# ipadm create-ip default0
# ipadm create-addr -a 10.2.3.4/24 default0
```

4. Create the VLANs over default0.

```
# dladm create-vlan -v 2 -l default0 orange0
# dladm create-vlan -v 3 -l default0 green0
# dladm create-vlan -v 4 -l default0 blue0
# dladm create-vlan -v 5 -l default0 white0
# dladm create-vlan -v 6 -l default0 yellow0
# dladm create-vlan -v 7 -l default0 red0
```

```
# dladm create-vlan -v 8 -l default0 cyan0
```

```
# dladm show-link
```

LINK	CLASS	MTU	STATE	OVER
net0	phys	1500	up	--
net1	phys	1500	up	--
net2	phys	1500	up	--
net3	phys	1500	up	--
default0	aggr	1500	up	net0 net1 net2 net3
orange0	vlan	1500	up	default0
green0	vlan	1500	up	default0
blue0	vlan	1500	up	default0
white0	vlan	1500	up	default0
yellow0	vlan	1500	up	default0
red0	vlan	1500	up	default0
cyan0	vlan	1500	up	default0

```
# dladm show-vlan
```

LINK	VID	SVID	PVLAN-TYPE	FLAGS	OVER
orange0	2	--	--	-----	default0
green0	3	--	--	-----	default0
blue0	4	--	--	-----	default0
white0	5	--	--	-----	default0
yellow0	6	--	--	-----	default0
red0	7	--	--	-----	default0
cyan0	8	--	--	-----	default0

5. Create IP interfaces over the VLAN links and assign IP addresses to the interfaces.

```
# ipadm create-ip orange0
```

```
# ipadm create-ip green0
```

```
# ipadm create-ip blue0
```

```
# ipadm create-ip white0
```

```
# ipadm create-ip yellow0
```

```
# ipadm create-ip red0
```

```
# ipadm create-ip cyan0
```

```
# ipadm create-addr -a 10.2.3.5/24 orange0
```

```
# ipadm create-addr -a 10.2.3.6/24 green0
```

```
# ipadm create-addr -a 10.2.3.7/24 blue0
```

```
# ipadm create-addr -a 10.2.3.8/24 white0
```

```
# ipadm create-addr -a 10.2.3.9/24 yellow0
```

```
# ipadm create-addr -a 10.2.3.10/24 red0
```

```
# ipadm create-addr -a 10.2.3.11/24 cyan0
```

Use Case: Configuring Flows for a VLAN

This section provides an example to show how to configure flows for a VLAN that is assigned to a zone. The zone has multiple services such as production, replication, and backup running over the same VLAN link. This example shows provisioning of the link's bandwidth among these services by using flows.

1. Check the available links and create the VLAN VNIC over a specific link.

```
root@solaris:~# dladm show-link
LINK      CLASS      MTU      STATE      OVER
net0      phys       1500     up         --
net1      phys       1500     up         --
net2      phys       1500     up         --
```

```
root@solaris:~# dladm create-vnic -v 101 -l net0 vnic1
```

2. Create a zone named zone1 and assign the VLAN VNIC to it.

```
root@solaris# zonecfg -z zone1
Use 'create' to begin configuring a new zone.
zonecfg:zone1> create -t SYSsolaris
zonecfg:zone1> add net
zonecfg:zone1:net> set physical=vnic1
zonecfg:zone1:net> end
zonecfg:zone1> verify
zonecfg:zone1> commit
zonecfg:zone1> exit
root@solaris# zoneadm -z zone1 boot
```

3. Log in to the zone and assign IP address for the VLAN VNIC.

```
root@solaris# zlogin zone1
root@zone1# ipadm create-ip vnic1
root@zone1:~# ipadm create-addr -a 192.168.110.1 vnic1
```

4. Configure flows for different services on the zone.

```
root@zone1:~# flowadm add-flow -a local_ip=192.168.110.1,local_port=80 \
-p maxbw=500M,priority=high -l vnic1 flow_prod
root@zone1:~# flowadm add-flow -a local_ip=192.168.110.1,local_port=138 \
-p maxbw=250M,priority=medium -l vnic1 flow_repl
root@zone1:~# flowadm add-flow -a local_ip=192.168.110.1,local_port=21 \
-p maxbw=250M,priority=low -l vnic1 flow_bkup
```

Since multiple services are running over the same VLAN, provisioning the bandwidth by using flows helps to manage the network resources efficiently.

5. Display the flows that are configured.

```
root@zone1:~# flowadm
```

FLOW	LINK	PROTO	LADDR	LPORT	RADDR	RPORT	DSFLD
flow_prod	vnic1	--	192.168.110.1	80	--	--	--
flow_repl	vnic1	--	192.168.110.1	138	--	--	--
flow_bkup	vnic1	--	192.168.110.1	21	--	--	--

6. Display the properties that are set for the flows.

```
root@zone1:~# flowadm show-flowprop
```

FLOW	PROPERTY	PERM	VALUE	DEFAULT	POSSIBLE
flow_prod	maxbw	rw	500	--	--
flow_prod	priority	rw	high	medium	low,medium,high
flow_prod	hwflow	r-	off	--	on,off
flow_repl	maxbw	rw	250	--	--
flow_repl	priority	rw	medium	medium	low,medium,high
flow_repl	hwflow	r-	off	--	on,off
flow_bkup	maxbw	rw	250	--	--
flow_bkup	priority	rw	low	medium	low,medium,high
flow_bkup	hwflow	r-	off	--	on,off

Use Case: Configuring Bandwidth for VLAN VNICs

This section shows an example to show how to configure bandwidth limit on the VLAN VNICs to provision the bandwidth of the physical link among the VLAN VNICs. In this example, there are three services running in different zones, which share the same physical link.

1. Check the available links and create the VNIC VLANs over a specific link.

```
root@solaris# dladm show-link
```

LINK	CLASS	MTU	STATE	OVER
net0	phys	1500	up	--

```
root@solaris# dladm create-vnic -v 10 -l net0 vnic_prod
root@solaris# dladm create-vnic -v 11 -l net0 vnic_repl
root@solaris# dladm create-vnic -v 12 -l net0 vnic_bkup
root@solaris# dladm show-vnic
```

LINK	OVER	SPEED	MACADDRESS	MACADDRTYPE	IDS
vnic_prod	net0	1000	2:8:20:35:b:9a	random	VID:10
vnic_repl	net0	1000	2:8:20:fa:94:57	random	VID:11
vnic_bkup	net0	1000	2:8:20:51:1c:4a	random	VID:12

2. Assign bandwidth to all the VNICs.

```
root@solaris# dladm set-linkprop -p maxbw=1G vnic_prod
root@solaris# dladm set-linkprop -p maxbw=500M vnic_repl
root@solaris# dladm set-linkprop -p maxbw=200M vnic_bkup
```

3. Create a zone named zone1 and assign a VLAN VNIC to it.

```
root@solaris# zonecfg -z zone1
```

Use 'create' to begin configuring a new zone.

```
zonecfg:zone1> create -t SYSsolaris
zonecfg:zone1> add net
zonecfg:zone1:net> set physical=vnic_prod
zonecfg:zone1:net> end
zonecfg:zone1> verify
zonecfg:zone1> commit
zonecfg:zone1> exit
root@solaris# zoneadm -z zone1 boot
```

4. Log in to zone1 and configure the VNIC with an IP address.

```
root@solaris# zlogin zone1
root@zone1# ipadm create-ip vnic_prod
root@zone1:~# ipadm create-addr -a 192.168.110.1 vnic_prod
vnic_prod/v4
root@solaris# zoneadm -z zone1 reboot
```

5. Create a zone named zone2 and assign a VLAN VNIC to it.

```
root@solaris# zonecfg -z zone2
Use 'create' to begin configuring a new zone.
zonecfg:zone2> create -t SYSsolaris
zonecfg:zone2> add net
zonecfg:zone2:net> set physical=vnic_repl
zonecfg:zone2:net> end
zonecfg:zone2> verify
zonecfg:zone2> commit
zonecfg:zone2> exit
root@solaris# zoneadm -z zone2 boot
```

6. Log in to zone2 and configure the VNIC with an IP address.

```
root@solaris# zlogin zone2
root@zone2# ipadm create-ip vnic_repl
root@zone2:~# ipadm create-addr -a 192.168.111.1 vnic_repl
vnic_repl/v4
```

7. Create a zone named zone3 and assign a VLAN VNIC to it.

```
root@solaris# zonecfg -z zone3
Use 'create' to begin configuring a new zone.
zonecfg:zone3> create -t SYSsolaris
zonecfg:zone3> add net
zonecfg:zone3:net> set physical=vnic_bkup
zonecfg:zone3:net> end
zonecfg:zone3> verify
zonecfg:zone3> commit
zonecfg:zone3> exit
root@solaris# zoneadm -z zone3 boot
```

8. Log in to zone3 and configure the VNIC with an IP address.

```
root@solaris# zlogin zone3
root@zone3# ipadm create-ip vnic_bkup
root@zone3:~# ipadm create-addr -a 192.168.112.1 vnic_bkup
vnic_bkup/v4
```

In the example, one VLAN VNIC is assigned for each service in the zones. The bandwidth of the VNIC is configured to provision the bandwidth of the physical link. This ensures that the available network resources are efficiently used.

Configuring Private Virtual Local Area Networks

This chapter discusses private VLANs (PVLANS) and describes the procedures to configure and modify PVLANS.

This chapter contains the following topics:

- “Overview of Private VLANs” on page 71
- “Configuring a Private VLAN” on page 75
- “Modifying Private VLANs” on page 76
- “Deleting a VLAN” on page 63
- “Assigning a PVLAN to a Zone” on page 77

Overview of Private VLANs

Private VLAN (PVLAN) technology enables you to divide a VLAN into sub-VLANs to isolate network traffic. A regular VLAN is a single broadcast domain. When you configure a PVLAN, the single broadcast domain is partitioned into smaller subdomains. The standard (IEEE 802.1 Q) VLAN is called the *primary VLAN* and the sub-VLANs are called the *secondary VLANs* when you configure a PVLAN. The secondary VLANs can be either isolated VLANs or community VLANs.

- **Isolated VLAN**

The ports that are associated with an isolated VLAN can communicate only with the primary VLAN and not with any other secondary VLAN. You can create only one isolated VLAN within a primary VLAN domain.

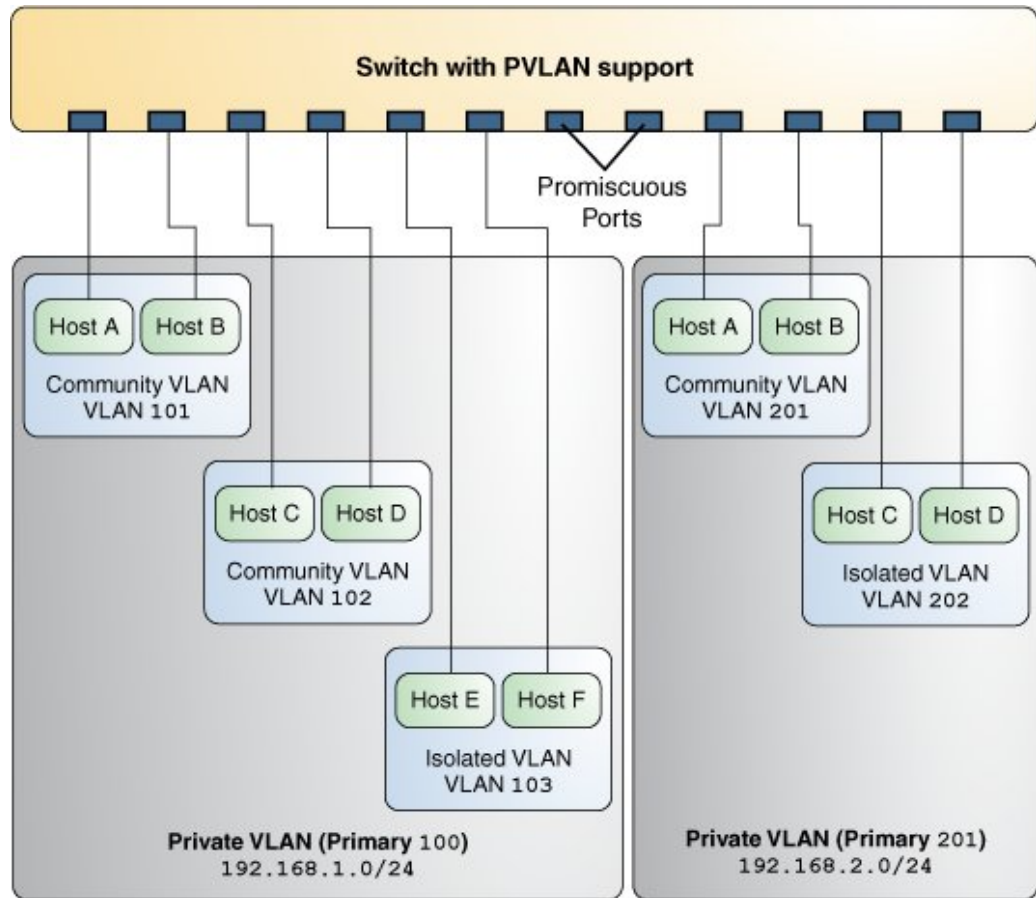
- **Community VLAN**

The ports associated with community VLAN can communicate with the primary VLAN and the other ports that are in the same community VLAN. You can create multiple community VLANs within a primary VLAN domain.

For information about the regular VLANs, see [Chapter 3, “Configuring Virtual Networks by Using Virtual Local Area Networks”](#).

The following figure shows a simple PVLAN configuration with two PVLANs configured on a switch that supports PVLAN.

FIGURE 9 Private VLAN



In the figure, the private VLAN with the primary VLAN ID 100 has three secondary VLANs with secondary VLAN IDs 101, 102, and 103. Two of the secondary VLANs are community VLANs with the hosts Host A and Host B and Host C and Host D, and the other is an isolated VLAN with the hosts Host E and Host F.

The ports associated with the community VLANs 101 and 102 can communicate with the primary VLAN 100 and also can communicate with each other, that is, Host A can communicate with Host B and Host C can communicate with Host D. However, Host A and Host B in community VLAN 101 cannot communicate with Host C and Host D in community

VLAN 102. The ports associated with the isolated VLAN can communicate only with the primary VLAN and not with each other.

For more information, see [“Configuring a Private VLAN” on page 75](#).

Advantages of Using PVLANS

Using PVLAN provides the following advantages:

- Conserves IP addresses because multiple VLANs can be created per IP subnet.
- Increases scalability because PVLANS provide a greater number of L2 isolation networks than regular VLANs. This increase helps to overcome the limitation of VLANs, which enables you to create a maximum of 4094 isolated networks.
- Provides better isolation compared to regular VLANs.

PVLAN Ports

PVLANS can span multiple switches similar to regular VLANs. A trunk port carries frames either from a primary VLAN or a secondary VLAN. The two different types of ports associated with the PVLAN are promiscuous trunk ports and secondary trunk ports.

- **Promiscuous trunk port** – All the frames going out from the promiscuous trunk port are tagged with the primary VLAN ID. This port is configured on the top-level switch uplink port. The switch maps the primary VLAN ID and the secondary VLAN ID.
- **PVLAN secondary trunk port** – All the frames going out from PVLAN secondary trunk port are tagged with the secondary VLAN ID.

Note - Even though the PVLAN ports are isolated at layer 2, the ports can still communicate with each other at layer 3 as long as the external policy gateway allows the communication.

Tagging the Outgoing Traffic

In Oracle Solaris, you must set the tag mode property for a PVLAN depending on whether the promiscuous trunk port is on the system or the switch. Therefore, you need to identify whether the PVLAN is configured on the switch and the secondary ports are configured on Oracle Solaris or the PVLAN is configured only on the system and not on the switch.

To send out the packets with the primary VLAN ID, you must set the tag mode to `primary`. To send out packets with the secondary VLAN ID, you must set the tag mode to `secondary`. The switch converts the secondary VLAN ID to the primary VLAN ID. By default, the packets are sent out with the primary VLAN ID.

You can tag the outgoing traffic with a primary VLAN ID or a secondary VLAN ID by configuring the `tagmode` property. The default value of this property is `primary`. If there are multiple PVLAN switches and you want a PVLAN to span these multiple switches, you must set the property `tagmode` to `secondary` by using the `dladm set-linkprop` command.

EXAMPLE 19 Setting the Tag Mode for a PVLAN

The following example shows how you can change the tag mode on a datalink.

```
# dladm set-linkprop -p pvlan-tagmode=secondary net0
# dladm show-linkprop -p pvlan-tagmode net0
```

LINK	PROPERTY	PERM	VALUE	EFFECTIVE	DEFAULT	POSSIBLE
net0	pvlan-tagmode	rw	secondary	secondary	primary	secondary, primary

Note - The value of `tagmode` depends on whether the switch supports PVLANS.

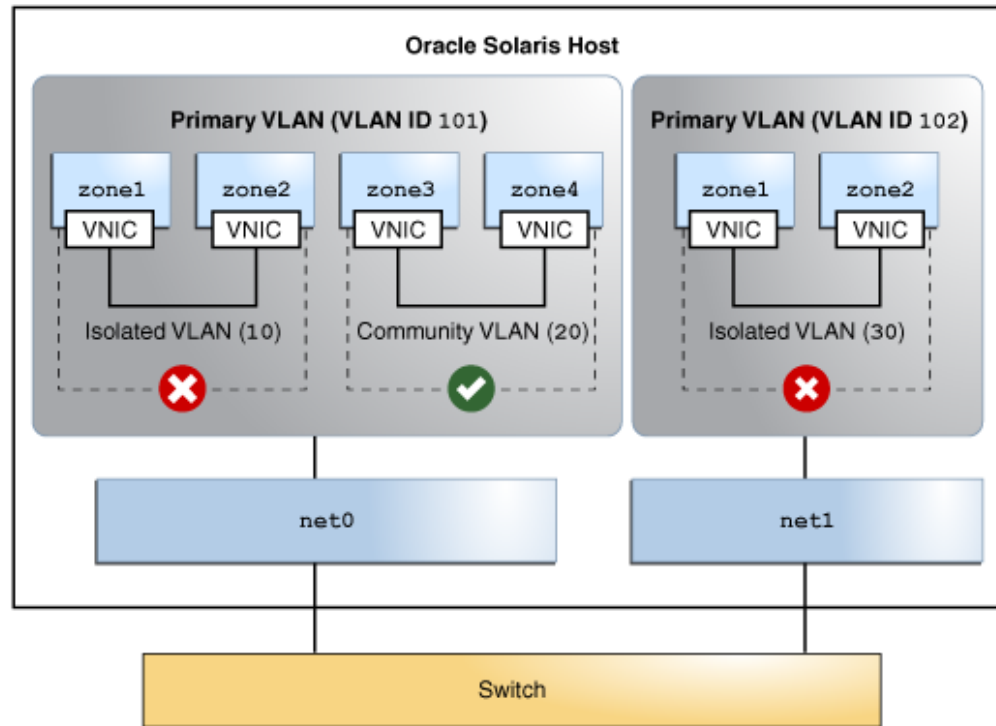
PVLAN Configuration Requirements

When configuring a PVLAN, note the following restrictions:

- The primary VLAN ID and secondary VLAN ID for community VLANs must be unique.
For example, if you have a community VLAN with a primary VLAN ID of 3 and a secondary VLAN ID of 100, you cannot create another community VLAN that uses either 3 or 100 as the secondary VLAN ID. That is, any combination containing the ID pair such as 4, 100 or 4, 3 is invalid.
- You can reuse the secondary VLAN ID of an isolated VLAN.
For example, if you have an isolated VLAN with a primary VLAN ID of 3 and a secondary VLAN ID of 100, you can reuse the VLAN ID 100 as the primary VLAN ID or the community secondary VLAN ID.

PVLANS With Zones

You can configure multiple private virtual networks within a single network unit such as a switch by combining VLANs and Oracle Solaris zones. With PVLAN you can provide network isolation between zones which are in the same VLAN without introducing any additional VLANs. The following figure shows a system with two physical NICs and two PVLANS are configured over them.

FIGURE 10 PVLAN With Zones

The figure shows two PVLANS with VLAN IDs 101 and 102. An isolated and a community VLAN are configured on Primary VLAN 101. Only one isolated VLAN is configured on primary VLAN 102. The zones in the isolated VLAN and the community VLAN on the Primary VLAN 101 cannot communicate with each other. However, the hosts within the community VLAN can communicate with each other.

For information about how to assign PVLAN to a zone, see [“Assigning a PVLAN to a Zone” on page 77](#).

Configuring a Private VLAN

You can either create a PVLAN by using the `dladm create-vlan` or configure a PVLAN VNIC to host the PVLAN traffic by using the `dladm create-vnic` command. For information

about configuring a PVLAN VNIC, see [“How to Configure VNICs as PVLANs” in *Managing Network Virtualization and Network Resources in Oracle Solaris 11.3*](#).

To create a PVLAN, use the following command:

```
dladm create-vlan [ -l link -v VLAN-ID[,PVLAN-SVID[,PVLAN-type]] ] [VLAN-link]
```

<i>link</i>	Specifies the Ethernet link over which the VLAN is created.
<i>VLAN-ID</i>	Primary ID associated with a VLAN.
<i>PVLAN-SVID</i>	PVLAN Secondary VLAN ID associated with the VLAN.
<i>PVLAN-type</i>	PVLAN type associated with the VLAN, which can be either <i>isolated</i> or <i>community</i> . The default value is <i>isolated</i> .
<i>VLAN-link</i>	Name of the VLAN.

EXAMPLE 20 Creating an Isolated PVLAN by Using the `dladm create-vlan` Command

The following example shows how you can create a PVLAN with a primary VLAN ID of 3, a secondary VLAN ID of 100, and the PVLAN type set to *isolated*.

```
# dladm create-vlan -v 3,100,isolated -l net0 vlan1
# dladm show-vlan
LINK    VID    SVID    PVLAN-TYPE  FLAGS    OVER
vlan1   3      100     isolated    -----  net0
```

Modifying Private VLANs

You can modify the primary VLAN ID, secondary VLAN ID, and the PVLAN type of PVLANs by using the `dladm modify-vlan` command to modify a PVLAN. You can use this command only if you created the PVLAN by using the `dladm create-vlan` command and not by using the `dladm create-vnic` command. For information about modifying PVLAN VNICs, see [“Modifying PVLAN VNICs” in *Managing Network Virtualization and Network Resources in Oracle Solaris 11.3*](#).

To modify the PVLAN, use the following command:

```
dladm modify-vlan [-v VID[,PVLAN-SVID[,PVLAN-type]] ] VLAN-link
```

You modify the primary VLAN ID, the secondary VLAN ID, and the PVLAN type tuple together.

The following example shows how you can set the primary VLAN ID to 15, the secondary VLAN ID to 103, and the PVLAN type to *community*.

```
# dladm show-vlan
LINK      VID    SVID    PVLAN-TYPE  FLAGS    OVER
vlan1     10     102     isolated    ----    net0
# dladm modify-vlan -v 15,103,community vlan1
# dladm show-vlan
LINK      VID    SVID    PVLAN-TYPE  FLAGS    OVER
vlan1     15     103     community    ----    net0
```

Deleting a Private VLAN

Use the `dladm delete-vlan` command to delete PVLAN configurations on your system.

Note - You must first delete any existing IP configurations on the PVLAN before you can delete the PVLAN. Deleting a PVLAN fails if IP interfaces exist over the PVLAN.

EXAMPLE 21 Deleting a VLAN Configuration

This example shows how to delete a VLAN configuration.

```
# dladm show-vlan
LINK      VID    SVID    PVLAN-TYPE  FLAGS    OVER
vlan1     15     103     community    ----    net0
vlan2     20     105     isolated    ----    net1
# ipadm delete-ip vlan1
# dladm delete-vlan vlan1
# dladm show-vlan
LINK      VID    SVID    PVLAN-TYPE  FLAGS    OVER
vlan2     20     105     isolated    ----    net1
```

Assigning a PVLAN to a Zone

You can assign a PVLAN to a zone for better isolation. For information about configuring a zone, see [Creating and Using Oracle Solaris Zones](#).

▼ How to Assign a PVLAN to a Zone

1. Become an administrator.

For more information, see [“Using Your Assigned Administrative Rights” in *Securing Users and Processes in Oracle Solaris 11.3*](#).

2. Create a VLAN.

```
# dladm create-vlan [-l link -v VLAN-ID[,PVLAN-SVID[,PVLAN-type]] [VLAN-link]
```

3. Verify the VLAN that you created.

```
# dladm show-vlan
```

4. Configure the zone by assigning a PVLAN.

```
global# zonecfg -z zone-name
zonecfg:zone2> add anet
zonecfg:zone2:anet> set vlan-id=VID[,PVLAN-SVID[,PVLAN-type]]
zonecfg:zone2:anet> end
zonecfg:zone2> verify
zonecfg:zone2> commit
zonecfg:zone2> exit
global# zoneadm -z zone reboot
```

PVLAN is assigned as the underlying link for the zone's anet.

Example 22 Assigning a PVLAN to a Zone

This example shows how to assign a PVLAN to a zone. In this example, the zone is assumed to have been created.

```
# dladm create-vlan -l net0 -v 100 vlan10
# dladm show-vlan
LINK      VID   SVID   PVLAN-TYPE  FLAGS   OVER
vlan10    100   --     ---         ----    net0
global# zonecfg -z zone1
zonecfg:zone2> add anet
zonecfg:zone2:anet> set vlan-id=100,200,community
zonecfg:zone2:anet> end
zonecfg:zone2> verify
zonecfg:zone2> commit
zonecfg:zone2> exit
global# zoneadm -z zone reboot
```

Administering Bridging Features

You can use bridging to connect separate network segments so that they communicate as if they were a single network segment. This chapter describes how to configure and administer bridged networks.

This chapter contains the following topics:

- [“Overview of Bridged Networks” on page 79](#)
- [“Creating a Bridge” on page 85](#)
- [“Modifying the Protection Type for a Bridge” on page 87](#)
- [“Adding Links to an Existing Bridge” on page 87](#)
- [“Removing Links From a Bridge” on page 88](#)
- [“Displaying Bridge Configuration Information” on page 89](#)
- [“Deleting a Bridge From the System” on page 91](#)
- [“Administering VLANs on Bridged Networks” on page 92](#)
- [“Debugging Bridges” on page 93](#)

Overview of Bridged Networks

Bridges connect various nodes in the network into a single network. The network segments share a single broadcast network and communicate as if they were a single network segment when connected. As a result, each node can reach the other nodes by using network protocols such as IP rather than using routers to forward traffic across network segments. If you do not use a bridge, you must configure IP routing to permit the forwarding of IP traffic between nodes.

Although you can use both bridging and routing to distribute information about the locations of resources on the network, they differ in several ways. Routing is implemented at the IP layer (L3) and uses routing protocols. No routing protocols are used on the datalink layer.

Bridging is used to distribute information about the locations of resources on the network. In a bridged network, the destinations of forwarded packets are determined by examining the

network traffic that is received on the links that are attached to the bridge. A bridged network uses protocols, such as Spanning Tree Protocol (STP) and Transparent Interconnection of Lots of Links (TRILL). For more information, see [“Bridging Protocols” on page 83](#).




Caution - Do not set the `local-mac-address?` property to false by using the `eeeprom` command on SPARC® based systems that use bridging. Doing so causes these systems to incorrectly use the same MAC address on multiple ports and on the same network.

Simple Bridged Network

The following figure shows a simple bridged network configuration.

FIGURE 11 Simple Bridged Network

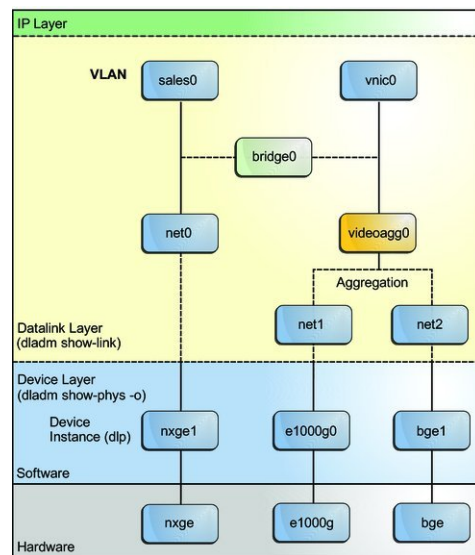


The bridge, `goldengate`, is an Oracle Solaris system that has bridging configured. The systems `sanfrancisco` and `sausalito` are physically connected to the bridge. Network A uses a hub that is physically connected to the bridge on one side and to three computer systems on the other side. The bridge ports are the links `net0`, `net1`, and `net2`.

How Oracle Solaris Bridges Are Implemented in the Network Stack

In Oracle Solaris, you can configure bridges on the datalink layer of the same network stack implementation, as shown in the following figure.


FIGURE 12 Bridges in the Network Stack for Oracle Solaris



Two interfaces, `net0` and `videoagg0`, are configured as a bridge, `bridge0`. Packets that are received on one interface are forwarded to the other interface. After the bridge configuration, both interfaces can still be used to configure VLANs and IP interfaces.

Bridged Network Ring

Bridged networks can be formed into rings that physically connect several bridges together. The following figure shows a bridged network ring configuration.

FIGURE 13 Bridged Network Ring

The figure shows a bridged network that is configured in a ring. The configuration shows three bridges. Two systems are physically connected to the westminister bridge. One system is physically connected to the waterloo bridge and one system is physically connected to the tower bridge. The bridges are physically connected to each other through the bridge ports.

This type of configuration can cause problems with old packets saturating the network links by endlessly looping around the ring. To protect against such looping conditions, Oracle Solaris bridges implement both the STP and TRILL protocols. Note that most hardware bridges also implement STP loop protection.

How a Bridged Network Works

When a packet is received by the bridge, its source address is examined. The source address of the packet associates the node from which the packet was sent with the link on which it is received. Thereafter, when a received packet uses that same address as the destination address, the bridge forwards the packet over the link to that address.

The link that is associated with a source address might be an intermediate link that is connected to another bridge in the bridged network. Over time, all of the bridges within the bridged

network “learn” which of the links sends a packet toward a given node. Therefore, destination address of the packet is used to direct the packet to its final destination by means of hop-by-hop bridging.

A local “link-down” notification indicates that all nodes on a given link are no longer reachable. In this situation, packet forwarding to the link is halted and all forwarding entries over the link are flushed. Older forwarding entries are also flushed over time. When a link is restored, packets that are received over the link are treated as new. The learning process begins again, based on the source address of a packet. This process enables the bridge to properly forward packets over that link when the address is used as the destination address.

Bridging Protocols

Bridged networks use the following protocols:

- **Spanning Tree Protocol (STP)**

STP is the default protocol that is used by the bridged networks. Bridging uses the STP mechanism to prevent network loops that potentially render the subnetworks unusable. To forward packets to their destinations, bridges must listen in promiscuous mode on every link that is attached to the bridge. Listening in promiscuous mode causes bridges to become vulnerable to the occurrences of forwarding loops, in which packets infinitely circle at full-line rate.



Caution - Do not configure a link into a bridge when the highest possible levels of performance are required. Bridging requires the underlying interfaces to be in promiscuous mode, which disables a number of important optimizations that are in the hardware (NIC), driver, and other layers of the system. The disabling of these performance enhancements is an unavoidable consequence of the bridging mechanism.

These performance issues only affect links that are configured to be part of a bridge. You can use a bridge on a system where some of the links of the system are not bridged and are hence not subject to those constraints.

- **Transparent Interconnection of Lots of Links (TRILL)**

Oracle Solaris supports the TRILL protection enhancement, which avoids loops without disabling links. TRILL helps to load-balance the traffic between several paths to the destination.

When STP is used for loop protection, the physical loop is mitigated by preventing one of the connections in the loop from forwarding packets. [Figure 13, “Bridged Network Ring,” on page 82](#) shows that the physical link between the westminster and tower bridges is not used to forward packets.

Unlike STP, TRILL does not shut down physical links to prevent loops. Instead, TRILL computes the shortest-path information for each TRILL node in the network and uses that information to forward packets to individual destinations.

You can use TRILL by specifying the `-P trill` option in the `dladm create-bridge` or `dladm modify-bridge` commands. For more information, see [“Creating a Bridge” on page 85](#) and [“Modifying the Protection Type for a Bridge” on page 87](#).

For information about STP, see IEEE 802.1D-1998. For information about TRILL, see the [Internet Engineering Task Force \(IETF\) TRILL draft documents \(http://tools.ietf.org/wg/trill\)](http://tools.ietf.org/wg/trill).

STP Daemon

Each bridge that you create by using the `dladm create-bridge` command is represented as an identically named Service Management Facility (SMF) instance of `svc:/network/bridge`. Each instance runs a copy of the `/usr/lib/bridged` daemon, which implements the STP.

For example, the following command creates a bridge called `pontevecchio`:

```
# dladm create-bridge pontevecchio
```

The system creates an SMF service instance called `svc:/network/bridge:pontevecchio` and an observability node called `/dev/net/pontevecchio0`. The observability node is intended for use with the `snoop` command and the `wireshark` packet analyzer. You can use the `d1stat` command to obtain the run time statistics of the bridge.

For safety purposes, all ports run standard STP by default. A bridge that does not run some form of bridging protocol, such as STP, can form long-lasting forwarding loops in the network. Because Ethernet has no hop-count or time-to-live (TTL) on packets, any such loops are fatal to the network.

When a particular port is not connected to another bridge (for example, because the port has a direct point-to-point connection to a system), you can administratively disable STP for that port. Even if all ports on a bridge have STP disabled, the STP daemon still runs for the following reasons:

- To handle any new ports that are added
- To implement BPDU guarding
- To enable or disable forwarding on the ports, as necessary

When a port has STP disabled, the bridged daemon continues to listen for BPDUs (BPDU guarding). The daemon uses `syslog` to flag any errors and disables forwarding on the port to indicate a serious network misconfiguration. The link is re-enabled when the link goes down and comes up again, or when you manually remove the link and add it again.

If you disable the SMF service instance for a bridge, the bridge stops on those ports as the STP daemon is stopped. If the instance is restarted, STP starts from its initial state.

TRILL Daemon

Each bridge that you create by using the `dladm create-bridge -P trill` command is represented by an identically named SMF instance of `svc:/network/bridge` and `svc:/network/routing/trill`. Each instance of `svc:/network/routing/trill` runs a copy of the `/usr/lib/trilld` daemon, which implements the TRILL protocol.

For example, the following command creates a bridge called `bridgeofsighs`:

```
# dladm create-bridge -P trill bridgeofsighs
```

The system creates two SMF services called `svc:/network/bridge:bridgeofsighs` and `svc:/network/routing/trill:bridgeofsighs`. In addition, the system creates an observability node called `/dev/net/bridgeofsighs0`.

Creating a Bridge

In Oracle Solaris, use the `dladm` command and the SMF feature to administer bridges. You can use SMF commands to enable, disable, and monitor bridge instances by using the fault-managed resource identifier (FMRI) of the instance, `svc:/network/bridge`. You can use the `dladm` command to create or destroy bridges, and to assign links to bridges or to remove links from them. The links that are assigned to the bridge must be an Ethernet type, which includes 802.3 and 802.11 media.

To create a bridge between links, you must create at least one bridge instance. Each bridge instance is separate. Bridges do not include a forwarding connection between them, and a link is a member of a maximum of one bridge.

The `dladm create-bridge` command creates a bridge instance and optionally assigns one or more network links to the new bridge. Because no bridge instances are present on the system by default, Oracle Solaris does not create bridges between network links by default.

To create a bridge, use the following command:

```
# dladm create-bridge [-P protect] [-p priority] [-d forward-delay] [-l link...] bridge-name
```

`-P protect` Specifies the protection method. It can be set to one of the following values.

- `stp` – STP protection method (the default)
- `trill` – TRILL protection method

<code>-p priority</code>	Specifies an IEEE STP priority value for a bridge to determine the root bridge node in the network. The default value is 32768. Valid values are from 0 (highest priority) to 61440 (lowest priority), in increments of 4096.
<code>-d forward-delay</code>	Specifies the STP forward delay parameter for the bridge. When the bridge that is created is the root node, all the bridges in the network use this timer to sequence the link states when a port is enabled. The default value is 15 seconds. Valid values are from 4 to 30 seconds.
<code>-l link</code>	Adds a link to the bridge. If any of the specified links cannot be added, the command fails and the bridge is not created.

bridge-name is an arbitrary string that must be a legal SMF service instance name. This name is an FMRI component that has no escape sequences, which means that white space, ASCII control characters, and the following characters cannot be present:

`; / ? : @ & = + $, % < > # "`

The name `default` and all names beginning with the `SUNW` string are reserved. Names that have trailing digits are reserved for the creation of observability devices, which are used for debugging. Because of the use of observability devices, the names of legal bridge instances are further constrained to be a legal `dlpi` name. The name must begin and end with an alphabetic character or an underscore character. The rest of the name can contain alphanumeric and underscore characters.

For more information about bridge creation options, see the description of the `dladm create-bridge` command in the [dladm\(1M\)](#) man page.

EXAMPLE 23 Creating a Bridge

The following example shows how to create the `brooklyn` bridge by connecting the `net0` and `net1` links.

```
# dladm create-bridge -P stp -d 12 -l net0 -l net1 brooklyn
# dladm show-bridge
BRIDGE      PROTECT ADDRESS          PRIORITY DESROOT
goldengate  stp      32768/8:0:20:bf:f 32768    8192/0:d0:0:76:14:38
brooklyn    stp      32768/8:0:20:e5:8 32768    8192/0:d0:0:76:14:38
```

The following example shows how to create the `westminister` bridge by connecting the `net0` and `net1` links.

```
# dladm create-bridge -P trill -l net0 -l net1 westminister
# dladm show-bridge
BRIDGE      PROTECT ADDRESS          PRIORITY DESROOT
goldengate  stp      32768/8:0:20:bf:f 32768    8192/0:d0:0:76:14:38
westminister trill    32768/8:0:20:e5:8 32768    8192/0:d0:0:76:14:38
```

Modifying the Protection Type for a Bridge

STP is a mechanism to prevent network loops that could render the subnetworks unusable. In addition to using STP for bridges, Oracle Solaris supports the TRILL protection enhancement. STP is used by default, but you can use TRILL by specifying the `-P trill` option for the bridging commands.

To modify the protection type from STP to TRILL or from TRILL to STP, use the following command:

```
# dladm modify-bridge -P protection-type bridge-name
```

The `-P protection-type` option specifies which protection type to use, either `stp` (the default) or `trill`.

EXAMPLE 24 Modifying the Protection Type for a Bridge

The following example shows how to modify the protection type for the `brooklyn` bridge from the default STP to TRILL.

```
# dladm modify-bridge -P trill brooklyn
```

The following example shows how to change the protection type for the `brooklyn` bridge from TRILL to STP.

```
# dladm modify-bridge -P stp brooklyn
```

Adding Links to an Existing Bridge

A link can be a member of at most one bridge. So, if you want to move a link from one bridge instance to another bridge instance, you must first remove the link from the current bridge before adding it to another bridge.

Links that are assigned to the same bridge must have the same MTU value. Although you can change the MTU value on an existing link, the bridge instance goes into maintenance state until you remove or change the assigned links so that the MTU values match before you restart the bridge.

Note - The links that are assigned to a bridge cannot be VLANs, VNICs, or tunnels. Only links that are acceptable as part of an aggregation or links that are aggregations can be assigned to a bridge.

To add a new link to an existing bridge, use the following command:

```
# dladm add-bridge -l new-link bridge-name
```

The following example shows how to add the net2 link to the existing bridge rialto.

```
# dladm add-bridge -l net2 rialto
```

Removing Links From a Bridge

Before you delete any bridge, all of its links must first be removed. To remove links, use the following command:

```
# dladm remove-bridge [-l link]... bridge-name
```

The following example shows how to remove the net0, net1, and net2 links from the bridge charles.

```
# dladm remove-bridge -l net0 -l net1 -l net2 charles
```

Setting Link Properties for a Bridge

You can set the following link properties for bridges:

default_tag	The default VLAN ID for untagged packets that are sent to a link and received from a link. The valid values are 0 to 4094. The default value is 1.
forward	Enables and disables traffic forwarding through the bridge. This property exists on all links except VNIC links. The valid values are 1 (true) and 0 (false). The default value is 1.
stp	Enables and disables STP and RSTP. Valid values are 1 (true) and 0 (false). The default value is 1, which enables STP and RSTP.
stp_cost	Represents the STP and RSTP cost values for using the link. Valid values are from 1 to 65535. The default value is 0, which is used to signal that cost is automatically computed by link type.
stp_edge	Specifies whether the port is connected to other bridges. Valid values are 1 (true) and 0 (false). The default value is 1.
stp_p2p	Specifies the connection mode type. Valid values are true, false, and auto. The default value is auto.

`stp_priority` Sets the STP and RSTP port priority value. Valid values are from 0 to 255. The default value is 128.

For more information, see the [dladm\(1M\)](#) man page.

To modify the link properties of a bridge, use the following command:

```
dladm set-linkprop -p prop=value link
```

EXAMPLE 25 Setting Link Properties for a Bridge

The following example shows how to disable traffic forwarding and set the connection mode type. To set the properties for the bridge, you must set the properties on the links that connect the bridge.

```
# dladm create-bridge -P stp -d 12 -l net0 -l net1 brooklyn
# dladm set-linkprop -p forward=0 net0
# dladm set-linkprop -p stp_p2p=true net1
```

The following example shows how to reset multiple properties of a bridge.

```
# dladm reset-linkprop -p default_tag,stp_priority brooklyn
```

Displaying Bridge Configuration Information

You can display bridge configuration information by using the `dladm show-bridge` command.

Displaying Information About Configured Bridges

You can use the `dladm show-bridge` and `d1stat show-bridge` commands to show different kinds of information about configured bridges.

Use the following command options:

- To view the list of bridges:

```
# dladm show-bridge
```

```
# dladm show-bridge
```

BRIDGE	PROTECT	ADDRESS	PRIORITY	DESROOT
goldengate	stp	32768/8:0:20:bf:f	32768	8192/0:d0:0:76:14:38
baybridge	stp	32768/8:0:20:e5:8	32768	8192/0:d0:0:76:14:38

- To show the link-related status of a bridge:

- ```
dladm show-bridge -l bridge-name
```

    - To show link-related statistics for the bridge:
  - ```
# dlstat show-bridge bridge-name
```

 - To show kernel forwarding entries for the bridge:
 - ```
dladm show-bridge -f bridge-name
```

    - To show TRILL information about the bridge:
  - ```
# dladm show-bridge -t bridge-name
```

 - To show statistics of each bridge and statistics of the links connected to each bridge:
- ```
dlstat show-bridge
```
- | BRIDGE  | LINK     | IPKTS | RBYTES  | OPKTS | OBYTES  | DROPS | FORWARDS |
|---------|----------|-------|---------|-------|---------|-------|----------|
| rbblue0 | --       | 1.93K | 587.29K | 2.47K | 3.30M   | 0     | 0        |
|         | simblue1 | 72    | 4.32K   | 2.12K | 2.83M   | 0     | --       |
|         | simblue2 | 1.86K | 582.97K | 348   | 474.04K | 0     | --       |
| stbred0 | --       | 975   | 976.69K | 3.44K | 1.13M   | 0     | 38       |
|         | simred3  | 347   | 472.54K | 1.86K | 583.03K | 0     | --       |
|         | simred4  | 628   | 504.15K | 1.58K | 551.51K | 0     | --       |
- To show all statistics of each bridge and statistics of the links connected to each bridge:
- ```
# dlstat show-bridge -o all
```

For more information about the `dladm show-bridge` command options, see the [dladm\(1M\)](#) man page and for information about the `dlstat show-bridge` command options, see the [dlstat\(1M\)](#) man page.

EXAMPLE 26 Displaying Bridge Information

The following examples show how to use the `dladm show-bridge` command with various options.

- The following command shows link-related status information for a single bridge instance, `tower`. To view configured properties, use the `dladm show-linkprop` command.

```
# dladm show-bridge -l tower
```

LINK	STATE	UPTIME	DESROOT
net0	forwarding	117	8192/0:d0:0:76:14:38
net1	forwarding	117	8192/0:d0:0:76:14:38

- The following command shows the kernel forwarding entries for the specified bridge, `avignon`:

```
# dladm show-bridge -f avignon
```

DEST	AGE	FLAGS	OUTPUT
------	-----	-------	--------

```

8:0:20:bc:a7:dc  10.860  --    net0
8:0:20:bf:f9:69  --      L    net0
8:0:20:c0:20:26  17.420  --    net0
8:0:20:e5:86:11  --      L    net1

```

- The following command shows the TRILL information about the specified bridge, key:

```

# dladm show-bridge -t key
NICK  FLAGS LINK      NEXTHOP
38628  --    london  56:db:46:be:b9:62
58753  L      --      --

```

Displaying Configuration Information About Bridge Links

You use the `dladm show-link` command with the `-o all` option to display the `BRIDGE` field in the output. If a link is a member of a bridge, this field identifies the name of the bridge of which it is a member. For links that are not part of a bridge, the field is blank if the `-p` option is used. Otherwise, the field shows `--`.

The observability node of the bridge also appears in the `dladm show-link` output as a separate link. For this node, the existing `OVER` field lists the links that are members of the bridge.

Use the following command to view configuration information about any link that is a member of a bridge.

```
# dladm show-link [-p]
```

The `-p` option produces output in a parseable format.

Deleting a Bridge From the System

Before you delete a bridge, you must first remove any links attached to the bridge.

▼ How to Delete a Bridge From the System

1. **Become an administrator.**

For more information, see [“Using Your Assigned Administrative Rights” in *Securing Users and Processes in Oracle Solaris 11.3*](#).

2. **Remove any links attached to the bridge.**

```
# dladm remove-bridge [-l link]... bridge-name
```

3. **Delete the bridge from the system.**

```
# dladm delete-bridge bridge-name
```

Example 27 Deleting a Bridge From the System

The following example shows how to remove the net0, net1, and net2 links from the coronado bridge, and then delete the bridge itself from the system.

```
# dladm remove-bridge -l net0 -l net1 -l net2 coronado
# dladm delete-bridge coronado
```

Administering VLANs on Bridged Networks

By default, VLANs that are configured on the system forward packets among all the ports on a bridge instance. When you invoke the `dladm create-vlan` or `dladm create-vnic -v` command and the underlying link is a part of a bridge, the command also enables packet forwarding of the specified VLAN on that bridge link. For more information about VLANs, see [Chapter 3, “Configuring Virtual Networks by Using Virtual Local Area Networks”](#).

To configure a VLAN on a link and disable packet forwarding to or from other links on the bridge, you must disable forwarding by setting the `forward` property for the VLAN with the `dladm set-linkprop` command. For more information, see [“Setting Link Properties for a Bridge” on page 88](#).

▼ How to Configure VLANs Over a Datalink That Is Part of a Bridge

Before You Begin This procedure assumes that the bridge already exists. For information about how to create a bridge, see [“Creating a Bridge” on page 85](#).

1. **Become an administrator.**

For more information, see [“Using Your Assigned Administrative Rights” in *Securing Users and Processes in Oracle Solaris 11.3*](#).

2. **List the link-related information of the bridge to determine the links that are a part of the bridge.**

```
# dladm show-bridge -l bridge-name
```

3. Create the VLAN over the link that is a part of the bridge.

```
# dladm create-vlan -l link -v vid VLAN-link
```

link Specifies the link on which the VLAN interface is being created.

Note - In this procedure, the link should part of the bridge that you have created.

vid Indicates the VLAN ID number.

VLAN-link Specifies the name of the VLAN.

4. Repeat this command for every VLAN you want to create. For every VLAN that you created, create an IP interface over the VLAN.

```
# ipadm create-ip interface
```

where *interface* is the VLAN name.

5. For each IP interface on a VLAN, configure a valid IP address.

```
# ipadm create-addr -a IP-address interface
```

VLANs and the STP and TRILL Protocols

VLANs are ignored by the standards-compliant STP. The bridging protocol computes just one loop-free topology by using tag-free BPDU messages, and uses this tree topology to enable and disable links. You must configure any duplicate links that are provisioned in your networks such that when those links are automatically disabled by STP, the configured VLANs are not disconnected. You must either run all VLANs everywhere on your bridged backbone or carefully examine all redundant links.

The TRILL protocol does not follow the complex STP rules. Instead, TRILL automatically encapsulates packets that have the VLAN tag intact and passes them through the network.

Debugging Bridges

Each bridge instance is assigned an *observability node*, which appears in the `/dev/net/` directory and is named with the bridge name plus a trailing `0`, for example, `/dev/net/bridgeofsighs0`.

The observability node is intended for use with the `snoop` command and the `wireshark` packet analyzer. This node operates like a standard Ethernet interface except for the transmission of packets, which are silently dropped. You cannot plumb IP on top of an observability node, and you cannot perform bind requests (`DL_BIND_REQ`) unless you use the `passive` option, which enables you only to receive packets and not to send them.

The observability node makes a single unmodified copy of every packet handled by the bridge. It is available to the user for monitoring and debugging. This behavior is similar to monitoring a port on a traditional bridge and is subject to the usual datalink provider interface (DLPI) promiscuous mode rules. You can also use the `pfmmod` command or features in the `snoop` command and the `wireshark` packet analyzer to filter packets based on the VLAN ID.

The delivered packets, which are the packets that are sent to the observability node, represent the data received by the bridge.

Note - When the bridging process adds, removes, or modifies a VLAN tag, the data shown by the `snoop` command and the `wireshark` packet analyzer describes the state prior to the processes taking place. This rare situation might be confusing if distinct `default_tag` values are used on different links.

To see the packets that are transmitted and received on a particular link after the bridging process is complete, run the `snoop` command on the individual links rather than on the bridge observability node.

You can also obtain statistics about how network packets use network resources on links by using the `dlstat` command. For information, see [Chapter 8, “Monitoring Network Traffic and Resource Usage”](#) in *Managing Network Virtualization and Network Resources in Oracle Solaris 11.3*.

Exchanging Network Connectivity Information With Link Layer Discovery Protocol

This chapter describes how to enable systems to exchange system and network connectivity information throughout the local network by using the Link Layer Discovery Protocol (LLDP).

This chapter contains the following topics:

- [“Overview of LLDP” on page 95](#)
- [“Information the LLDP Agent Advertises” on page 98](#)
- [“Enabling LLDP on the System” on page 101](#)
- [“Specifying TLV Units and Values for the LLDP Packet of an Agent” on page 105](#)
- [“Disabling LLDP” on page 109](#)
- [“Monitoring LLDP Agents” on page 110](#)

Overview of LLDP

LLDP is used by systems in a local area network (LAN) to exchange configuration and management information with each other. With this protocol, a system can advertise connectivity and management information to other systems on the network. This information can include system capabilities, management addresses, and other information relevant to network operations. This protocol also enables systems to receive similar information about other systems that are on the same local network.

On any LAN, individual components such as systems and switches are not configured in isolation. To host network traffic efficiently, the configuration of systems on the network must be coordinated with each other.

When you manually configure each system, switch, and other components, ensuring compatibility among the components is a challenge. The manual configuration of systems is risky and can easily cause misconfigurations, particularly if different administrators work independently on different systems. A better alternative is to use LLDP, which enables systems to transmit their individual configuration information to peer systems and helps to detect any misconfigurations.

Oracle Solaris supports the use of LLDP to promote the exchange of system and network connectivity information between systems on the network, which reduces the risk of misconfigured network resources.

In this release, LLDP is used by the network diagnostics service to automatically detect problems that could lead to limited or degraded network connectivity, or both. Enabling the LLDP service enhances the ability to perform network diagnostics on your Oracle Solaris system. For more information about network diagnostics, see [Chapter 5, “Performing Network Diagnostics With the network-monitor Transport Module Utility” in *Troubleshooting Network Administration Issues in Oracle Solaris 11.3*](#).

In Oracle Solaris, LLDP is also used to exchange data center bridging exchange protocol (DCBX) Type-Length-Value (TLV) units. DCBX provides configuration information about DCB features such as priority-based flow control (PFC) and enhanced transmission selection (ETS). For more information about DCB, see [Chapter 7, “Managing Converged Networks by Using Data Center Bridging”](#).

With LLDP, the system administrator can easily detect faulty system configurations, particularly in complex networks such as virtual local area networks (VLANs) and link aggregations. Information about the network topology can be obtained readily without having to trace physical connections between servers, switches, and other devices that comprise the network.

Components of an LLDP Implementation

LLDP is implemented with the following components:

- **LLDP package** – You install this package to enable the LLDP. This package includes the LLDP daemon, command-line utilities, the service manifest and scripts, and other components that are required for LLDP to operate.
- **LLDP service** – You can enable the LLDP service by using the `svcadm` command. This service uses the fault management resource identifier (FMRI) of the service management facility (SMF) service instance, `svc:/network/lldp:default`, to manage the LLDP daemon, `lldpd`. This LLDP service is responsible for starting, stopping, restarting, or refreshing the `lldpd` daemon. This service is automatically enabled after you install the LLDP package.
- **lldpadm command** – You can use this command to administer LLDP on individual links and to configure the operating mode of LLDP, to specify TLV units that are transmitted, and to configure DCBX TLV units. For information about TLV units, see [“Information the LLDP Agent Advertises” on page 98](#).

You must use this command to set the per-agent LLDP properties and global LLDP properties and to get LLDP information for a specific agent or its peer.

The `lldpadm` subcommands are described in the following sections. For more information about the `lldpadm` command, see the [`lldpadm\(1M\)`](#) man page.

- **LLDP daemon** – The LLDP services manage LLDP agents on the system. They also interact with `snmpd`, the daemon for the Simple Network Management Protocol (SNMP), to retrieve LLDP information that is received on the system through SNMP.
- **LLDP agents** – LLDP agents are the LLDP instances that are associated with a physical datalink on which LLDP is enabled. LLDP agents transmit information about the datalink to its peer and also receive information from the peer. You can configure an LLDP agent to advertise specific information about the associated physical datalink. You can enable LLDP only on physical datalinks.

Information Sources of the LLDP Agent

The LLDP agent transmits and receives LLDP data units (LLDPDUs). The agent manages and stores information that is contained in these LLDPDUs in the following types of data stores:

- **Local management information base (MIB)** – This data store contains network information that pertains to a system's specific link on which the LLDP agent is enabled. A local MIB contains both common and unique information. For example, the chassis ID is common information that is shared among all the LLDP agents on the system. However, port IDs for the system's datalinks are different. Therefore, each agent manages its own local MIB.
- **Remote MIB** – Information in this data store is received from LLDP agents of peer hosts.

LLDP Agent Modes

The LLDP agent operates in the following modes:

- **Transmit only (txonly)** – The LLDP agent does not process incoming LLDPDUs. Therefore, the remote MIB is empty.
- **Receive only (rxonly)** – The agent processes only incoming LLDPDUs and stores the information in remote MIBs. However, no information from the local MIB is transmitted.
- **Transmit and receive (both)** – The agent transmits local information and processes incoming LLDPDUs and therefore maintains both local and remote MIBs.
- **Disabled (disable)** – The agent does not exist.

For information about setting agent modes, see [“How to Enable LLDP for Specific Ports” on page 103](#).

Information the LLDP Agent Advertises

The LLDP agent transmits system and connectivity information in the LLDP packets or LLDPDUs. These packets contain information units that are individually formatted in TLV format. The information units are also called *TLV units*.

Mandatory TLV Units

Certain TLV units are mandatory and are included in the LLDP packets by default when LLDP is enabled. You cannot use the `lldpadm` command to exclude any of these units.

The following TLV units are mandatory:

- Chassis ID – Information that is generated by the `hostid` command
- Port ID – MAC address of the physical NIC
- TTL (time to live)
- End of protocol data unit (PDU)

Multiple LLDP agents can be enabled on a single system depending on the number of links. The chassis ID and port ID combination uniquely identifies an agent and distinguishes it from other agents on the system.

EXAMPLE 28 Displaying the Chassis ID and Port ID

The following example displays the chassis ID and port ID for an LLDP agent.

```
# hostid
004e434e

# dladm show-phys -m net4
LINK           SLOT   ADDRESS              INUSE CLIENT
net4           primary 0:1b:21:87:8b:b4    yes  net4

# lldpadm show-agent -l net4
AGENT          CHASSISID          PORTID
net4           004e434e           00:1b:21:87:8b:b4
```

Oracle Solaris LLDP agents use `hostid` as the chassis ID and the port's MAC address as the port ID.

Optional TLV Units

Optional TLV units can be added to an LLDP packet. These optional TLV units enable vendors to insert vendor-specific TLV units to be advertised. LLDP enables defining additional TLV

units by using organization unique identifiers (OUIs). An OUI identifies the category for a TLV unit depending on whether the OUI follows the IEEE 802.1 or IEEE 802.3 standard. The LLDP agent properties can be configured to enable or disable the transmission of these optional TLV units.

The following table lists each TLV group, its corresponding name, and the TLV units for each property, and their descriptions. You configure any one of these properties to specify the TLV units to be included in the packets when LLDP is enabled.

TABLE 2 Optional TLV Units for an LLDP Agent

TLV Group	TLV Name	TLV units	Description
Basic management	basic-tlv	sysname, portdesc, syscapab, sysdesc, mgmtaddr	Specifies the system name, port description, system capability, system description, and management address to be advertised.
802.1 OUI	dot1-tlv	vlanname, pvid, linkaggr, pfc, appln, evb, etscfg, etsreco	Specifies the following to be advertised: VLAN name, port VLAN ID, link aggregation, TLV units for priority-based flow control, application, enhanced transmission selection, and edge virtual bridging.
802.3 OUI	dot3-tlv	max-framesize	Specifies the maximum frame size to be advertised.
Oracle specific OUI (which is defined as 0x0003BA)	virt-tlv	vnic	Specifies the VNIC to be advertised if a virtual network is configured.

TLV Unit Properties

Each TLV unit has properties that you can further configure with specific values. If the TLV unit is enabled as an LLDP agent's property, then that TLV unit is advertised in the network only with the specified values. For example, consider the TLV unit `syscapab`, which advertises a system's capabilities. These capabilities can potentially include support for routers, bridges, repeaters, telephones, and other devices. However, you can set `syscapab` so that only those capabilities that are actually supported on your specific system, such as routers and bridges, are advertised.

The procedure for configuring TLV units depends on whether you are configuring *global TLV units* or *per-agent TLV units*. For information about how to configure TLV units, see [“Specifying TLV Units and Values for the LLDP Packet of an Agent” on page 105](#).

Global TLV units apply to all LLDP agents on the system. The following table lists the global TLV units and their corresponding possible configurations.

TABLE 3 Global TLV Units and Their Properties

TLV Unit	Property Name	Possible Property Values	Value Description
syscapab	supported	other, repeater, bridge, wlan-ap, router, telephone, docsis-cd, station, cvlan, sylvan, tpmr	Represents the primary supported functions of the system. Default values are router, station, and bridge.
	enabled	Subset of the values listed for supported	Represents the enabled functions of the system.
mgmtaddr	ipaddr	ipv4 or ipv6	Specifies the type of IP addresses that are associated with the local LLDP agent. The addresses are used to reach higher layer entities and assist in discovery by network management. Only one type can be specified.

TLV units that are specific to LLDP agent are managed on a per-agent basis. With per-agent TLV units, the values that you provide are used when the TLV unit is enabled for transmission by a specific LLDP agent.

The following table lists the TLV values and their corresponding possible configurations for an LLDP agent.

TABLE 4 Per-Agent TLV Units and Their Properties

TLV Unit	Property Name	Possible Property Values	Value Description
pfc	willing	on, off	Sets an LLDP agent to accept or reject configuration information from a remote system that pertains to priority-based flow control.
appln	apt	Values are taken from the information that is defined in the Application Priority Table.	Configures the Application Priority Table. This table contains the list of application TLV units and their corresponding priorities. The application is identified by the id/selector pair. The contents of the table use the following format: id/selector/priority For more information, see “Application Priority Configurations” on page 126 .
etscfg	willing	on, off	Sets an LLDP agent to accept or reject configuration information from a remote system that pertains to enhanced transmission selection.

For information about per-agent TLV units, see [Chapter 7, “Managing Converged Networks by Using Data Center Bridging”](#).

Enabling LLDP on the System

You can configure LLDP to exchange system information with other systems or peers on the network.

The SMF property `auto-enable-agents` controls the way in which you can enable LLDP agents on the system. With this property, you can choose to enable LLDP globally across all the physical links or only one physical link at a time.

The SMF property `auto-enable-agents` can have one of the following three possible values:

- `yes` enables LLDP on all ports in the transmit and receive (both) mode provided that no previous LLDP configuration exists on a port. If a configuration exists on a port, then that port's configuration is retained. For example, if a port has been previously configured with LLDP in the `rxonly` mode, then the LLDP service will not switch the agent to run in the transmit and receive (both) mode. LLDP on that port continues to be in the `rxonly` mode. This is the default value of the SMF property `auto-enable-agents`.
- `force` enables LLDP in the transmit and receive (both) mode on all ports and overrides any existing LLDP configurations on any port. For example, if a previous LLDP configuration on a port runs in the `rxonly` mode, the LLDP agent is switched to run in the transmit and receive (both) mode, which is the default LLDP mode.
- `no` disables automatic enabling of LLDP on all ports except those with existing LLDP configurations. On these ports, the existing LLDP configuration is retained.

Note - Every time you customize the `auto-enable-agents` property, you must restart the LLDP service for the new value to become effective.

▼ How to Install the LLDP Package

By default, LLDP is enabled and ready to be used after you have completed installing the LLDP package.

1. **Become an administrator.**

For more information, see [“Using Your Assigned Administrative Rights” in *Securing Users and Processes in Oracle Solaris 11.3*](#).

2. **Install the package.**

```
# pkg install lldp
```

3. **Determine whether the LLDP service has started.**

```
# svcs lldp
STATE      STIME    FMRI
online     Jul_10   svc:/network/lldp:default
```

If the LLDP service is disabled, start the service with the following command:

```
# svcadm enable svc:/network/lldp:default
```

▼ How to Enable LLDP Globally

Before You Begin In order to enable LLDP, you must first install the LLDP package. For more information, see [“How to Install the LLDP Package” on page 101](#).

- 1. Become an administrator.**

For more information, see [“Using Your Assigned Administrative Rights” in *Securing Users and Processes in Oracle Solaris 11.3*](#).

- 2. Change the SMF `auto-enable-agents` property to yes if it is set to no.**

```
# svccfg -s svc:/network/lldp:default setprop lldp/auto-enable-agents = "yes"
```

By default, this property is set to yes.

- 3. Restart the LLDP service.**

```
# svcadm restart svc:/network/lldp:default
```

- 4. (Optional) Customize global TLV units.**

```
# lldpadm set-tlvprop -p property=value global-TLV
```

where *property* refers to the property of the global TLV unit.

Next Steps For an explanation of global TLV units, see [“TLV Unit Properties” on page 99](#).

To display a list of global TLVs, type `lldpadm show-tlvprop` or refer to [Table 3, “Global TLV Units and Their Properties,” on page 100](#).

For instructions about how to define TLV values, see [“How to Define TLV Units” on page 107](#).

For information about the `lldpadm` command, see the [`lldpadm\(1M\)`](#) man page.

▼ How to Enable LLDP for Specific Ports

Before You Begin In order to enable LLDP, you must first install the LLDP package. For more information, see [“How to Install the LLDP Package” on page 101](#).

1. Become an administrator.

For more information, see [“Using Your Assigned Administrative Rights” in *Securing Users and Processes in Oracle Solaris 11.3*](#).

2. Change the SMF `auto-enable-agents` property to `no` if it is set to `yes`.

```
# svccfg -s svc:/network/lldp:default setprop lldp/auto-enable-agents = "no"
```

By default, this property is set to `yes`.

3. Restart the LLDP service if you have changed the SMF property `auto-enable-agents` in step 2.

```
# svcadm restart svc:/network/lldp:default
```

4. Enable LLDP agents on selected ports or links.

```
# lldpadm set-agentprop -p mode=value agent
```

where *agent* is the LLDP agent and is identified by the physical link on which the agent is enabled. For example, if you enable LLDP on `net0`, the agent is `net0`.

The property `mode` can be set to one of four possible values that represent the LLDP agent's modes of operation: `txonly`, `rxonly`, `both`, and `disable`. For an explanation of these values, see [“LLDP Agent Modes” on page 97](#).

5. Specify the TLV units that the LLDP agent can advertise.

```
# lldpadm set-agentprop -p property=value agent
```

For an explanation of the properties of the LLDP agent, see [“Information the LLDP Agent Advertises” on page 98](#).

To display a list of the other properties of the LLDP agent, type `lldpadm show-agentprop` or refer to [Table 2, “Optional TLV Units for an LLDP Agent,” on page 99](#).

For instructions about how to specify TLV units for LLDP packet of an agent, see [“How to Specify TLV Units for the LLDP Packet of an Agent” on page 105](#).

6. (Optional) Customize the per-agent TLV units.

```
# lldpadm set-agenttlvprop -p property=value -a agent per-agent-TLV
```

where *property* refers to the property of the per-agent TLV unit.

For an explanation of per-agent TLV units, see [“TLV Unit Properties” on page 99](#).

To display a list of per-agent TLVs, type `lldpadm show-agenttlvprop` or refer to [Table 4, “Per-Agent TLV Units and Their Properties,” on page 100](#).

For instructions about how to define TLV values, see [“How to Define TLV Units” on page 107](#).

For information about the `lldpadm` command, see the `lldpadm(1M)` man page.

Example 29 Customizing the auto-enable-agents SMF Property

The following example shows the different way in which LLDP is enabled if you change the value of the SMF property `auto-enable-agents`. For example, given a system with four ports, LLDP is configured on two ports as follows:

- `net0`: both mode
- `net1`: rxonly mode
- `net2` and `net3`: none

If the SMF property `auto-enable-agents` has the default value `yes`, LLDP is automatically enabled on `net2` and `net3`. You can display the LLDP configuration as follows:

```
# lldpadm show-agentprop -p mode
```

AGENT	PROPERTY	PERM	VALUE	EFFECTIVE	DEFAULT	POSSIBLE
net0	mode	rw	both	--	disable	txonly, rxonly, both, disable
net1	mode	rw	rxonly	--	disable	txonly, rxonly, both, disable
net2	mode	rw	both	--	disable	txonly, rxonly, both, disable
net3	mode	rw	both	--	disable	txonly, rxonly, both, disable

If you switch the SMF property to `no`, the configuration changes when you restart the service.

```
# svccfg -s svc:/network/lldp:default setprop lldp/auto-enable-agents = "no"
# svcadm restart svc:/network/lldp:default
# lldpadm show-agentprop -p mode
```

AGENT	PROPERTY	PERM	VALUE	EFFECTIVE	DEFAULT	POSSIBLE
net0	mode	rw	both	--	disable	txonly, rxonly, both, disable
net1	mode	rw	rxonly	--	disable	txonly, rxonly, both, disable
net2	mode	rw	disable	--	disable	txonly, rxonly, both, disable
net3	mode	rw	disable	--	disable	txonly, rxonly, both, disable

In the sample output, `net2` and `net3`, whose LLDP modes were previously automatically enabled, are now flagged as disabled. However, no change occurs on `net0` and `net1`, whose LLDP agents were previously configured.

Example 30 Enabling LLDP on Multiple Datalinks

This example shows how to enable LLDP selectively. A system has two datalinks, `net0` and `net1`. On `net0`, to set the agent to transmit and receive LLDP packets, and on `net1`, to set the agent to only transmit LLDP packets, type the following commands:

```
# svccfg -s svc:/network/lldp:default setprop lldp/auto-enable-agents = "no"
# svcadm restart svc:/network/lldp:default
# lldpadm set-agentprop -p mode=both net0
# lldpadm set-agentprop -p mode=txonly net1
# lldpadm show-agentprop -p mode
```

AGENT	PROPERTY	PERM	VALUE	EFFECTIVE	DEFAULT	POSSIBLE
net0	mode	rw	both	--	disable	txonly, rxonly, both, disable
net1	mode	rw	txonly	--	disable	txonly, rxonly, both, disable

Specifying TLV Units and Values for the LLDP Packet of an Agent

You can specify TLV units such as `dot1-tlv` and `basic-tlv` as property values for an LLDP agent. You can further configure these property values. To specify TLV units, use the `lldpadm set-agentprop` command. For more information, see [lldpadm\(1M\)](#) man page. When the TLV unit is specified as a property of an LLDP agent, then that TLV unit is advertised in the network only with the values that you have specified for the TLV units. For information about TLV units and properties see, [“TLV Unit Properties” on page 99](#).

▼ How to Specify TLV Units for the LLDP Packet of an Agent

This procedure explains how to specify TLV units to be advertised in an LLDP packet that an agent transmits.

- 1. Become an administrator.**

For more information, see [“Using Your Assigned Administrative Rights” in *Securing Users and Processes in Oracle Solaris 11.3*](#).

- 2. (Optional) Identify the LLDP agent property that can contain the TLV unit that you want to add by displaying the TLV units.**

```
# lldpadm show-agentprop agent
```

This command helps you to see the TLV units that are already set for each property. If you do not specify a property, this command displays all the LLDP agent properties and their TLV

values. For a list of agent properties, see [Table 2, “Optional TLV Units for an LLDP Agent,” on page 99](#).

3. Add or remove the TLV unit from the property.

```
# lldpadm set-agentprop -p property[+|-]=value[,...] agent
```

You can use qualifiers to add (+) or remove (-) values from the list of values for properties that accept multiple values.

If you do not use the add (+) or remove (-) qualifiers, then the value that you set replaces all the values that were previously defined for the property.

4. (Optional) Display the new values for the property.

```
# lldpadm show-agentprop -p property agent
```

Example 31 Adding Optional TLV Units to an LLDP Packet

In the following example, the LLDP agent has `net0` configured to advertise VLAN information in its LLDP packet. The LLDP packet is further configured to include system capabilities, link aggregation, and virtual NIC information as items that the LLDP can advertise. Later, the VLAN description is removed from the packet.

1. Display the existing agent properties.

```
# lldpadm show-agentprop net0
```

AGENT	PROPERTY	PERM	VALUE	EFFECTIVE	DEFAULT	POSSIBLE
net0	mode	rw	both	--	disable	txonly, rxonly, both, disable
net0	basic-tlv	rw	sysname, sysdesc	--	none	none, portdesc, sysname, sysdesc, syscapab, mgmtaddr, all
net0	dot1-tlv	rw	vlanname, pvid, pfc	--	none	none, vlanname, pvid, linkaggr, pfc, appln, evb, etscfg, etsreco, all
net0	dot3-tlv	rw	max-framesize	--	none	none, max-framesize, all
net0	virt-tlv	rw	none	--	none	none, vnic, all

The output displays the existing, default, and possible values for each property of the LLDP agent.

2. Set system capabilities, link aggregation, and network virtualization information as items to advertise over the network.

```
# lldpadm set-agentprop -p basic-tlv+=syscapab,dot1-tlv+=linkaggr,virt-tlv=vnic net0
```

3. Remove the VLAN description from the packet.

```
# lldpadm set-agentprop -p dot1-tlv==vlanname net0
```

4. Display the agent properties.

```
# lldpadm show-agentprop -p net0
```

AGENT	PROPERTY	PERM	VALUE	EFFECTIVE	DEFAULT	POSSIBLE
net0	mode	rw	both	--	disable	txonly, rxonly, both, disable
net0	basic-tlv	rw	sysname, sysdesc, syscapab	--	none	none, portdesc, sysname, sysdesc, syscapab, mgmtaddr, all
net0	dot1-tlv	rw	pvid, pfc linkaggr	--	none	none, vlanname, pvid, linkaggr, pfc, appln, evb, etscfg, etsreco, all
net0	dot3-tlv	rw	max-framesize	--	none	none, max-framesize, all
net0	virt-tlv	rw	vnuc	--	none	none, vnuc, all

▼ How to Define TLV Units

This procedure explains how to provide values for specific TLV units.

Tip - You can reset the TLV properties to their default values by using the `lldpadm reset-tlvprop` command for global TLV units and `lldpadm reset-agenttlvprop` command for per-agent TLV units.

1. **Become an administrator.**

For more information, see [“Using Your Assigned Administrative Rights” in *Securing Users and Processes in Oracle Solaris 11.3*](#).

2. **Configure a global or a per-agent TLV unit.**

- **To configure a global TLV unit, set the appropriate TLV property to contain the values that you want to advertise.**

```
# lldpadm set-tlvprop -p TLV-property=value[,value,value,...] TLV-name
```

where *TLV-name* is the name of the global TLV unit and *TLV-property* is a property of that TLV unit. You can assign multiple values to the property. For a list of global TLV units and their properties, see [Table 3, “Global TLV Units and Their Properties,” on page 100](#).

- **To configure a per-agent TLV unit, configure the appropriate TLV property of the LLDP agent to contain the values that you want the agent to advertise.**

```
# lldpadm set-agenttlvprop -p TLV-property[+|-]=value[,value,value,...] -a agent TLV-name
```

where *TLV-name* is the name of the agent TLV unit and *TLV-property* is a property of that TLV unit. You can assign multiple values to the property. For a list of per-agent TLV units and their properties, see [Table 4, “Per-Agent TLV Units and Their Properties,” on page 100](#).

You can use qualifiers to add (+) or remove (-) values from the list of values for properties that accept multiple values.

3. (Optional) Display the values of the TLV property that you have configured.

■ **To display the global TLV property values:**

```
# lldpadm show-tlvprop
```

■ **To display the values of the TLV property of an agent:**

```
# lldpadm show-agenttlvprop
```

Example 32 Defining TLV Values for the syscapab and mgmtaddr TLV Units

In the following example, specific information about the capabilities of the system to be advertised in the LLDP packet and management IP address is configured.

1. Configure both supported and enabled properties of the syscapab TLV unit.

```
# lldpadm set-tlvprop -p supported=bridge,router,repeater syscapab
# lldpadm set-tlvprop -p enabled=router syscapab
```

2. Specify the management IP address for the mgmtaddr TLV unit.

```
# lldpadm set-tlvprop -p ipaddr=192.168.1.2 mgmtaddr
```

3. Display the TLV values of the agent properties.

```
# lldpadm show-tlvprop
```

TLVNAME	PROPERTY	PERM	VALUE	DEFAULT	POSSIBLE
syscapab	supported	rw	bridge, router, repeater	bridge,router, station	other,router, repeater,bridge, wlan-ap,telephone, docis-cd,station, cvlan,svlan,tpmr
syscapab	enabled	rw	router	none	bridge,router, repeater
mgmtaddr	ipaddr	rw	192.162.1.2	none	--

The output includes the default values of the TLV units and also the possible values that can be set for the property.

For information about configuring per-agent TLV properties, see [Chapter 7, “Managing Converged Networks by Using Data Center Bridging”](#).

Disabling LLDP

This section describes how to disable LLDP selectively on individual ports.

▼ How to Disable LLDP

To disable LLDP across all of the system's interfaces, perform the following steps.

1. **Become an administrator.**

For more information, see [“Using Your Assigned Administrative Rights” in *Securing Users and Processes in Oracle Solaris 11.3*](#).

2. **Change the SMF LLDP property to `no`, which disables automatic enabling of LLDP on all ports except those with existing LLDP configurations.**

```
# svccfg -s svc:/network/lldp:default setprop lldp/auto-enable-agents = "no"
```

3. **Restart the LLDP service.**

```
# svcadm restart svc:/network/lldp:default
```

4. **Disable LLDP on each port whose previous LLDP configuration is retained.**

- **To disable LLDP by changing the mode of the agent:**

```
# lldpadm set-agentprop -p mode=disable agent
```

where *agent* is the LLDP agent and is identified by the physical link on which the agent is enabled. For example, if you enable LLDP on `net0`, the agent is `net0`.

- **To disable LLDP by removing the LLDP configuration from the port:**

```
# lldpadm reset-agentprop -p mode agent
```

In this command, you do not set a value for the mode property.



Caution - If `auto-enable-agents` that is set to `no` is switched back to `yes`, LLDP behaves differently than if the agent's mode on that port were simply disabled.

Monitoring LLDP Agents

The `lldpadm show-agent` command displays complete information that is advertised by an LLDP agent. Relative to a given system, the advertisement can be information about the local system that is transmitted to the rest of the network or information that is received by the system from other systems on the same network.

Displaying the Advertised Information

The information can be either local or remote. *Local* information comes from the local LLDP agent. *Remote* information comes from other LLDP agents on the network that is received by the local LLDP agent.

Use the `lldpadm show-agent` command to display the advertised information.

```
# lldpadm show-agent -[l|r][v] agent
```

- `-l` displays local information advertised by the local LLDP agent.
- `-r` displays remote information received by the LLDP agent.
- `-v` displays detailed local or remote information.

EXAMPLE 33 Displaying Advertised LLDP Agent Information

The following example shows how to display the information that is being advertised locally or remotely by an LLDP agent. By default, the information is displayed in short form. By using the `-v` option, you can obtain verbose or detailed information.

To display local information that is advertised by the LLDP agent:

```
# lldpadm show-agent -l net0
AGENT  CHASSISID  PORTID
net0   004bb87f    00:14:4f:01:77:5d
```

To display remote information that is advertised by the LLDP agent:

```
# lldpadm show-agent -r net0
AGENT  SYSNAME  CHASSISID  PORTID
net0   hostb      0083b390   00:14:4f:01:59:ab
```

To display local information in the verbose mode, use the `-v` option:

```
# lldpadm show-agent -l -v net4
Agent: net4
Chassis ID Subtype: Local(7)
Chassis ID: 00843300
```

```

Port ID Subtype: MacAddress(3)
Port ID: 00:1b:21:89:03:d0
Port Description: --
Time to Live: 21 (seconds)
System Name: --
System Description: --
Supported Capabilities: --
Enabled Capabilities: --
Management Address: --
Maximum Frame Size: --
Port VLAN ID: --
VLAN Name/ID: vlan1/22
VNIC PortID/VLAN ID: 02:08:20:63:2d:9d,02:08:20:e5:6c:af/21
Aggregation Information: --
PFC Willing: On
PFC Cap: 8
PFC MBC: False
PFC Enable: 4
PFC Pending: True
Application(s)(ID/Sel/Pri): --
ETS Willing: On
ETS Configured CBS: 0
ETS Configured TCS: 8
ETS Configured PAT: 0,1,2,3,4,5,6,7
ETS Configured BAT: 40,20,0,40,0,0,0,0
ETS Configured TSA: 2,2,2,2,2,2,2,2
ETS Recommended PAT: 0,1,2,3,4,5,6,7
ETS Recommended BAT: 40,20,0,40,0,0,0,0
ETS Recommended TSA: 2,2,2,2,2,2,2,2
EVB Mode: Station
EVB GID (Station): Not Supported
EVB ReflectiveRelay REQ: Not Requested
EVB ReflectiveRelay Status: RR Not Enabled
EVB GID (Bridge): Not Supported
EVB ReflectiveRelay Capable (RRCAP): Not Supported
EVB ReflectiveRelay Control (RRCTR): Not Enabled
EVB max Retries (R): 0
EVB Retransmission Exponent (RTE): 0
EVB Remote or Local(ROL) and
Resource Wait Delay (RWD): Local
EVB Resource Wait Delay (RWD): 0
EVB Remote or Local (ROL) and
Reinit Keep Alive (RKA): Local
EVB Reinit Keep Alive (RKA): 0
Next Packet Transmission: 4 (seconds)

```

To display remote information in the verbose mode, use the -v option:

```
# lldpadm show-agent -r -v net4
```

```

Agent: net4
Chassis ID Subtype: Local(7)
Chassis ID: 00843300

```

```
Port ID Subtype: MacAddress(3)
Port ID: 00:1b:21:89:03:d0
Port Description: --
Time to Live: 21 (seconds)
System Name: --
System Description: --
Supported Capabilities: --
Enabled Capabilities: --
Management Address: --
Maximum Frame Size: --
Port VLAN ID: --
VLAN Name/ID: vlan1/22
VNIC PortID/VLAN ID: 02:08:20:63:2d:9d,02:08:20:e5:6c:af/21
Aggregation Information: --
PFC Willing: On
PFC Cap: 8
PFC MBC: False
PFC Enable: 4
Application(s)(ID/Sel/Pri): --
ETS Willing: On
ETS Configured CBS: 0
ETS Configured TCS: 8
ETS Configured PAT: 0,1,2,3,4,5,6,7
ETS Configured BAT: 40,20,0,40,0,0,0,0
ETS Configured TSA: 2,2,2,2,2,2,2,2
ETS Recommended PAT: 0,1,2,3,4,5,6,7
ETS Recommended BAT: 40,20,0,40,0,0,0,0
ETS Recommended TSA: 2,2,2,2,2,2,2,2
EVB Mode: Station
EVB GID (Station): Not Supported
EVB ReflectiveRelay REQ: Not Requested
EVB ReflectiveRelay Status: RR Not Enabled
EVB GID (Bridge): Not Supported
EVB ReflectiveRelay Capable (RRCAP): Not Supported
EVB ReflectiveRelay Control (RRCTR): Not Enabled
EVB max Retries (R): 0
EVB Retransmission Exponent (RTE): 0
EVB Remote or Local(ROL) and
Resource Wait Delay (RWD): Local
EVB Resource Wait Delay (RWD): 0
EVB Remote or Local (ROL) and
Reinit Keep Alive (RKA): Local
EVB Reinit Keep Alive (RKA): 0
Information Valid Until: 19 (seconds)
```


Displaying LLDP Statistics

You can display LLDP statistics to obtain information about LLDP packets that are being advertised by the local system or by remote systems. The statistics refer to significant events that involve LLDP packet transmission and reception.

- To display all statistics about LLDP packet transmission and reception:

```
# lldpadm show-agent -s agent
```

- To display selected statistics information, use the `-o` option:

```
# lldpadm show-agent -s -o field[,field,...]agent
```

where *field* refers to any field name in the output of the `show-agent -s` command.

EXAMPLE 34 Displaying LLDP Packet Statistics

This example shows how to display information about LLDP packet advertisement.

```
# lldpadm show-agent -s net0
AGENT IFRAMES IERR IDISCARD OFRAMES OLENERR TLVDISCARD TLVUNRECOG AGEOUT
net0      9      0          0      14          0          4          5          0
```

This output provides the following information:

- AGENT specifies the name of the LLDP agent, which is identical to the datalink on which the LLDP agent is enabled.
- IFRAMES, IERR, and IDISCARD display information about packets being received, incoming packets with errors, and incoming packets that are dropped.
- OFRAMES and OLENERR refer to outgoing packets and packets that have length errors.
- TLVDISCARD and TLVUNRECOG display information about TLV units that are discarded and TLV units that are not recognized.
- AGEOUT refers to packets that have timed out.

The example indicates that out of 9 frames received into the system, 5 TLV units are unrecognized, possibly because of noncompliance with standards. The example also shows that 14 frames were transmitted by the local system to the network.

EXAMPLE 35 Displaying Selected LLDP Packet Statistics

This example shows how to display selected statistics information.

```
# # lldpadm show-agent -s -o iframes,oframes net4
IFRAMES OFRAMES
```

0 10

Managing Converged Networks by Using Data Center Bridging

Traditionally, different networks are used for traffic management based on the application requirements, and load distribution of network traffic that is based on the available bandwidth. For example, local area network (LAN) uses Ethernet and storage area network (SAN) uses fibre channel. However, data center bridging enhances the Ethernet making it more suitable for running different types of traffic, which is converged traffic, and also to support features such as losslessness. DCB enables efficient network infrastructure by consolidating SAN and LAN and thereby reducing operational and management costs in large deployments.

This chapter contains the following topics:

- [“Overview of Data Center Bridging” on page 115](#)
- [“Priority-Based Flow Control” on page 117](#)
- [“Enhanced Transmission Selection” on page 118](#)
- [“Enabling DCBX” on page 119](#)
- [“Setting the Mode of Operation for DCB” on page 120](#)
- [“Customizing Priority-Based Flow Control for DCB” on page 121](#)
- [“Displaying PFC Configuration Information” on page 124](#)
- [“Application Priority Configurations” on page 126](#)
- [“Customizing Enhanced Transmission Selection for DCB” on page 127](#)
- [“Recommending ETS Configuration to the Peer” on page 130](#)
- [“Displaying ETS Configuration Information” on page 132](#)

Overview of Data Center Bridging

Data center bridging is used to manage the bandwidth, relative priority, and flow control of multiple traffic types when sharing the same network link, for example, when sharing a datalink between networking and storage protocols. Fibre channel can be dedicated to host this type of traffic. However, using dedicated links to service only fibre channel traffic can be costly.

Therefore, fibre channel over Ethernet (FCoE) is more commonly used. DCB addresses the sensitivity of fibre channels to packet loss while traversing an Ethernet network.

DCB distinguishes traffic based on priorities, which are also called class of service (CoS) priorities. The host and the next hop use the DCB exchange protocol (DCBX) to negotiate a network configuration, such as no traffic loss and minimum bandwidth share, based on priorities. This process enables packets from different applications on the host and in the network to be treated according to their priorities and the corresponding configuration to be negotiated by using DCBX.

Oracle Solaris supports the IEEE 802.1qaz DCBX specification and also the pre-standard Converged Enhanced Ethernet (CEE) DCBX specification v1.01 to enable interoperability with a diverse set of switches when using DCB. For more information about how to choose different versions for negotiations, see [“Setting the Mode of Operation for DCB” on page 120](#).

Each packet in a DCB network has a VLAN header that contains a DCB 3-bit priority value, which is a DCB priority. This IEEE 802.1p priority value differentiates each Ethernet packet in the network from the other packets. Depending on the priority values of the packets, you can configure DCB to allocate specific bandwidth to the packets. For example, all packets with a priority 1 must have PFC enabled and all packets with a priority 2 must have PFC disabled and a bandwidth share of 10%.

You can configure DCB features such as priority-based flow control (PFC) and enhanced transmission selection (ETS) based on priorities. For more information about PFC and ETS, see [“Priority-Based Flow Control” on page 117](#) and [“Enhanced Transmission Selection” on page 118](#).

The DCB cos datalink property enables you to specify the CoS or priority of the datalink. The cos value that is set on a primary datalink does not apply to the VNICs that are created over this physical link. For information about customizing PFC based on the cos property, see [“Customizing Priority-Based Flow Control for DCB” on page 121](#). For information about customizing ETS based on the cos property, see [“Customizing Enhanced Transmission Selection for DCB” on page 127](#).

In Oracle Solaris, LLDP is used to exchange DCBX type-length-value (TLV) units. For more information about LLDP, see [Chapter 6, “Exchanging Network Connectivity Information With Link Layer Discovery Protocol”](#). Provided that the underlying network interface card (NIC) supports DCB features such as priority-based flow control and enhanced transmission selection, configuration information for these features can be shared with peers on the network, as follows:

- PFC prevents packet loss by implementing a mechanism that pauses traffic flow for packets with a defined class of service (CoS). For more information about CoS, see the description of the cos link property in the [dladm\(1M\)](#) man page.
- ETS enables bandwidth sharing among packets based on the defined CoS. See [“Enhanced Transmission Selection” on page 118](#).

Considerations When Using DCB

Note the following considerations for using DCB:

- DCB is *only* supported on Intel Niantic physical NICs.
To verify whether the NIC supports DCB, issue the following command:

```
# dladm show-linkprop -p ntcs agent
```

A property value that is greater than zero (0) indicates that the NIC supports DCB.

- DCB ports that are configured in DCB mode cannot be aggregated (trunk or DLMP mode).
- Because DCB supports only the IEEE and CEE versions of DCBX (not the CIN version), external bridges must support the IEEE or CEE version to interoperate with Oracle Solaris DCB.
- DCB supports ETS configuration and recommendation TLVs.
- DCB is only supported in the eight traffic class configuration.
- DCB does not support congestion notification (CN).

Priority-Based Flow Control

Priority-based flow control (PFC) extends the standard IEEE 802.3x PAUSE frame to include IEEE 802.1p CoS values. With PFC, instead of halting all traffic on the link when a PAUSE frame is sent, traffic is paused only for those CoS values that are enabled in the PFC frame. A PFC frame is sent for the enabled cos property value for which traffic needs to be paused. The sending host stops traffic for that cos property value while traffic for other disabled cos property values are unaffected. After a time interval that is specified in the PFC frame, transmission resumes for the paused packets.

Pausing based on CoS values ensures that packets are not dropped for that cos property value. For packets without any defined CoS value or with CoS values that do not have PFC enabled, no PAUSE frames are sent. Therefore, traffic continues to flow, and packets might be dropped during traffic congestion. Handling loss of packets depends on the protocol stack, for example, TCP.

Two types of DCB information exist on the host: local DCB information and remote DCB information. For the PFC features to be effective, the local and remote types of DCB information for PFC on the host must be symmetric. The local host must be able to match the DCB information that it receives from the peer. If you enable DCB on your system, DCB synchronizes the DCB information with the peer.

In most cases, the default configuration for PFC is sufficient. This configuration is automatically set up when you enable LLDP. However, you can adjust different options when


configuring PFC. For more information, see [“Customizing Priority-Based Flow Control for DCB” on page 121](#) and [“Displaying PFC Configuration Information” on page 124](#).

Enhanced Transmission Selection

ETS is a DCB feature that enables you to allocate bandwidth on a NIC to applications based on their DCB priority.

The following figure shows the ETS feature of DCB in a network.

FIGURE 14 Enhanced Transmission Selection in DCB



The host in the figure has different types of traffic, such as FCoE and iSCSI, that share the link bandwidth. In the figure, the priority and bandwidth as shown in the following table are assigned for different types of traffic.

Traffic	Priority	Bandwidth
FCoE	3	60%
iSCSI	4	30%
Ethernet (non iSCSI) traffic	0	10%

The DCBX ETS TLV with the corresponding ETS bandwidth allocation is as follows:

Priority	0	1	2	3	4	5	6	7
----------	---	---	---	---	---	---	---	---

Bandwidth Allocation Percentage	10	0	0	60	30	0	0	0
--	----	---	---	----	----	---	---	---

To use the ETS feature, the NIC must support the feature and run in the DCB mode. When you enable LLDP, the default configuration for the ETS feature is automatically set up if the underlying link supports DCB. However, you can modify the default configuration. For more information, see [“Customizing Enhanced Transmission Selection for DCB” on page 127](#), [“Recommending ETS Configuration to the Peer” on page 130](#), and [“Displaying ETS Configuration Information” on page 132](#).

Enabling DCBX

Support for DCBX is automatically enabled when you enable LLDP. This procedure provides alternative manual steps in case certain automatic processes fail.

▼ How to Enable the Data Center Bridging Exchange Feature Manually

Before You Begin Ensure that you install LLDP. For more information about enabling LLDP, see [“Enabling LLDP on the System” on page 101](#).

1. Become an administrator.

For more information, see [“Using Your Assigned Administrative Rights” in *Securing Users and Processes in Oracle Solaris 11.3*](#).

2. Verify that the LLDP service is running.

```
# svcs lldp
```

If the LLDP service is disabled, start the service with the following command:

```
# svcadm enable svc:/network/lldp:default
```

3. Ensure that the LLDP agent is running on Rx and Tx modes.

```
# lldpadm show-agentprop -p mode agent
```

If the LLDP agent is not enabled on both modes, type the following command:

```
# lldpadm set-agentprop -p mode=both agent
```

For the other possible configurations of the LLDP agents, see [“Enabling LLDP on the System” on page 101](#).

4. Verify that the underlying NIC supports DCB.

```
# dladm show-linkprop -p ntcs agent
```

A property value that is greater than zero (0) indicates that the NIC supports DCB.

Setting the Mode of Operation for DCB

Oracle Solaris hosts can use either the IEEE DCBX or CEE DCBX protocol to exchange information about DCB features with their peers that are directly connected, for example, the first hop switch. The exchange enables negotiating and configuring DCB features both on the host and the peer.

The following table shows the DCBX applications that are supported for the IEEE DCBX and CEE DCBX versions.

TABLE 5 Applications Supported for IEEE DCBX and CEE DCBX Versions

Applications	IEEE DCBX	CEE DCBX
PFC	Yes	Yes
Application TLVs	Yes	Yes
ETS configuration	Yes	No
ETS recommendation	Yes	No

You can select `ieee`, `cee`, or `auto` as the mode of operation for DCB depending on the standard that the switch supports. The default mode is `auto`, which by default operates in IEEE mode and switches to CEE mode when the host receives a CEE packet from the peer. To select the mode, you need to set the property `dcbx-version` by using the `lldpadm` command.

If you explicitly set the mode, the mode does not transition to CEE or IEEE when the CEE or IEEE packets are received.

The transition from IEEE to CEE happens only once. Subsequently, when the peer changes its DCBX mode, Oracle Solaris will not switch mode automatically.

Some switches initiate exchange of information regardless of the peer. However, some switches might listen to the peer and respond only if it supports the DCBX version. For example, if you configure the DCBX mode of an Oracle Solaris host as `auto`, then by default it sends IEEE DCBX packets. If this host is connected to a switch that does not support IEEE DCBX, then

the switch might not respond even if it supports CEE DCBX version. In such cases, you must explicitly configure the DCBX mode as *cee*.

▼ How to Set the Mode of Operation for DCB

1. Become an administrator.
2. (Optional) Display the current DCBX mode.

```
# lldpadm show-agentprop -p dcbx-version net0
```

AGENT	PROPERTY	PERM	VALUE	EFFECTIVE	DEFAULT	POSSIBLE
net0	dcbx-version	rw	auto	ieee	auto	auto,ieee,cee

3. Set the mode of the operation.

```
# lldpadm set-agentprop -p dcbx-version=DCBX-mode net0
```

where *DCBX-mode* can be set to one of the following values:

- *auto* – The default DCBX mode. When the mode is *auto*, DCBX operates in IEEE mode and switches to CEE mode when the host receives a CEE packet from the peer.
- *ieee* – Uses only the IEEE protocol to exchange information.
- *cee* – Uses only the CEE protocol to exchange information.

4. (Optional) Display the current mode.

```
# lldpadm show-agentprop -p dcbx-version net0
```

AGENT	PROPERTY	PERM	VALUE	EFFECTIVE	DEFAULT	POSSIBLE
net0	dcbx-version	rw	auto	cee	auto	auto,ieee,cee

Tip - To reset the DCBX mode to the default setting, use the following command:

```
# lldpadm reset-agentprop -p dcbx-version net0
```

Customizing Priority-Based Flow Control for DCB

PFC and ETS are enabled by default only if the underlying NIC is in DCB mode. If you prefer to use only PFC, then you must remove *etscfg* from the *dot1 -tlv* property of the LLDP agent by using the following command:

```
# lldpadm set-agentprop -p dot1-tlv==etscfg net0
```

For a list of possible values for `dot1 -tlv`, refer to [Table 2, “Optional TLV Units for an LLDP Agent,” on page 99](#).

Setting the PFC-Related Datalink Properties

The PFC feature of DCB provides the following datalink properties:

- `pfcmap` – Provides information about priority definitions and mappings. The `pfcmap` property refers to an 8-bit mask (0–7) that represent priorities. The lowest bit represents priority 0 while the highest bit represents priority 7. Each bit in this mask signifies whether PFC is enabled for a corresponding priority. By default, `pfcmap` is set to 1111111, which means that PFC is enabled on all the priorities.
- `pfcmap-rmt` – Specifies the operative PFC mapping on the remote peer. This property is read-only.

In a DCB network, when a receiver is unable to keep up with the incoming rate of traffic, it sends a PFC frame to the sender requesting the sender to pause traffic for priorities that have PFC enabled. For any packet transmitted over a link, DCB sends a PFC frame to the sending host if traffic congestion accumulates on the receiving host. To send PFC frames properly, the communicating hosts must have symmetric DCB configuration information. A system can automatically adjust its PFC configurations to match the PFC configurations on the remote peer. You can determine the operative PFC mapping on the local host by using the `dladm show-linkprop` command that displays the `EFFECTIVE` value for the `pfcmap` property. For more information, see [“Displaying Datalink Properties” on page 124](#).

▼ How to Customize Priority-Based Flow Control for DCB

1. **Become an administrator.**

For more information, see [“Using Your Assigned Administrative Rights” in *Securing Users and Processes in Oracle Solaris 11.3*](#).

2. **Ensure that the `flowctrl` property of the datalink is set to `pfc`.**

```
# dladm show-linkprop -p flowctrl datalink
```

If the property is not set to `pfc` or `auto`, use the following command:

```
# dladm set-linkprop -p flowctrl=pfc datalink
```

3. **Set the `pfcmap` property to a value other than the default value 11111111.**

```
# dladm set-linkprop -p pfcmap=value datalink
```

For example, to enable priority only on CoS priority 6, type the following command:

```
# dladm set-linkprop -p pfcmap=01000000 net0
```

Setting the PFC TLV Units

DCB uses PFC TLV units to exchange PFC information between hosts. Either hosts must have the same value for the `pfcmap` property or at least one host must be willing to accept the configuration of its peer. You can set the `pfcmap` property by using the `dladm set-linkprop` command.

The TLV property `willing` indicates whether the host is ready to accept the peer's configuration if the host configuration differs from the peer's configuration. By default, the property value of `willing` is set to `on`, which indicates that the host will accept the peer's configuration.

To verify that the host can synchronize its PFC information with the PFC information on the remote peer, you must determine whether the `willing` is set to `on` by using the following command:

```
# lldpadm show-agenttlvprop -p willing -a agent pfc
```

If the PFC TLV property `willing` is set to `off`, type the following command to set the property `willing` to `on` and enable synchronization.

```
# lldpadm set-agenttlvprop -p willing=on -a agent pfc
```

where `agent` is the datalink on which the agent is enabled.

EXAMPLE 36 Enabling Synchronization Between the Host and the Peer

To enable synchronization for a `net0` datalink, type the following command:

```
# lldpadm set-agenttlvprop -p willing=on -a net0 pfc
```

```
# dladm show-linkprop -p pfcmap,pfcmap-rmt net0
```

LINK	PROPERTY	PERM	VALUE	EFFECTIVE	DEFAULT	POSSIBLE
net0	pfcmap	rw	11111111	00010000	11111111	00000000-11111111
net0	pfcmap-rmt	r-	--	00010000	11111111	--

In the example, the `pfcmap` and `pfcmap-rmt` properties have the value `00010000`. This indicates that the local host has synchronized with the peer. Hence, PFC is enabled with a priority of 4 on both the host and the peer.

For more information, see the [lldpadm\(1M\)](#) man page.

Displaying PFC Configuration Information

This section describes commands to display information related to PFC after LLDP and DCB are configured and provides examples to show the use of these commands.

Displaying Datalink Properties

The following command displays the priority definitions and the effective PFC mappings on the datalink:

```
# dladm show-linkprop -p pfcmap,pfcmap-rmt datalink
```

On a datalink with matching PFC information between the local and remote peers, the values of the EFFECTIVE column for pfcmap and pfcmap-rmt properties are identical regardless of the value set for the pfcmap property. If the ability to synchronize is disabled on the local host, then the EFFECTIVE field for the pfcmap property reflects the value of the pfcmap property for the local host.

EXAMPLE 37 Displaying PFC-Related Datalink Properties

This example shows how to display the status of physical datalink properties that are related to priority-based flow control.

```
# dladm show-linkprop -p pfcmap,pfcmap-rmt net0
LINK  PROPERTY      PERM  VALUE      EFFECTIVE  DEFAULT  POSSIBLE
net0   pfcmap          rw    11111111   11111111   11111111  00000000-11111111
net0   pfcmap-rmt      r-    --         --         --        --
```

In the example, the value field for the pfcmap property has the value of 11111111. This value indicates that the PFC mapping on the local host has the default value where all eight priorities are enabled. The EFFECTIVE values for the pfcmap and pfcmap-rmt properties are 11111111 and --. These mismatched values for the EFFECTIVE field indicate that the local host has not synchronized its PFC information with the remote peer.

You can use the `lldpadm show-agenttlvprop` command to verify the value of the willing property and the `lldpadm show-agent -r` command to check the PFC TLV information from the peer.

Displaying the Capability of the Local Host to Synchronize PFC Information

The following command displays the PFC TLV property that controls a host's capability to synchronize its PFC mapping with a peer.

```
# lldpadm show-agenttlvprop -a agent pfc
```

where *agent* is identified by the datalink on which LLDP is enabled.

EXAMPLE 38 Displaying the Capability of the Local Host to Synchronize PFC Information

This example shows how to display the current status of the host's ability to adjust to PFC configurations of the peer.

```
# lldpadm show-agenttlvprop -a net0 pfc -p willing
AGENT  TLVNAME  PROPERTY  PERM  VALUE  EFFECTIVE  DEFAULT  POSSIBLE
net0    pfc        willing   rw    off    --         on       on,off
```

For more information, see [“Setting the PFC TLV Units” on page 123](#).

Displaying PFC Mapping Information Between Host and Peer

The PFC Pending value returns a True status if PFC information between the host and the peer does not converge. After the mismatch is resolved, the status of PFC Pending returns to False.

The following command alerts you to a mismatch of PFC mapping information between the local host and the peer.

```
# lldpadm show-agent -lv -o "PFC Pending" agent
```

```
# lldpadm show-agent -lv -o "PFC Pending" agent
```

EXAMPLE 39 Verifying Symmetry of PFC Information Between Host and Peer

The following example shows how to verify in actual running time whether PFC information is synchronized between the host and peer, or whether a mismatch occurs.

```
# lldpadm show-agent -lv -o "PFC Pending" net0
PFC Pending: True
```

To display all the information that the agent advertises, use the `-v` (verbose) option of the `lldpadm show-agent` command:

```
# lldpadm show-agent -v net0
```

Displaying Priority Definitions

The following command displays PFC information on the physical link with regards to enabled priorities on the NIC:

```
# dladm show-phys -D pfc datalink
```

EXAMPLE 40 Displaying CoS Priority Definitions

This example shows how to display the current priority definitions on a specific physical link based on the value of the `pfcmap` property. For example, assume that `pfcmap` is configured as `01000000`. To display the corresponding priority mappings on the physical link, you would proceed as follows:

```
# dladm show-phys -D pfc net0
LINK COS   PFC   PFC_EFFECT CLIENTS
net0  0      YES   NO      net0,vnic1
      1      YES   YES     vnic2
      2      YES   NO     vnic3
      3      YES   NO     vnic4
      4      YES   NO     vnic5
      5      YES   NO     vnic6
      6      YES   NO     vnic7
      7      YES   NO     vnic8
```

For the physical link `net0`, priority is enabled for all VNIC clients configured over the datalink. However, the local host adjusts its PFC mapping to the PFC mapping on the peer, as shown by the values of the `PFC_EFFECT` field, where priority is disabled on `COS 0` and `2-7`. As a result, no PFC frames would be exchanged for traffic on any VNIC except `vnic2` regardless of the availability of resources. With this configuration, packet drops are allowed on traffic that flows on all VNICs except `vnic2`. For traffic on `vnic2`, PFC PAUSE frames are sent when traffic congestion occurs to prevent packet loss on this client.

Application Priority Configurations

DCBX exchanges the priority information that is associated with the application. Application TLV units that are exchanged through DCBX contain information about the priority to be

used for an application on the host. The priority is defined in the Application Priority Table. Each entry in the table contains the name of the application and the priority assigned to the application. When you set a priority for an application, all the DCB settings of the PFC and ETS of that priority are applicable to the application. The application TLV uses the table data for exchanging application priority information with the other hosts.

By default, the application feature accepts priority mappings from the peer. The willing property of the application TLV `appln` is similar to PFC and enables information exchange between the peers.

Entries on the table uses the following format:

protocol-ID/selector/priority

The pair *protocol-ID/selector* identifies the application. The priority for a corresponding application is identified by a priority that contains a value from 0 to 7.

To exchange this information about the priority of an application with other hosts, you set an application TLV as follows:

```
# lldpadm set-agenttlvprop -p property=value -a agent appln
```

For example, for FCoE traffic, the protocol ID is 0x8906 and the selector ID is 1. Suppose that the priority 4 is assigned to this application. Based on [Table 4, “Per-Agent TLV Units and Their Properties,” on page 100](#) that lists the parameters for setting an application TLV you would type the following command:

```
# lldpadm set-agenttlvprop -p apt=8906/1/4 -a net0 appln
# lldpadm show-agenttlvprop -a net0 appln -p apt
```

AGENT	TLVNAME	PROPERTY	PERM	VALUE	EFFECTIVE	DEFAULT	POSSIBLE
net0	appln	apt	rw	8906/1/4	8906/1/4	--	--

Customizing Enhanced Transmission Selection for DCB

The default configuration is automatically set up when LLDP is enabled and DCB is supported by the underlying link. In this default configuration, the cos value 0 is assigned all of the bandwidth. However, you can use the `dladm set-linkprop` command to configure the cos values on a datalink to assign part of the bandwidth to that datalink.

ETS configuration and recommendation TLVs are enabled by default for a NIC. For a list of possible values for `dot1-tlv`, refer to [Table 2, “Optional TLV Units for an LLDP Agent,” on page 99](#).

If you want to remove the `pfc` TLV, type the following command:

```
# lldpadm set-agenttlvprop -p dot1-tlv==pfc agent
```

Setting the ETS-Related Datalink Properties

The properties of datalinks that refer to PFC information apply to the prevention of packet loss based on the priorities defined for the packets. The ETS properties relate to assigning shares of the underlying link's bandwidth based on priorities.

DCB provides the following ETS-related properties:

- **cos** – Specifies the class of service or priority of the datalink. The value of this property ranges from 0 to 7. The default value is 0. The cos value is set in the VLAN tag of the packets that are transmitted over this link.
- **etsbw-lcl** – Indicates the ETS bandwidth that is allocated on the transmit (Tx) side for the datalink. This property is configurable only if the underlying physical NIC has DCB capabilities and supports ETS and the link's cos property is not set to 0. You set a value for this property on a datalink by specifying the percentage of total bandwidth of the underlying physical link. The sum of the values for the etsbw-lcl property for all the datalinks over the same physical NIC must not exceed 100%.

The bandwidth percentage that is defined on etsbw-lcl is not reserved only for that datalink. If the allocated bandwidth is not used, then it can be used by other datalinks on that physical NIC. Further, the bandwidth allocation is enforced only on the transmission side of the host's traffic.

- **etsbw-rmt-advice** – Specifies the recommended ETS bandwidth value sent to the peer. By default, the locally configured value of the etsbw-lcl property is recommended to the peer. However, you can recommend a value that is different from the etsbw-lcl property by explicitly configuring the etsbw-rmt-advice datalink property.

Configuring the etsbw-rmt-advice property is useful if the bandwidth assignment for a datalink is asymmetrical, which means that the receive (Rx) and transmit (Tx) bandwidth are different. When you explicitly set the etsbw-rmt-advice property, transmission of the ETS recommendation DCBX TLV starts automatically.

- **etsbw-lcl-advice** – Specifies the recommended bandwidth share for the datalink, which is sent by the peer to the local host. This is a read-only property.
- **etsbw-rmt** – Specifies the bandwidth share that is configured on the peer for the datalink. This is a read-only property.

To set the priority and to allocate a bandwidth to the VNIC, use the following commands:

- To set the priority to the VNIC:

```
# dladm set-linkprop -p cos=value VNIC
```

- To allocate a percentage of the bandwidth of the underlying physical link to a VNIC:


```
# dladm set-linkprop -p etsbw-lcl=value VNIC
```

The value that you assign to the `etsbw-lcl` property represents a percentage of the total bandwidth capacity of the underlying link. The sum of all the allocated bandwidth values that you assign to the clients must not exceed 100 percent.

- To explicitly recommend a bandwidth that is sent to the peer:

```
# dladm set-linkprop -p etsbw-rmt-advice=value VNIC
```

You can determine the actual bandwidth share that is implemented on the local host's datalink and the bandwidth share that is configured on the peer's datalink by using the `dladm show-linkprop` command. The value in the `EFFECTIVE` field of the output for the `etsbw-lcl` and `etsbw-rmt` properties shows the actual bandwidth share implemented. For more information, see [“Displaying ETS Configuration Information” on page 132](#).

For the appropriate bandwidth to be used for packets with specific priorities, symmetric or synchronized ETS information between the communicating hosts is preferable. Specifically, should be the local system that is able to adjust its bandwidth share to the value of `etsbw-lcl-advice`. An Oracle Solaris system can automatically adjust its ETS configurations to match the ETS recommendation from the peer.

Setting ETS TLV Units

The ETS TLV (`etscfg`) configuration determines how the host responds to ETS recommendations from the peer. This TLV unit has only one configurable property, `willing`. By default, this property is set to `on` and enables the local host to synchronize its ETS configuration with the ETS recommendation of the remote peer.

To verify that the host can synchronize its ETS information with the ETS information of the remote peer, use the following command:

```
# lldpadm show-agenttlvprop -p willing -a agent etscfg
```

If the `willing` property is set to `off`, type the following command to establish synchronization:

```
# lldpadm set-agenttlvprop -p willing=on -a agent etscfg
```

To prevent synchronization of information for a specific agent, set the `willing` property to `off` as follows:

```
# lldpadm set-agenttlvprop -p willing=off -a agent etscfg
```

where *agent* is the datalink on which the agent is enabled.

Recommending ETS Configuration to the Peer

The configured ETS bandwidth values (`etsbw-lcl`) for each priority can be recommended to the peer so that the peer can configure the same values. You must enable the `etsreco` property in the `dot1-tlv` type of the LLDP agent on the NIC to recommend the ETS bandwidth values. The recommended values can be the same as the locally configured ETS values or you can also explicitly configure the recommended values by setting the new datalink property `etsbw-rmt-advice` by using the `dladm set-linkprop` command. Configuring the `etsbw-rmt-advice` property is useful if the assigned bandwidth for a datalink is asymmetrical, which means that the receive (Rx) and transmit (Tx) bandwidth are different.

By default, the configured values of the `etsbw-lcl` property are used to recommend values to the peer. However, you can recommend a different ETS value by setting a different value for the `etsbw-rmt-advice` property. For example, if the network traffic is more on the Tx, then you can configure a higher ETS value for the `etsbw-lcl` property (Tx on the host) and a lower value configured for the `etsbw-rmt-advice` property (Rx to the host).

EXAMPLE 41 Recommending an ETS Configuration to the Peer

1. Ensure that the `etsreco` property is enabled by displaying the `dot1-tlv` type property of the LLDP agent for `net5`.

```
# lldpadm show-agentprop -p dot1-tlv net5
```

AGENT	PROPERTY	PERM	VALUE	EFFECTIVE	DEFAULT	POSSIBLE
net5	dot1-tlv	rw	etsreco,etscfg	etsreco, etscfg	none	none,vlanname,pvid, linkaggr,pfc,appln, evb,etscfg,etsreco, all

2. Allocate a share of 20% of the underlying link's bandwidth for `vnic1`.

```
# dladm set-linkprop -p etsbw-lcl=20 vnic1
```

```
# dladm show-linkprop -p etsbw-lcl vnic1
```

LINK	PROPERTY	PERM	VALUE	EFFECTIVE	DEFAULT	POSSIBLE
vnic1	etsbw-lcl	rw	20	20	0	--

By default, the same value is recommended for the peer.

```
# dladm show-linkprop -p etsbw-rmt-advice vnic1
```

LINK	PROPERTY	PERM	VALUE	EFFECTIVE	DEFAULT	POSSIBLE
vnic1	etsbw-rmt-advice	rw	--	20	0	--

3. Display the information exchanged by LLDP.

```
# lldpadm show-agent -l -v net5
```

4. Display the bandwidth recommended to the peer.

```
# dladm show-phys -D ets -r net5
```

LINK	COS	ETSBW_RMT_EFFECT	ETSBW_RMT_ADVICE	CLIENTS
--	0	0	80	net5
	1	0	0	--
	2	0	0	--
	3	0	20	vnic1
	4	0	0	--
	5	0	0	--
	6	0	0	--
	7	0	0	--

By default, the values configured for the `etsbw-lcl` property for each priority on `net5` are sent as the recommended values to the peer. The `ETSBW_RMT_ADVICE` shows the values recommended to the peer. The output also shows that the peer has not configured any ETS bandwidth on its end. You can also display bandwidth recommended to the peer by using the `lldpadm show-agent` command.

5. Recommend a different value to the peer.

```
# dladm set-linkprop -p etsbw-rmt-advice=10 vnic1
# dladm show-linkprop -p etsbw-rmt-advice vnic1
```

LINK	PROPERTY	PERM	VALUE	EFFECTIVE	DEFAULT	POSSIBLE
vnic1	etsbw-rmt-advice	rw	10	10	0	--

6. Display the bandwidth recommended to the peer.

```
# dladm show-phys -D ets -r net5
```

LINK	COS	ETSBW_RMT_EFFECT	ETSBW_RMT_ADVICE	CLIENTS
--	0	0	90	net5
	1	0	0	--
	2	0	0	--
	3	0	10	vnic2
	4	0	0	--
	5	0	0	--
	6	0	0	--
	7	0	0	--

The `ETSBW_RMT_EFFECT` field shows value `0` for `vnic2`, which indicates that the peer has not set any bandwidth on its end, even though you have recommended bandwidth values. This situation means that peer might not have enabled LLDP or does not support ETS.

Displaying ETS Configuration Information

You can use the following commands to display information about the ETS configuration:

- `# dladm show-linkprop -p etsbw-lcl,etsbw-rmt,etsbw-lcl-advice,etsbw-rmt-advice datalink`

This command displays the information related to ETS on a physical link.

- `# dladm show-phys -D ets phys-link`

This command displays the local and remote ETS configuration on the physical link with regard to bandwidth allocation and distribution across the link.

- `# lldpadm show-agenttlvprop -a agent etscfg`

where *agent* is the datalink on which LLDP is enabled. This command displays the ETS TLV property that controls the capability of a host to synchronize ETS information with a peer.

EXAMPLE 42 Displaying ETS-Related Datalink Properties

This example shows how to display the status of datalink properties that are related to ETS before synchronization is enabled.

```
# dladm show-linkprop -p cos,etsbw-lcl,etsbw-rmt,etsbw-lcl-advice, \
etsbw-rmt-advice vnic1
```

LINK	PROPERTY	PERM	VALUE	EFFECTIVE	DEFAULT	POSSIBLE
vnic1	cos	rw	2	2	0	0-7
vnic1	etsbw-lcl	rw	10	10	0	--
vnic1	etsbw-rmt	r-	20	20	--	--
vnic1	etsbw-lcl-advice	r-	20	20	--	--
vnic1	etsbw-rmt-advice	rw	10	10	0	--

The output shows that the host has set and recommended an ETS value of 10% for vnic1 with a cos value of 2. However, the peer has set and recommended an ETS value of 20% with a cos value of 2 for vnic1. Because the synchronization is not enabled (*willing* is not enabled) the host has not accepted the peer's recommendation, which is reflected in the EFFECTIVE value of the *etsbw-lcl* property (locally configured value).

EXAMPLE 43 Displaying the Capability of the Local Host to Synchronize ETS Information

This example shows how to display the current status of the local host's ability to adjust to the ETS configurations of the peer.

```
# lldpadm show-agenttlvprop -a net0 etscfg
```

AGENT	TLVNAME	PROPERTY	PERM	VALUE	EFFECTIVE	DEFAULT	POSSIBLE
-------	---------	----------	------	-------	-----------	---------	----------

```
net0    etscfg    willing    rw    off    --    on    on,off
```

To enable synchronization, type the following commands:

```
# lldpadm set-agenttlvprop -p willing=on -a net0 etscfg
```

```
# dladm show-linkprop -p cos,etsbw-lcl,etsbw-rmt, \
etsbw-lcl-advice,etsbw-rmt-advice vnic1
```

LINK	PROPERTY	PERM	VALUE	EFFECTIVE	DEFAULT	POSSIBLE
vnic1	cos	rw	2	2	0	0-7
vnic1	etsbw-lcl	rw	10	20	0	--
vnic1	etsbw-rmt	r-	20	20	--	--
vnic1	etsbw-lcl-advice	r-	20	20	--	--
vnic1	etsbw-rmt-advice	rw	10	10	0	--

Because the synchronization is enabled (the property `willing` is enabled), the host has accepted the peer's recommendation, which is reflected in the `EFFECTIVE` value of `etsbw-lcl`.

The following example shows effective ETS values on the host and the peer for each priority value on the physical link.

```
# dladm show-phys -D ets net4
```

LINK	COS	ETSBW_LCL_EFFECT	ETSBW_RMT_EFFECT	ETSBW_LCL_SOURCE	CLIENTS
net4	0	0	30	local	net4
	1	0	0	local	--
	2	0	0	local	--
	3	0	0	local	--
	4	0	70	local	--
	5	0	0	local	--
	6	0	0	local	--
	7	0	0	local	--

`ETSBW_LCL_EFFECT` Displays the effective ETS bandwidth as a percentage for the priority.

`ETSBW_RMT_EFFECT` Displays the effective ETS bandwidth as a percentage for the priority value on the peer.

`ETSBW_LCL_SOURCE` Indicates the source for the `ETSBW_LCL_EFFECT` value. This value could be either `local`, which is the configured value, or `remote`, which is the recommended value.

The following example shows the local ETS information including the locally configured values, local effective values, and the values recommended by the peer.

```
# dladm show-phys -D ets -l net5
```

LINK	COS	ETSBW_LCL	ETSBW_LCL_EFFECT	ETSBW_LCL_ADVICE	CLIENTS
--	0	80	80	0	net5
	1	0	0	0	--
	2	0	0	0	--
	3	20	20	0	vnic2

4	0	0	0	--
5	0	0	0	--
6	0	0	0	--
7	0	0	0	--

Because the peer has not recommended any values, the local effective value (ETSBW_LCL_EFFECT) is set by using the local configured value (ETSBW_LCL).

The following example displays information about the peer.

```
# dladm show-phys -D ets -r net5
```

LINK	COS	ETSBW_RMT_EFFECT	ETSBW_RMT_ADVICE	CLIENTS
--	0	0	20	net5
	1	0	0	--
	2	0	0	--
	3	0	80	vnic2
	4	0	0	--
	5	0	0	--
	6	0	0	--
	7	0	0	--

The output shows that the remote peer does not have any value set in the ETSBW_RMT_EFFECT field even though the host had recommended the peer to set 80% for priority 3.

Link Aggregations and IPMP: Feature Comparison

Link aggregation and IPMP are different technologies that achieve improved network performance and maintain network availability.

The following table presents a general comparison between link aggregation and IPMP.

Feature	Link Aggregation	IPMP
Network technology type	Layer 2 (link layer).	Layer 3 (IP layer).
Configuration tool	dladm	ipadm
Link-based failure detection	Supported.	Supported.
Probe-based failure detection	Trunk: Based on LACP, targeting the immediate peer host or switch. DLMP: Supported. ICMP-based, targeting any defined systems in the same subnet as DLMP addresses, across multiple levels of intervening Layer 2 switches.	ICMP-based, targeting any defined system in the same IP subnet as test addresses across multiple levels of intervening Layer 2 switches.
Use of standby interfaces	Trunk: Not supported. DLMP: Not supported.	Supported. Standby interfaces can be configured.
Span multiple switches	Trunk: Supported. However, requires switch vendor extensions. DLMP: Supported.	Supported.
Switch configuration	Trunk: Required. DLMP: Not required.	Not required.
Back-to-back configuration	Trunk: Supported. DLMP: Not supported.	Not supported.
Media types supported	Ethernet-specific.	Broadcast-capable.
Load-spreading support	Trunk: Supported and controlled by the administrator by using the dladm command. Inbound load spreading is supported.	Supported. Controlled by the kernel. Inbound load spreading is indirectly affected by the source address selection.

Feature	Link Aggregation	IPMP
	DLMP: Supported across clients and VNICs of the aggregation. However, load spreading by individual clients and VNICs over the aggregation is not supported.	
Level of support when integrating with VNICs	Excellent support. Aggregation is configured in the control domain or the global zone only and is transparent to the zones.	Supported. However, VNIC properties such as bandwidth limit, dedicated Rx or Tx rings, and link protection cannot be enforced on an IPMP group. Requires multiple VNICs to be assigned to the zones and needs to be configured in every zone.
User defined flows for resource management	Supported.	Not supported.
Link protection	Supported.	Not supported.
Protocol requirements	None.	None.

In link aggregations, incoming traffic is spread over the multiple links that comprise the aggregation in trunk mode. Therefore, networking performance is enhanced as more NICs are installed to add links to the aggregation.

DLMP aggregations span multiple switches. As a Layer 2 technology, aggregations integrate well with other Oracle Solaris virtualization technologies.

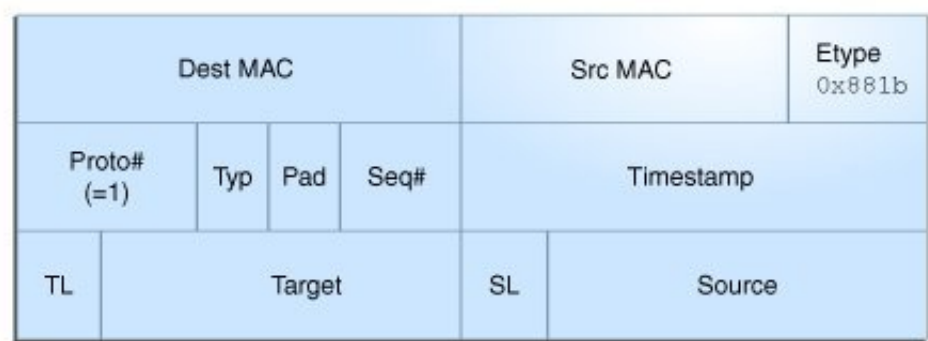
IPMP's traffic uses the IPMP interface's data addresses as they are bound to the available active interfaces. If, for example, all the data traffic is flowing between only two IP addresses but not necessarily over the same connection, then adding more NICs will not improve performance with IPMP because only two IP addresses remain usable.

♦ ♦ ♦ B APPENDIX B

Packet Format of Transitive Probes

The transitive probe packet is a proprietary protocol packet with an Ethernet type, ETHERTYPE_ORCL (0x881b). For more information about the transitive probes, see [“Probe-Based Failure Detection” on page 26](#). For an example to display the statistics of the probes, see [Example 7, “Displaying Probe-Related Information,” on page 39](#). The following figure shows the transitive probe packet format.

FIGURE 15 Transitive Probe Packet



Field	Description
Dest MAC	Destination MAC address.
Src MAC	Source MAC address.
Proto#	Protocol number. The Layer 2 payload of ETHERTYPE_ORCL packets must start with a 16-bit protocol number, which is 1 for the transitive probe packet.
Typ	Probe packet type. The probe packet type is 0 for request and 1 for response.
Pad	Padding (all-zero).
Seq#	Probe sequence number.

Field	Description
Timestamp	Probe timestamp.
TL	Target information length. For Ethernet, target length is 6 bits (Ethernet MAC address length).
Target	Target port MAC address.
SL	Source information length. For Ethernet, source length is 6 bits.
Source	Source port MAC address.

Index

A

- adding
 - links to an external bridge, 87
- administering VLANs on bridged networks, 92
- aggregations *See* link aggregations
- application priority configurations, 126
- application TLV units, 126, 126
 - See also* PFC
- assigning a pvlan to a zone, 77
- auto-enable-agents, 101

B

- basic-tlv, 99
- bridged network ring, 81
- bridged networks, 15, 79
 - adding links to an existing bridge, 87
 - administering VLANs on bridges, 92
 - bridged network ring, 81
 - creating a bridge, 85
 - debugging bridges, 93
 - deleting a bridge, 91
 - displaying configuration information, 89
 - displaying configuration information about bridge links, 91
 - example of creating a bridge, 86
 - example of deleting a bridge from the system, 92
 - example of displaying bridge information, 90
 - example of modifying the protection type for a bridge, 87
 - how a bridge network works, 82
 - modifying the protection type, 87
 - network stack, 81
 - overview, 79
 - protocols, 83
 - removing links, 88

- setting link properties, 88
- simple bridged network, 80
- STP daemon, 84
- TRILL daemon, 85
- VLANs and STP and TRILL protocols, 93

bridges

- adding links to an existing bridge, 87
- configuring a VLAN over a link that is part of a bridge, 92
- creating, 85
- deleting, 91
- naming bridges, 86
- removing links from, 88

bridging protocols, 83

C

- class of service *See* CoS
- community VLAN, 71
- configuring
 - PVLAN, 75
- CoS
 - priority definitions, 116
- creating
 - bridges, 85
 - link aggregations, 28
 - PVLAN, 75
 - VLAN over a link that is part of a bridge, 92
 - VLANs, 53
 - VLANs on legacy device, 59
 - VLANs over a link aggregation, 57
- customizing
 - PFC for DCB, 122

D

data center bridging *See* DCB

datalink multipathing aggregations *See* DLMP aggregations

DCB, 16

- configuring ETS, 127

- considerations, 117

- cos property, 116

- customizing PFC, 121

- enabling DCBX, 119

- enhanced transmission selection (ETS), 116

- overview, 115

- priority-based flow control (PFC), 116, 117

DCBX protocol, 95, 115

defining

- LLDP TLV, 107

deleting

- bridge, 91

- PVLAN, 77

- VLANs, 63

disabling

- LLDP, 109

displaying

- aggregated port information, 39

- bridge configuration information, 89

- datalink properties, 124, 132

- ETS configuration information, 132

- IP address state of aggregation, 40

- LLDP advertised information, 110

- LLDP statistics, 113

- PFC configuration information, 124

- PFC mapping information, 125

- PFC synchronization status, 125

- priority definitions, 126

- probe-ip property values, 40

- probe-related information, 39

- state of the aggregated port, 40

- synchronization status, 132

- VLAN information, 60

- willing property value, 125, 132

dladm command

- add-aggr, 32

- add-bridge, 87

- create-aggr, 28

- create-bridge, 85

- create-vlan, 53, 75

- delete-aggr, 41

- delete-bridge, 91

- delete-vlan, 63, 77

- modify-aggr, 33

- modify-bridge, 87

- modify-vlan, 61, 76

- remove-bridge, 88

- show-aggr, 28

- show-bridge, 89

- show-linkprop, 124

- show-vlan, 60

DLMP aggregations, 23

- advantages, 23

- configuring probe-based failure detection, 35

- example of configuring probe-based failure detection, 37

- example of creating DLMP aggregation, 31

- failure detection, 26

- how DLMP aggregation works, 24

- link-based failure detection, 26

- monitoring probe-based failure detection, 38

- port failure, 25

- probe-based failure detection, 26

- switching to trunk aggregations, 42

- topology, 24

dot1-tlv, 99

dot3-tlv, 99

E

enabling

- DCBX, 119

- LLDP for specific ports, 103

- LLDP globally, 102

- synchronization of PFC information, 123

enhanced transmission selection *See* ETS

ETS, 116, 118

- bandwidth share, 128

- configuring, 127

- displaying information, 132

- ETS TLV units, 129

- example of displaying ETS-related datalink properties, 132

- example of displaying the capability to synchronize ETS information, 132
 - example of recommending ETS configuration to the peer, 130
 - local and remote information, 128
 - properties, 128
 - recommending ETS configuration to the peer, 130
 - setting ETS-related datalink properties, 128
- examples**
- creating a VLAN, 53
 - creating a VLAN with zones, 54, 55
 - creating multiple VLANs over a link aggregation, 58
 - deleting a VLAN configuration, 64, 77
 - migrating multiple VLANs, 63
- F**
- failure detection in DLMP aggregation, 26
 - link-based failure detection, 26
 - probe-based failure detection, 26
- G**
- global TLV units, 100
- H**
- high-availability
 - DLMP aggregations, 23
- I**
- ICMP probing, 27
 - installing
 - LLDP package, 101
 - IPMP
 - link aggregations, comparison with, 135
 - isolated VLAN, 71
- L**
- LACP
 - definition of, 22
 - LACPDU, 22
 - modes, 22
 - using a switch with, 22
 - link aggregation
 - example of adding a link to an aggregation, 32
 - example of deleting a link aggregation, 41
 - Link Aggregation Control Protocol (LACP) *See* LACP
 - link aggregations, 14
 - adding datalinks, 32
 - benefits, 18
 - combined use with VLANs, 64
 - creating, 28
 - deleting, 41
 - DLMP aggregations, 23
 - example of removing a link from an aggregation, 33
 - feature comparison of trunk and DLMP aggregations, 45
 - features, 18
 - IPMP, comparison with, 135
 - overview, 17
 - removing, 33
 - requirements, 27
 - switching between DLMP and trunk aggregations, 42
 - trunk aggregations, 19
 - Link Layer Discovery Protocol *See* LLDP
 - link layer discovery protocol *See* LLDP
 - link properties
 - setting for a bridge, 88
 - link state notification, 28
 - link-based failure detection, 26
 - LLDP, 15, 95
 - agent modes, 97
 - agents, 97
 - auto-enable-agents, 101
 - components in Oracle Solaris, 96
 - defining TLV values, 107
 - disabling, 109
 - displaying advertised information, 110
 - displaying statistics, 113
 - enabling, 101
 - enabling for specific ports, 103
 - enabling globally, 102
 - example of adding optional TLV units, 106
 - example of customizing the auto-enable-agents SMF property, 104
 - example of defining TLV values, 108

- example of displaying advertised information, 110
- example of displaying selected statistics, 113
- example of displaying statistics, 113
- example of displaying the chassis ID and port ID, 98
- example of enabling LLDP on multiple datalinks, 105
- global TLV units, 96, 99
- installing, 101
- lldpd daemon, 97
- management information base (MIB), 97
- monitoring agents, 110
- optional TLV units, 98
- package, 96
- per-agent TLV units, 96, 99
- SMF property for, 101
- specifying agent TLV units, 105
- TLV units, 98, 99

LLDP SMF service, 96

lldpadm command, 96

- reset-agentprop, 109
- set-agentprop, 103, 109
- set-agenttlvprop, 103, 107, 121, 123, 126
- set-tlvprop, 102, 107
- show-agent, 110, 113
- show-agenttlvprop, 107, 124
- show-tlvprop, 107

LLDPDUs, 97

load balancing

- trunk aggregations, 22

local MIB, 97

M

management information base (MIB), 97

managing network datalinks

- bridged networks, 15
- DCB, 16
- features and components, 14
- introduction, 13
- link aggregations, 14
- LLDP, 15
- VLANs, 15
- what's new, 13

mandatory TLV units, 98

migrating

- VLAN, 61

mode of operation for DCB

- setting, 120

modifying

- protection type of a bridge, 87
- PVLAN, 76
- trunk aggregation, 33
- VLAN ID of a VLAN, 61

monitoring

- LLDP agents, 110

N

network stack

- bridge implementation, 81

O

optional TLV units, 98

P

PAUSE frames, 117

per-agent TLV units, 100

PFC, 116, 117

- CoS priority mappings, 122
- customizing, 121
- customizing PFC for DCB, 122
- datalink properties, related, 122
- displaying datalink properties, 124
- displaying information, 124
- example of displaying capability to synchronize PFC information, 125
- example of displaying CoS priority definitions, 126
- example of displaying PFC-related datalink properties, 124
- example of enabling synchronization between host and the peer, 123
- example of verifying symmetry of PFC information, 125
- local and remote information, 122
- PAUSE frames, 117
- pfcmap, 117, 122
- pfcmap-rmt, 122

- synchronized information, 122
- VNIC clients, 126
- PFC mapping, 117
- PFC TLV units, 123
- priority-based flow control *See* PFC
- private virtual local area network, 71
- private VLAN
 - configuring, 75
 - deleting, 77
 - modifying, 76
- private VLANs, 71
- probe-based failure detection, 26
 - configuring, 34
 - displaying aggregated port information, 39
 - displaying probe-ip property values, 40
 - displaying probe-related information, 39
 - displaying the IP address state of aggregation, 40
 - displaying the state of the aggregated port, 40
 - example of configuring probe-based failure detection, 37
 - ICMP probing, 27
 - monitoring, 38
 - transitive probing, 27
- promiscuous trunk port, 73
- protocol data units (PDUs), 97
- protocols
 - DCBX, 115
 - LLDP, 95
 - STP, 81, 83
 - TRILL, 81, 83
- PVLAN
 - advantages, 73
 - assigning to a zone, 77
 - community VLAN, 71
 - configuration requirements, 74
 - configuring, 75
 - deleting, 77
 - isolated VLAN, 71
 - modifying, 76
 - overview, 71
 - ports, 73
 - promiscuous trunk port, 73
 - secondary trunk port, 73
 - with zone, 74
- PVLAN ports, 73
- PVLAN secondary trunk port, 73

PVLANs, 71

R

- recommending
 - ETS configuration to the peer, 130
- remote MIB, 97
- removing
 - link from an aggregation, 33
 - links from a bridge, 88

S

- setting
 - ETS TLV units, 129
 - ETS-related datalink properties, 128
 - mode of operation for DCB, 120
 - PFC TLV units, 123
 - PFC-related datalink properties, 122
 - tagmode, 74
- simple bridged network, 80
- specifying
 - TLV units for LLDP packet of an agent, 105
- STP
 - setting as bridge protection type, 87
- STP daemon, 84
- STP protocol, 83
 - contrasted with TRILL, 83
- switching between DLMP and trunk aggregations, 41

T

- tag mode, 74
- TLV property
 - willing, 123
- topology discovery
 - with LLDP, 96
- transitive probing, 27
- TRILL, 83
 - setting as bridge protection type, 87
- TRILL daemon, 85
- TRILL protocol
 - contrasted with STP, 83
- trunk aggregations, 19
 - back-to-back, 21

- example of creating trunk aggregation, 30
- example of modifying a trunk aggregation, 34
- Link Aggregation Control Protocol (LACP), 22
- load balancing policy, 22
- modifying, 33
- prerequisites, 28
- switching to DLMP aggregations, 42
- unique features, 20
- using a switch, 20
- when to use, 19
- type -length-value (TLV) units, 98
- PVLANS, 13

V

- virt-tlv, 99
- virtual local area networks *See* VLANs
- VLANs, 15, 47
 - combined use with link aggregations, 64
 - configuration, 53
 - creating over link aggregations, 57
 - deleting, 63
 - displaying information, 60
 - example to create a VLAN, 53
 - example to create a VLAN with zones, 54, 55
 - example to create multiple VLANs over a link aggregation, 58
 - example to delete a VLAN configuration, 64, 77
 - legacy devices, on, 59
 - MAC addresses on, 62
 - migrating, 61
 - modifying VLAN IDs, 61
 - overview, 47
 - planning a VLAN configuration, 52
 - STP and TRILL protocols, 93
 - topologies, 48
 - used with zones, 51
 - VLAN names, 48
 - when to use VLANs, 48
 - workgroups, 48

W

- what's new
 - CEE DCBX, 14
 - private virtual local area networks, 13