

Strategies for Network Administration in Oracle® Solaris 11.3



Part No: E54834
November 2016

Part No: E54834

Copyright © 2012, 2016, Oracle and/or its affiliates. All rights reserved.

This software and related documentation are provided under a license agreement containing restrictions on use and disclosure and are protected by intellectual property laws. Except as expressly permitted in your license agreement or allowed by law, you may not use, copy, reproduce, translate, broadcast, modify, license, transmit, distribute, exhibit, perform, publish, or display any part, in any form, or by any means. Reverse engineering, disassembly, or decompilation of this software, unless required by law for interoperability, is prohibited.

The information contained herein is subject to change without notice and is not warranted to be error-free. If you find any errors, please report them to us in writing.

If this is software or related documentation that is delivered to the U.S. Government or anyone licensing it on behalf of the U.S. Government, then the following notice is applicable:

U.S. GOVERNMENT END USERS. Oracle programs, including any operating system, integrated software, any programs installed on the hardware, and/or documentation, delivered to U.S. Government end users are "commercial computer software" pursuant to the applicable Federal Acquisition Regulation and agency-specific supplemental regulations. As such, use, duplication, disclosure, modification, and adaptation of the programs, including any operating system, integrated software, any programs installed on the hardware, and/or documentation, shall be subject to license terms and license restrictions applicable to the programs. No other rights are granted to the U.S. Government.

This software or hardware is developed for general use in a variety of information management applications. It is not developed or intended for use in any inherently dangerous applications, including applications that may create a risk of personal injury. If you use this software or hardware in dangerous applications, then you shall be responsible to take all appropriate fail-safe, backup, redundancy, and other measures to ensure its safe use. Oracle Corporation and its affiliates disclaim any liability for any damages caused by use of this software or hardware in dangerous applications.

Oracle and Java are registered trademarks of Oracle and/or its affiliates. Other names may be trademarks of their respective owners.

Intel and Intel Xeon are trademarks or registered trademarks of Intel Corporation. All SPARC trademarks are used under license and are trademarks or registered trademarks of SPARC International, Inc. AMD, Opteron, the AMD logo, and the AMD Opteron logo are trademarks or registered trademarks of Advanced Micro Devices. UNIX is a registered trademark of The Open Group.

This software or hardware and documentation may provide access to or information about content, products, and services from third parties. Oracle Corporation and its affiliates are not responsible for and expressly disclaim all warranties of any kind with respect to third-party content, products, and services unless otherwise set forth in an applicable agreement between you and Oracle. Oracle Corporation and its affiliates will not be responsible for any loss, costs, or damages incurred due to your access to or use of third-party content, products, or services, except as set forth in an applicable agreement between you and Oracle.

Documentation Accessibility

For information about Oracle's commitment to accessibility, visit the Oracle Accessibility Program website at <http://www.oracle.com/pls/topic/lookup?ctx=acc&id=docacc>.

Access to Oracle Support

Oracle customers that have purchased support have access to electronic support through My Oracle Support. For information, visit <http://www.oracle.com/pls/topic/lookup?ctx=acc&id=info> or visit <http://www.oracle.com/pls/topic/lookup?ctx=acc&id=trs> if you are hearing impaired.

Référence: E54834

Copyright © 2012, 2016, Oracle et/ou ses affiliés. Tous droits réservés.

Ce logiciel et la documentation qui l'accompagne sont protégés par les lois sur la propriété intellectuelle. Ils sont concédés sous licence et soumis à des restrictions d'utilisation et de divulgation. Sauf stipulation expresse de votre contrat de licence ou de la loi, vous ne pouvez pas copier, reproduire, traduire, diffuser, modifier, accorder de licence, transmettre, distribuer, exposer, exécuter, publier ou afficher le logiciel, même partiellement, sous quelque forme et par quelque procédé que ce soit. Par ailleurs, il est interdit de procéder à toute ingénierie inverse du logiciel, de le désassembler ou de le décompiler, excepté à des fins d'interopérabilité avec des logiciels tiers ou tel que prescrit par la loi.

Les informations fournies dans ce document sont susceptibles de modification sans préavis. Par ailleurs, Oracle Corporation ne garantit pas qu'elles soient exemptes d'erreurs et vous invite, le cas échéant, à lui en faire part par écrit.

Si ce logiciel, ou la documentation qui l'accompagne, est livré sous licence au Gouvernement des Etats-Unis, ou à quiconque qui aurait souscrit la licence de ce logiciel pour le compte du Gouvernement des Etats-Unis, la notice suivante s'applique:

U.S. GOVERNMENT END USERS. Oracle programs, including any operating system, integrated software, any programs installed on the hardware, and/or documentation, delivered to U.S. Government end users are "commercial computer software" pursuant to the applicable Federal Acquisition Regulation and agency-specific supplemental regulations. As such, use, duplication, disclosure, modification, and adaptation of the programs, including any operating system, integrated software, any programs installed on the hardware, and/or documentation, shall be subject to license terms and license restrictions applicable to the programs. No other rights are granted to the U.S. Government.

Ce logiciel ou matériel a été développé pour un usage général dans le cadre d'applications de gestion des informations. Ce logiciel ou matériel n'est pas conçu ni n'est destiné à être utilisé dans des applications à risque, notamment dans des applications pouvant causer des dommages corporels. Si vous utilisez ce logiciel ou matériel dans le cadre d'applications dangereuses, il est de votre responsabilité de prendre toutes les mesures de secours, de sauvegarde, de redondance et autres mesures nécessaires à son utilisation dans des conditions optimales de sécurité. Oracle Corporation et ses affiliés déclinent toute responsabilité quant aux dommages causés par l'utilisation de ce logiciel ou matériel pour ce type d'applications.

Oracle et Java sont des marques déposées d'Oracle Corporation et/ou de ses affiliés. Tout autre nom mentionné peut correspondre à des marques appartenant à d'autres propriétaires qu'Oracle.

Intel et Intel Xeon sont des marques ou des marques déposées d'Intel Corporation. Toutes les marques SPARC sont utilisées sous licence et sont des marques ou des marques déposées de SPARC International, Inc. AMD, Opteron, le logo AMD et le logo AMD Opteron sont des marques ou des marques déposées d'Advanced Micro Devices. UNIX est une marque déposée d'The Open Group.

Ce logiciel ou matériel et la documentation qui l'accompagne peuvent fournir des informations ou des liens donnant accès à des contenus, des produits et des services émanant de tiers. Oracle Corporation et ses affiliés déclinent toute responsabilité ou garantie expresse quant aux contenus, produits ou services émanant de tiers, sauf mention contraire stipulée dans un contrat entre vous et Oracle. En aucun cas, Oracle Corporation et ses affiliés ne sauraient être tenus pour responsables des pertes subies, des coûts occasionnés ou des dommages causés par l'accès à des contenus, produits ou services tiers, ou à leur utilisation, sauf mention contraire stipulée dans un contrat entre vous et Oracle.

Accessibilité de la documentation

Pour plus d'informations sur l'engagement d'Oracle pour l'accessibilité à la documentation, visitez le site Web Oracle Accessibility Program, à l'adresse <http://www.oracle.com/pls/topic/lookup?ctx=acc&id=docacc>.

Accès aux services de support Oracle

Les clients Oracle qui ont souscrit un contrat de support ont accès au support électronique via My Oracle Support. Pour plus d'informations, visitez le site <http://www.oracle.com/pls/topic/lookup?ctx=acc&id=info> ou le site <http://www.oracle.com/pls/topic/lookup?ctx=acc&id=trs> si vous êtes malentendant.

Contents

- Using This Documentation 7
- 1 Summary of Oracle Solaris Network Administration 9
 - Highlights of Network Administration in Oracle Solaris 9
 - Basic Network Configuration in Oracle Solaris 10
 - Key Oracle Solaris Network Administration Features 11
 - Network Administration Within the Oracle Solaris Network Protocol Stack 14
 - Network Administration by Functional Area 16
 - Summary of Network Virtualization in Oracle Solaris 19
 - Network Virtualization Building Blocks 19
 - Network Virtualization Strategies 21
 - Creating a Highly Available Virtual Network Stack for Cloud Environments 22
 - Features for Managing Network Resources in Oracle Solaris 24
 - Features for Managing Network Security in Oracle Solaris 25
- 2 Network Configuration Scenarios 27
 - Basic Network Configuration Scenario 27
 - Configure Datalinks, IP Interfaces, and IP Addresses 28
 - Configure Naming Services Through SMF 29
 - Set a System's Host Name 31
 - Combining Aggregations With VNICs for High Availability 31
 - Setting Up an EVS Virtual Tenant Network 33
 - Perform Preliminary Tasks Prior to Creating an EVS Virtual Tenant Network 36
 - Create an EVS Virtual Tenant Network (vswitch) 37
 - Combining Network Virtualization With Oracle VM Server for SPARC to Create a Cloud Environment 39
 - Objectives for Creating and Deploying a Cloud Environment 40
 - Configure a Virtual Network on Oracle VM Server for SPARC Service and Guest Domains 42

Create an EVS Switch to Deploy the Cloud Workload 44

Create Oracle Solaris Zones on the Oracle VM Server for SPARC Guest
Domains 46

Index 49

Using This Documentation

- **Overview** – Provides information about networking strategies and describes how to use networking features to administer network configuration in the Oracle Solaris operating system (OS).
- **Audience** – System administrators.
- **Required knowledge** – Basic understanding of network administration concepts and practices.

Product Documentation Library

Documentation and resources for this product and related products are available at <http://www.oracle.com/pls/topic/lookup?ctx=E53394>.

Feedback

Provide feedback about this documentation at <http://www.oracle.com/goto/docfeedback>.

Summary of Oracle Solaris Network Administration

This chapter provides a summary of network administration in the Oracle Solaris release, including specific information about the network virtualization features that are supported.

For examples of network configuration scenarios, including basic network configuration of a system, as well as network virtualization use cases, see [Chapter 2, “Network Configuration Scenarios”](#).

For a shortcut to commonly used networking commands, see [Oracle Solaris 11.3 Network Administration Cheatsheet](#).

This chapter contains the following topics:

- [“Highlights of Network Administration in Oracle Solaris” on page 9](#)
- [“Network Administration Within the Oracle Solaris Network Protocol Stack” on page 14](#)
- [“Network Administration by Functional Area” on page 16](#)
- [“Summary of Network Virtualization in Oracle Solaris” on page 19](#)
- [“Features for Managing Network Resources in Oracle Solaris” on page 24](#)
- [“Features for Managing Network Security in Oracle Solaris” on page 25](#)

Highlights of Network Administration in Oracle Solaris

Users rely on various networking technologies to communicate, share, store, and process information. One of the primary goals of network administration is to establish and maintain reliable, secure, and efficient data communications on systems that are running the Oracle Solaris release. See [“Basic Network Configuration in Oracle Solaris” on page 10](#).

Beyond the basic configuration that is required to connect systems to the network, Oracle Solaris also supports several advanced networking technologies, including features that provide support for following functional areas:

- High availability
- Network security
- Network storage

- Network virtualization
- Observability, monitoring, and debugging
- Performance and efficiency
- Resource management

Most of these features are designed to address the complexities of modernized network environments by enabling you to use a more modular and layered approach to administering the various aspects of network configuration. For more information, see [“Key Oracle Solaris Network Administration Features” on page 11](#) and [“Network Administration by Functional Area” on page 16](#).

Basic Network Configuration in Oracle Solaris

Basic network configuration of a system evolves in two stages: assembling the hardware and then configuring the daemons, files, and services that implement the network protocol stack. For more information about how the various networking components are configured within the network protocol stack, see [“Network Administration Within the Oracle Solaris Network Protocol Stack” on page 14](#).

For an example of the information that is described in this section, see [“Basic Network Configuration Scenario” on page 27](#).

The basic network configuration process typically involves the following tasks:

- You first customize the physical datalinks on the system. Each datalink represents a link object in the second layer (L2) of the Open Systems Interconnection (OSI) model. In this release, generic names are automatically assigned to datalinks by using the `net0`, `net1`, `netN` naming convention. The name that is assigned to each datalink depends on the total number of network devices that are on that system. For more information, see [Chapter 2, “Administering Datalink Configuration in Oracle Solaris” in *Configuring and Managing Network Components in Oracle Solaris 11.3*](#).
- After you customize the datalinks on a system, you then configure IP interfaces and addresses over each datalink. This configuration takes place at the network layer (L3) of the OSI model. You obtain unique IP addresses to communicate to public networks on the Internet. See [Chapter 3, “Configuring and Administering IP Interfaces and Addresses in Oracle Solaris” in *Configuring and Managing Network Components in Oracle Solaris 11.3*](#).

Oracle Solaris supports both IPv4 and IPv6 configuration. You can choose whether to deploy a purely IPv4 network, an IPv6 network, or a network that uses a combination of both types of IP addresses. Deploying an IPv4 or IPv6 network requires some advanced planning. For more information about deploying a physical network in an organized and cost-effective manner, see [Planning for Network Deployment in Oracle Solaris 11.3](#).
- Naming services and other system-wide network settings are fundamental to any computing network. These services perform lookups of stored information such as host names and addresses, user names, passwords, access permissions, and so on. The information is made available so that users can log in to their system, access resources, and be granted

permissions. Naming service information is centralized in the form of files, maps, and database files to make network administration more manageable. In this release, naming services are managed through the Service Management Facility (SMF). For more information about configuring system-wide network settings on an Oracle Solaris client, see [Chapter 4, “Administering Naming and Directory Services on an Oracle Solaris Client” in *Configuring and Managing Network Components in Oracle Solaris 11.3*](#).

- Network administration might also involve configuring systems that perform specific functions within your network, for example, routers, IP tunnels, and so on. For additional information, see [Configuring an Oracle Solaris 11.3 System as a Router or a Load Balancer](#) and [Administering TCP/IP Networks, IPMP, and IP Tunnels in Oracle Solaris 11.3](#).

Before beginning the task of configuring a client system on the network, see [“Information That Is Required to Configure Client Systems on the Network” in *Configuring and Managing Network Components in Oracle Solaris 11.3*](#).

Key Oracle Solaris Network Administration Features

Oracle Solaris supports several networking features that you can use for different purposes. The following are some of the key features that are supported in this release. This list is not exhaustive:

- **Aggregation** – Is an L2 entity that is used to ensure that a system has continuous access to the network. Link aggregations increase the availability and reliability of network connectivity by enabling you to pool multiple datalink resources that you administer as a single unit. See [Chapter 2, “Configuring High Availability by Using Link Aggregations” in *Managing Network Datalinks in Oracle Solaris 11.3*](#).

The following types of aggregations are supported:

- **Datalink Multipathing (DLMP)** – Is a type of link aggregation that supports multiple switches and provides continuous connectivity to its datalinks. When a switch fails, the aggregation continues to provide connectivity to its datalinks by using the other switches. This type of link aggregation does not require switch configuration. Using DLMP aggregation can help overcome some of the disadvantages of using trunk aggregation. See [“Datalink Multipathing Aggregations” in *Managing Network Datalinks in Oracle Solaris 11.3*](#).
- **Trunk aggregation** – Is a link aggregation mode that is based on the IEEE 802.3ad standard and works by enabling multiple flows of traffic to be spread across a set of aggregated ports. IEEE 802.3ad requires switch configuration and switch-vendor proprietary extensions in order to work across multiple switches. See [“Trunk Aggregations” in *Managing Network Datalinks in Oracle Solaris 11.3*](#).
- **Bridging** – Is an L2 technology that connects multiple datalinks on a network into a single network. For bridging, Oracle Solaris supports the Spanning Tree Protocol (STP) and TRansparent Interconnection of Lots of Links (TRILL) protocols. See [Chapter 5, “Administering Bridging Features” in *Managing Network Datalinks in Oracle Solaris 11.3*](#).

- **Edge Virtual Bridging (EVB)** – Is an L2 technology that enables hosts to exchange virtual link information with an external switch. EVB offloads the enforcement of traffic service level agreements (SLAs) to the switch. See [Chapter 4, “Administering Server-Network Edge Virtualization by Using Edge Virtual Bridging” in *Managing Network Virtualization and Network Resources in Oracle Solaris 11.3*](#).
- **Data Center Bridging (DCB)** – Is an L2 technology that is used to manage the bandwidth, relative priority, and flow control of multiple traffic types that share the same network link, for example, when sharing a datalink between networking and storage protocols. See [Chapter 7, “Managing Converged Networks by Using Data Center Bridging” in *Managing Network Datalinks in Oracle Solaris 11.3*](#).
- **Elastic Virtual Switch (EVS)** – Is an L2 technology that expands network virtualization capabilities by enabling you to manage virtual switches across multiple hosts. With the Oracle Solaris EVS feature, you can deploy virtual networks that span multiple hosts within either a multi-tenant cloud environment or a large deployment. See [Chapter 6, “Administering Elastic Virtual Switches” in *Managing Network Virtualization and Network Resources in Oracle Solaris 11.3*](#).
- **Etherstub** – Is a pseudo Ethernet NIC that is configured at the datalink layer (L2) of the Oracle Solaris network protocol stack. You can create virtual interface cards (VNICs) over etherstubs instead of physical links for the purpose of constructing a private virtual network that is isolated from other virtual networks on the system, as well as from the external network. See [“How to Configure VNICs and Etherstubs” in *Managing Network Virtualization and Network Resources in Oracle Solaris 11.3*](#).
- **Flows** – Are a subset of packets that are identified by common attributes. These attributes consist of packet header information such as IP addresses, protocol type, and transport port numbers. You can observe flows individually, as well as assign flows their own SLAs, for example, bandwidth control and priority. You administer flows at the L2, L3, and L4 layers of the Oracle Solaris network protocol stack. For more information, see [“Features for Managing Network Resources in Oracle Solaris” on page 24](#).
- **Integrated Load Balancer (ILB)** – Is an L3 and L4 technology that enables a system to spread the load of network processing amongst available resources. ILB can be used to improve reliability and scalability and to minimize the response time of network services. Load balancing involves using multiple systems to deal with high demands of a network by balancing the load between multiple systems. Support for ILB in Oracle Solaris includes stateless Direct Server Return (DSR) and Network Address Translation (NAT) modes of operation for IPv4 and IPv6, as well as server monitoring capabilities through health checks. See [“Features of ILB” in *Configuring an Oracle Solaris 11.3 System as a Router or a Load Balancer*](#).
- **IP network multipathing (IPMP)** – Is an L3 technology that ensures that a system has continuous access to the network. With IPMP, you configure multiple IP interfaces into an *IPMP group*. The IPMP group functions like an IP interface with data addresses for sending or receiving network traffic. If an underlying interface in the group fails, the data addresses are redistributed amongst the remaining underlying active interfaces in the group.

The IPMP model and administrative interface has undergone some changes in Oracle Solaris 11. To familiarize yourself with the new model, see [“What’s New in IPMP” in *Administering TCP/IP Networks, IPMP, and IP Tunnels in Oracle Solaris 11.3*](#).

Link aggregations work similarly to IPMP for improving network performance and availability, but at the datalink layer (L2). Aggregations are recommended when combining features for high availability in a virtualized environment. For a comparative analysis, see [Appendix A, “Link Aggregations and IPMP: Feature Comparison,” in *Managing Network Datalinks in Oracle Solaris 11.3*](#).

- **IP tunnel** – Is an L3 technology that provides a means for transporting data packets between domains when the protocol in those domains is not supported by intermediary networks. See [Chapter 4, “About IP Tunnel Administration” in *Administering TCP/IP Networks, IPMP, and IP Tunnels in Oracle Solaris 11.3*](#).
- **Link Layer Discovery Protocol (LLDP)** – Is an L2 technology that is used by systems in a local area network (LAN) to exchange configuration and management information with each other. LLDP enables a system to advertise connectivity and management information to other systems on the network. See [Chapter 6, “Exchanging Network Connectivity Information With Link Layer Discovery Protocol” in *Managing Network Datalinks in Oracle Solaris 11.3*](#).
- **Virtual local area network (VLAN)** – Is an L2 technology that enables you to divide a LAN into subnetworks without having to add a physical network environment. A VLAN is a subdivision of a LAN at the datalink layer of the network protocol stack. For more information, see [Chapter 3, “Configuring Virtual Networks by Using Virtual Local Area Networks” in *Managing Network Datalinks in Oracle Solaris 11.3*](#).
- **Virtual eXtensible area network (VXLAN)** – Is an L2 and L3 technology that works by overlaying a datalink (L2) network on top of an IP (L3) network. VXLANs address the 4K limitation that is imposed when using VLANs. Typically, VXLANs are used in a cloud infrastructure to isolate multiple virtual networks. You can manage VXLANs by using the EVS feature. For more information, see [Chapter 3, “Configuring Virtual Networks by Using Virtual Extensible Local Area Networks” in *Managing Network Virtualization and Network Resources in Oracle Solaris 11.3*](#).
- **Virtual network interface card (VNIC)** – Is an L2 entity or virtual network device that behaves just like a physical NIC when configured. You configure a VNIC over an underlying datalink to share it among multiple Oracle Solaris zones or VMs. See [“Configuring the Components of a Virtual Network” in *Managing Network Virtualization and Network Resources in Oracle Solaris 11.3*](#).

In this release, you can also manage network devices that support single root I/O virtualization (SR-IOV). For more details, see [“Using Single Root I/O Virtualization With VNICs” in *Managing Network Virtualization and Network Resources in Oracle Solaris 11.3*](#).

- **Virtual Router Redundancy Protocol (VRRP)** – Is an L3 technology that provides high availability of IP addresses, such as those that are used for routers and load balancers. Oracle Solaris supports both L2 and L3 VRRP. L3 VRRP removes the need to configure unique VRRP virtual MAC addresses for VRRP routers, thereby providing better support for VRRP over IPMP, InfiniBand interfaces, and zones. For more information, see [Chapter 3, “Using Virtual Router Redundancy Protocol” in *Configuring an Oracle Solaris 11.3 System as a Router or a Load Balancer*](#).

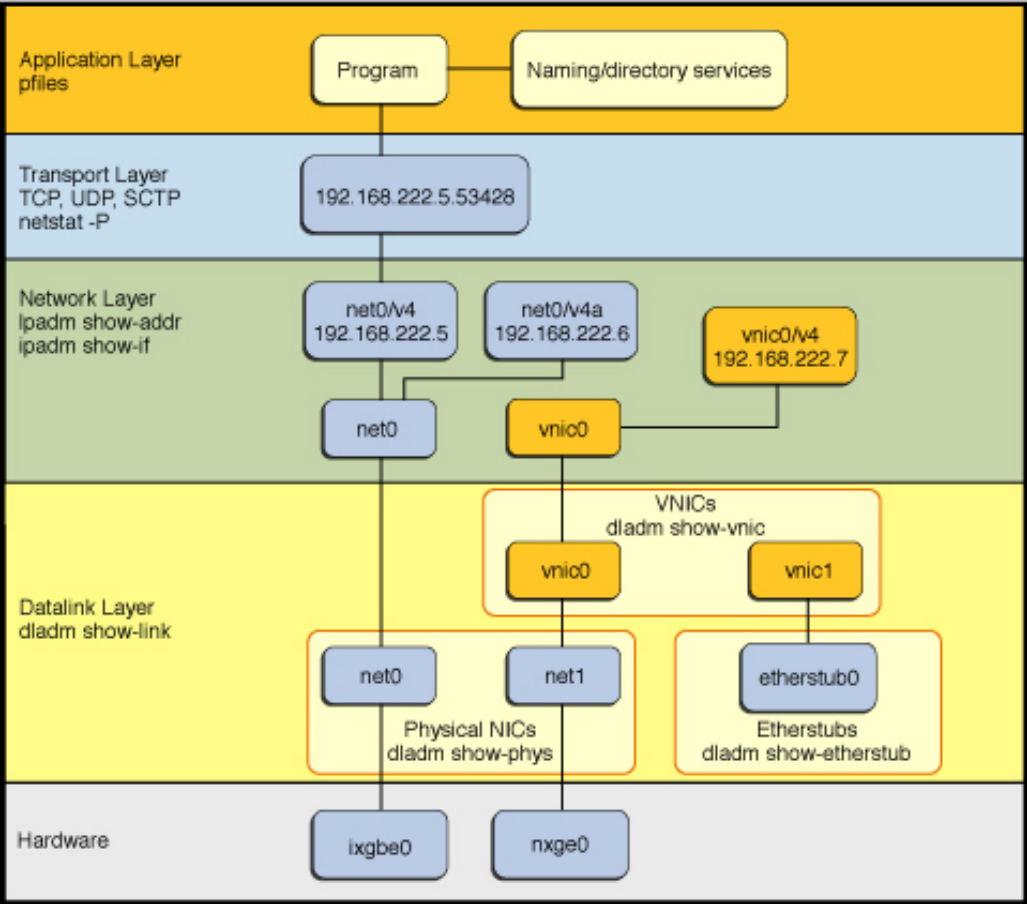
- **Virtual switch** – Is an L2 technology that simulates the capabilities of a physical network switch. A virtual switch is implicitly created whenever you create a VNIC on top of an underlying datalink. Virtual switches provide a method for virtual machines and zones to transfer packets. You can manage virtual switches by using the EVS feature. For more information, see [“Network Virtualization Components” in *Managing Network Virtualization and Network Resources in Oracle Solaris 11.3*](#).

Network Administration Within the Oracle Solaris Network Protocol Stack

The following figure shows the layers of the Oracle Solaris network protocol stack and where within the stack both the physical and virtual interfaces are administered. This information can be helpful when planning which networking strategy or strategies to deploy at your site. Knowing which layer of the network protocol stack a particular feature is configured is also helpful when troubleshooting network configuration issues, detecting network connectivity problems, and diagnosing performance issues, such as packet loss.

The information in [Table 1, “Networking Features by Network Protocol Stack Layer,” on page 15](#) provides additional details about where each feature is administered within the Oracle Solaris network protocol stack. For more information about administering network configuration at each layer of the network protocol stack, as well as examples that show how to observe network traffic usage at each layer of the stack, see [Chapter 2, “Using Observability Tools to Monitor Network Traffic Usage” in *Troubleshooting Network Administration Issues in Oracle Solaris 11.3*](#).

FIGURE 1 Physical and Virtual Network Administration Within the Network Protocol Stack



The following table further describes which layer of the Oracle Solaris network protocol stack, each networking feature is administered. Note that some features are administered on more than one layer of the stack.

Note - Only those layers of the network protocol stack that pertain to the various network administration features described in this document are shown.

TABLE 1 Networking Features by Network Protocol Stack Layer

Network Protocol Stack Layer	Feature or Technology
Transport (L4)	■ Firewalls

Network Protocol Stack Layer	Feature or Technology
	<ul style="list-style-type: none">■ Flows■ Pluggable congestion control■ Socket filtering
Protocol or network (L3)	<ul style="list-style-type: none">■ DHCP■ Flows■ IP interfaces and IP addresses■ IP tunnels■ IPMP■ ILB■ Routing■ VNIs■ VRRP■ VXLANs
Datalink (L2)	<ul style="list-style-type: none">■ Aggregations (DLMP and trunking)■ EVB■ Flows■ LLDP■ Physical datalinks■ Network virtualization features:<ul style="list-style-type: none">■ DCB■ Etherstubs■ EVS■ Virtual switches■ VLANs■ VNICs■ VXLANs

Network Administration by Functional Area

Oracle Solaris network administration features are designed to meet specific networking needs by providing support in the following functional areas: high availability, network virtualization, performance, resource management, security, and storage. Knowing which functional area a particular feature supports is helpful for evaluating which networking strategy or strategies to implement at your site.

The following table describes the various network administration features that are supported in Oracle Solaris according to functional area. Information about the administrative interface that is used to administer the feature and at which layer of the network protocol stack the feature is administered is also provided.

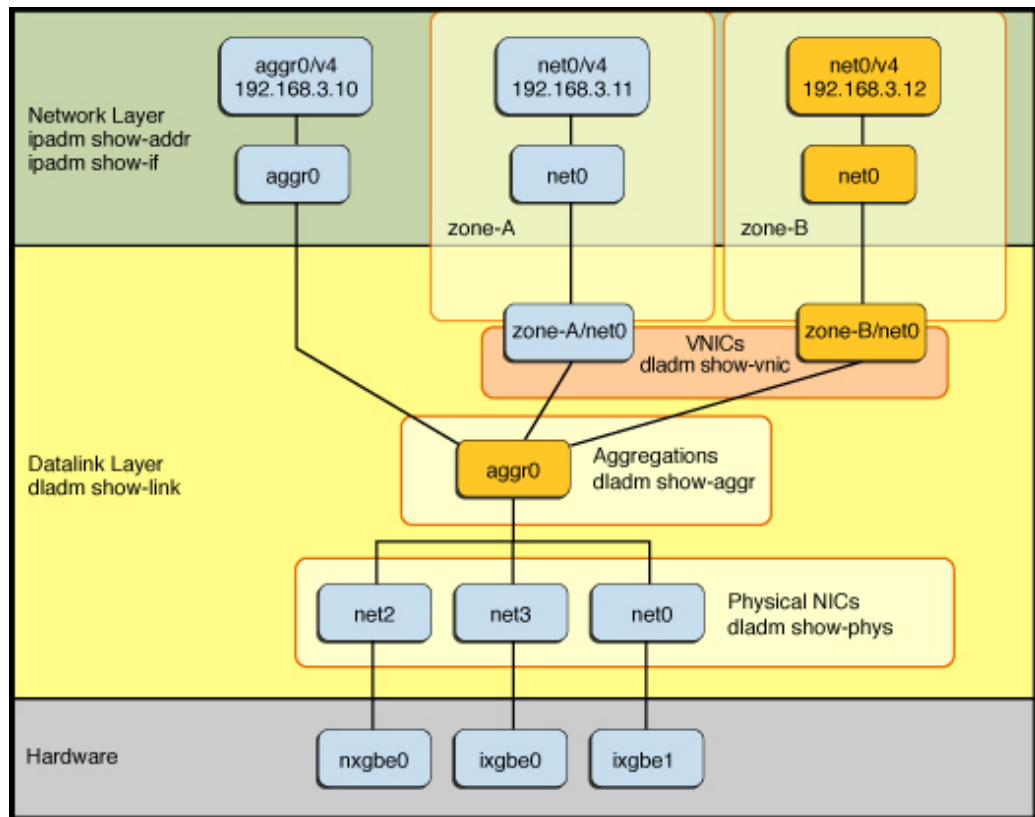
TABLE 2 Networking Features by Functional Area

Feature	Functional Area	Administrative Interface	Network Protocol Stack Layer
Aggregations (DLMP and trunking)	High availability	dladm (create-aggr, delete-aggr, modify-aggr, add-aggr, remove-aggr)	L2
Bridging protocols: ■ STP ■ TRILL	High availability, network virtualization	dladm (create-bridge, delete-bridge, modify-bridge, add-bridge, remove-bridge, show-bridge)	L2
DCB	Network storage, performance	lldpadm, dladm	L2
Etherstubs	Network virtualization	dladm (create-etherstub, delete-etherstub, show-etherstub)	L2
EVB	Network virtualization	dladm	L2
EVS	Network virtualization	evsadm, evsstat, dladm	L2, L3
Firewalls	Security	Packet filtering with ipf and ipnat	L3, L4
Flows	Observability, resource management, security	flowadm, flowstat	L2, L3, L4
ILB	Performance	ilbadm (create-servergroup, add-server, delete-servergroup, enable-server, disable-server, show-server, show-servergroup, remove-server)	L3
IPMP	High availability	ipadm (create-ipmp interface, delete-ipmp interface, add-ipmp interface, remove-ipmp interface)	L3
IP tunnels	IP connectivity	dladm (create-iptun, modify-iptun, delete-iptun, show-iptun); ipadm (to create the IP address over the tunnel)	L2, L3
LLDP	Observability, network storage, network virtualization	lldpadm	L2
Pluggable congestion control	Performance	ipadm set-prop <i>property</i>	L4
Routing	IP connectivity	route (route -p display; netstat); routeadm	L3
Socket filtering	Security	soconfig (-F)	L4
VLANs	Network virtualization	dladm (create-vlan, modify-vlan, delete-vlan, show-vlan)	L2
VNI	IP connectivity	ipadm (create-vni, delete-vni)	L3
VNICs	Network virtualization	dladm (create-vnic, modify-vnic, delete-vnic, show-vnic)	L2

Feature	Functional Area	Administrative Interface	Network Protocol Stack Layer
VRRP	High availability	dladm, vrrpadm	L3
VXLANs	Network virtualization	dladm (create-vxlan, show-vxlan, delete-vxlan)	L2, L3

In many cases, you can obtain optimal results by using a combination of networking features. For example, the following figure shows how you might combine multiple networking features for high availability.

FIGURE 2 Combining the Use of Aggregations With VNICs



In the figure, multiple physical datalinks (`net0`, `net2`, and `net3`) are combined into a single link aggregation (`aggr0`). The aggregation datalink is then directly configured from IP in the global zone through the `aggr0` and `aggr0` IP interface and IP address, respectively. For another example, see [“Combining Aggregations With VNICs for High Availability”](#) on page 31.

You can also virtualize the aggregation datalink by using it as the underlying link for the VNICs. In this figure, two VNICs are configured and then assigned to two non-global zones. This particular configuration makes the VNICs highly available because any failures of the underlying physical NICs that occur are automatically handled by the link aggregation layer and are transparent to the zones.

Summary of Network Virtualization in Oracle Solaris

With server virtualization becoming more mainstream in the IT industry, the focus is shifting to a deployment model that uses network virtualization to support the sharing of network traffic amongst multiple virtual machines (VMs) or zones. Along with a rise in the adoption of cloud architectures that rely upon virtualization for deploying workloads, network virtualization is playing an even more critical role in the overall network administration strategy in Oracle Solaris.

Virtual environments require a high degree of availability, insulation, performance, and separation. Oracle Solaris provides several features that meet these requirements. Additionally, Oracle Solaris network virtualization features are tightly integrated with other Oracle Solaris features (subsystems). For example, when configuring a zones environment, you can create VNICs (anets) that are automatically configured when the zone boots. For information about working with Oracle Solaris zones, see [Introduction to Oracle Solaris 11 Virtual Environments](#).

Network virtualization is also tightly integrated with Oracle Solaris resource management features, which are used to limit the CPUs in a zones environment. For more information about network virtualization and resource management features in Oracle Solaris, see [Managing Network Virtualization and Network Resources in Oracle Solaris 11.3](#).

For information about Oracle VM, including Oracle VM Server for x86, Oracle VM Server for SPARC (previously called Sun Logical Domains, or LDoms), and Oracle VM Manager, see the documentation at <http://www.oracle.com/technetwork/documentation/vm-sparc-194287.html>.

Oracle also provides the Oracle Enterprise Manager Ops Center for managing some aspects of network virtualization, for example, the ability to create virtual private networks inside a virtual data center. For more information about Oracle Enterprise Manager Ops Center, refer to the Oracle Help Center at <http://docs.oracle.com/en/>.

For more information and examples, see the scenarios that are described in [Chapter 2, “Network Configuration Scenarios”](#).

Network Virtualization Building Blocks

Network virtualization in Oracle Solaris includes the following key building blocks:

- VNICs

When a datalink such as a physical NIC or a link aggregation needs to be shared by multiple VMs or zones, you can carve it up into *virtual NICs* or VNICs. These VNICs appear on the system as any other NIC and are administered exactly the same way as a physical NIC. Each VNIC has its own MAC address that you can configure with additional attributes, such as a *VLAN ID*, thus enabling the VNIC to be easily integrated into an existing network infrastructure. For higher availability, you can also create VNICs on top of link aggregations, which you can then assign individual bandwidth limits, thereby enabling them to consume just their assigned share of bandwidth. VNICs have a rich set of configurable features. For more information, see [“Building Virtual Networks” in *Managing Network Virtualization and Network Resources in Oracle Solaris 11.3*](#).

- **Virtual switching**

The Oracle Solaris virtual network stack includes built-in virtual switching capabilities that simulate the capabilities of a physical network switch. You can use virtual switches within a single machine to enable zones and VMs to communicate with each other. Virtual switches are automatically instantiated when multiple VNICs are created on top the same datalink. In addition to being able to create VNICs on top of physical NICs or aggregations, you can also create virtual switches on top of an *etherstub*. This capability enables you to create fully virtualized networks that are independent of the physical hardware. For more information, see [“Configuring the Components of a Virtual Network” in *Managing Network Virtualization and Network Resources in Oracle Solaris 11.3*](#).

- **Elastic Virtual Switch (EVS) feature of Oracle Solaris**

The EVS feature is an L2 technology that expands network virtualization capabilities by enabling you to directly manage virtual switches. You can create EVS switches to deploy multiple virtual networks that span multiple hosts, within either a multi-tenant cloud environment or a large deployment. You can also optionally configure an EVS switch with virtual ports, IP subnets, and Service Level Agreements (SLAs). Additionally, you can connect any Oracle Solaris VNIC to an EVS switch or a virtual port. Such VNICs automatically inherit their network configuration from EVS. This capability enables you to more cleanly separate the network configuration from a zone or VM configuration.

You manage and observe EVS switches through a central *controller*. The elastic virtual switches are then automatically deployed on the various hosts, as needed. Hence, the term *elastic* is used to describe these switches. The EVS architecture is tightly integrated with various other network virtualization features, including the VXLAN feature. See [Chapter 3, “Configuring Virtual Networks by Using Virtual Extensible Local Area Networks” in *Managing Network Virtualization and Network Resources in Oracle Solaris 11.3*](#). Together, you can use these two features to create a large number of virtual networks. Also, because EVS switches are transport-agnostic, you can use them with other types of network fabrics, such as traditional VLANs.

EVS switches are also supported in a zones environment. An *anet* VNIC resource can connect to an EVS switch by using the appropriate *zonecfg* properties. See [Creating and Using Oracle Solaris Zones](#) and the [zonecfg\(1M\)](#) man page for more information.

For more information about the EVS feature, see [Chapter 6, “Administering Elastic Virtual Switches”](#) in *Managing Network Virtualization and Network Resources in Oracle Solaris 11.3*.

The EVS feature introduces new administrative commands. For more information, see the [evsadm\(1M\)](#) and [evsstat\(1M\)](#) man pages. See also the [dladm\(1M\)](#) man page.

Network Virtualization Strategies

You can deploy Oracle Solaris network virtualization features for the following purposes:

- **Workload consolidation**

It is a common practice in large deployments to consolidate multiple workloads onto a single system. This type of workload consolidation is typically accomplished by using virtualization on multiple VMs or zones or by using a combination of both methods. To provide network access to these entities, Oracle Solaris network virtualization features provide a way for the physical NICs that are on a system to be virtualized into multiple VNICs. By virtualizing the physical NICs, you eliminate the need to have separate physical NICs for each VM or zone. The VM or zone shares the physical NIC. As is the case for other virtualized resources, it is important to control the share of the network resources to which each virtual machine is entitled. To accomplish this task, you can configure bandwidth limits on the individual VNICs. Using resource control along with VNICs can further improve the use of resources in a multiple virtual network stack.

- **Private virtual networks**

You can also use network virtualization features to build private virtual networks for the following purposes:

- **Security** – Create a private virtual network behind a virtual firewall to better insulate virtual machines from the physical network and also to better insulate the physical network from the virtual network.
- **Testing and simulation** – Create a private virtual *network within a box* to test different features or simulate a feature's behavior under a given network load before actually implementing that new feature or network configuration.
- **Network consolidation** – Consolidate multiple hosts, network functions, and various network devices such as routers, firewalls, load balancers, and so on, within a box.

- **Cloud networking**

A *cloud architecture* is a network administration approach that uses a *utility computing model* to deploy workloads. In this administrative model, multiple tenants share the same cloud and therefore must be isolated from one another. A cloud architecture is highly dynamic.

Oracle Solaris provides several network virtualization features that are ideal for this type of environment. For example, you could use the EVS feature to build virtual network topologies that span multiple hosts while providing a single point of control and observability.

With EVS, a cloud administrator can easily provision, control, and observe a per-tenant, virtual network. This type of configuration includes the capability for satisfying the most demanding requirements for agility and security in modern cloud environments. For more information about how you would set this type of scenario up, see [“Setting Up an EVS Virtual Tenant Network” on page 33](#).

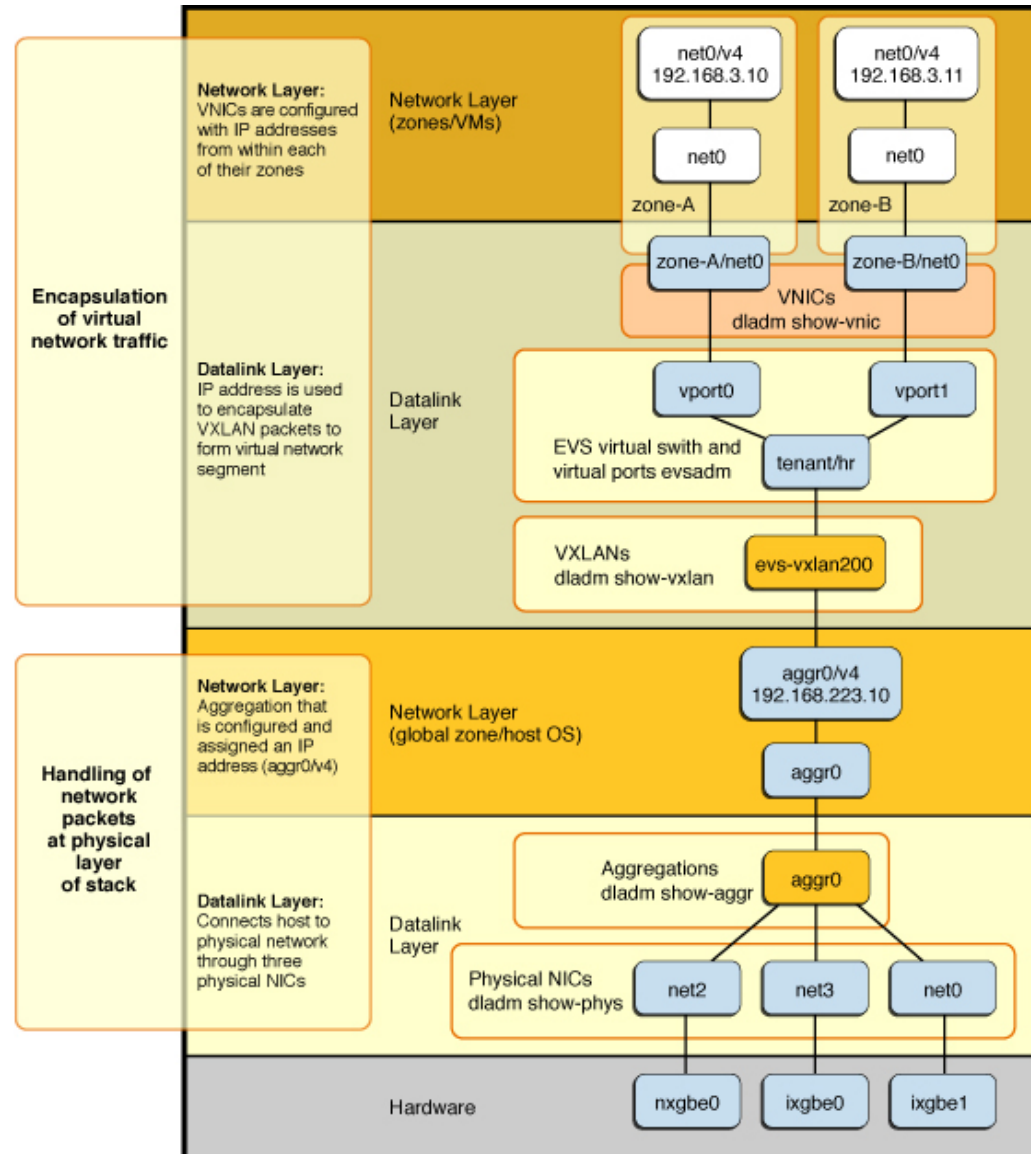
For background information, see [Chapter 5, “About Elastic Virtual Switches” in *Managing Network Virtualization and Network Resources in Oracle Solaris 11.3*](#).

Creating a Highly Available Virtual Network Stack for Cloud Environments

The following figure shows how you might combine multiple network virtualization features, for example, link aggregations, VNICs, VXLANs, and an EVS switch, to provide a highly available and integrated virtual network stack for cloud environments.

In this figure the datalink and network layers of the network protocol stack appear twice. This layering is the result of using VXLANs which provide virtual network segments that are encapsulated within IP packets. The datalink and network layers therefore appear twice in the resulting network stack: once to show the handling of packets at the physical layer and then again to show the encapsulated virtual network traffic within these layers of the stack.

The datalink layer that is shown at the lower level of the figure (just above the hardware layer) is used to connect the host to the physical network through three physical NICs, which are aggregated to provide high availability. The resulting aggregation is configured at the network layer and assigned an IP address (aggr0/v4). The same IP address is then used to encapsulate VXLAN packets which form a virtual network segment. In Oracle Solaris, VXLANs are configured through datalinks and then used through VNICs. These VNICs are then configured with IP addresses from within their zones, as depicted in the datalink and network layers that are shown in the top portion of the figure.

FIGURE 3 Combining Aggregations With VXLANs, VNICs, and an EVS Switch

The figure represents the following configuration:

1. Starting from the hardware layer, multiple physical NICs (net0, net2, and net3) are aggregated to form a highly available link aggregation called aggr0.

2. The aggregation is then configured with an IP address, `aggr0/v4` (`192.168.223.10`).
3. An EVS virtual switch `tenant/hr` is created on top of the IP interface `aggr0`. In this figure, EVS is configured to use a VXLAN.

The new `vxlان0` datalink is connected to a virtual L2 network that overlays the IP network.

4. Assuming that EVS assigned the virtual switch a VXLAN ID of 200, EVS automatically creates a VXLAN datalink called `evs-vxlان200`, which is associated with the `tenant/hr` virtual switch.
5. The EVS switch has two virtual ports (`vport0` and `vport1`), which are connected to two VNICs that are used by two zones. The VNICs appear in the zone as datalinks named `net0` and are visible from the global zone as `zone-A/net0` and `zone-B/net0`.

For some examples of how you would deploy these features, see [Chapter 2, “Network Configuration Scenarios”](#).

Features for Managing Network Resources in Oracle Solaris

Network resource management in Oracle Solaris consists of setting datalink properties that pertain specifically to how network resources are allocated. By setting these properties, you can determine how much of a given resource can be used for networking processes. For example, a link can be associated with a specific number of CPUs that are reserved exclusively for networking processes. Or, a link can be allotted a given bandwidth to process a specific type of network traffic.

Procedures for allocating resources apply to both virtual networks and traditional (physical) networks. For example, you can use the `dladm set-linkprop` command to set properties that are related to network resources. This same syntax is used for both physical and virtual datalinks.

Network resource management is comparable to creating dedicated lanes for traffic. When you combine different resources to cater to specific types of network packets, those resources form a network lane specifically for those network packets.

Use network resource management features to accomplish the following:

- Provision a network
- Establish service level agreements
- Bill clients
- Diagnose security problems

You can also use flows to manage network resources. A *flow* is a customized way of categorizing packets to further control how resources are used to process those packets. Network packets can be categorized according to an attribute. Packets that share an attribute constitute a flow and are labeled with a specific *flow name*. You can then assign specific resources to the flow.

The commands that you use to allocate network resources depend on whether you are working with datalinks or flows.

- For datalinks, use the `dladm` command with the appropriate subcommand, depending on whether you are setting the property during link creation or afterwards.
- For flows, use the `flowadm` command with the appropriate subcommand. Managing resources on flows parallels the methods for managing resources on datalinks.

You can use the `flowadm add-flow` command to configure flows on a datalink based on a single attribute or a combination of attributes. Configuring flows based on a combination of attributes enables you to selectively organize network packets that are received from different ports, transport protocols, and IP addresses.

The set of defined attributes that characterizes the flows constitutes the system's *flow control policy*.

For complete instructions, see [Chapter 7, “Managing Network Resources”](#) in *Managing Network Virtualization and Network Resources in Oracle Solaris 11.3* and the `dladm(1M)` and `flowadm(1M)` man pages.

Features for Managing Network Security in Oracle Solaris

Oracle Solaris provides several security features that enable you to protect and secure your network. The following table briefly describes several key network security features.

TABLE 3 Oracle Solaris Network Security Features

Features and Methods Used to Secure the Network	Description	For More Information
Link protection	The link protection mechanism provides protection from basic threats to a network, for example, IP, DHCP, and MAC spoofing, as well as L2 frame spoofing and Bridge Protocol Data Unit (BPDU) attacks.	Chapter 1, “Using Link Protection in Virtualized Environments” in <i>Securing the Network in Oracle Solaris 11.3</i>
Network parameter tuning	Tuning network parameters ensures that the network is secure and prevents malicious attacks, for example, various types of denial-of-service (DoS) attacks.	Chapter 2, “Tuning Your Network” in <i>Securing the Network in Oracle Solaris 11.3</i>
Secure Sockets Layer (SSL) protocol for web server communications	The SSL protocol encrypts and accelerates web server communications on your Oracle Solaris system. SSL provides confidentiality, message integrity, and endpoint authentication between two applications.	Chapter 3, “Web Servers and the Secure Sockets Layer Protocol” in <i>Securing the Network in Oracle Solaris 11.3</i>
OpenBSD Packet Filter (PF)	PF is a network firewall that captures and evaluates inbound packets for	Chapter 4, “OpenBSD Packet Filter Firewall in Oracle Solaris”

Features and Methods Used to Secure the Network	Description	For More Information
	entry to and exit from the system. PF provides stateful packet inspection. PF can match packets by IP address and port number, as well as the receiving network interface.	in <i>Securing the Network in Oracle Solaris 11.3</i>
IP Filter feature of Oracle Solaris	Packet filtering provides basic protection against attacks on your network. The IP Filter feature of Oracle Solaris is a firewall that provides stateful packet filtering and network address translation (NAT). IP Filter also includes stateless packet filtering and the ability to create and manage address pools.	Chapter 6, “IP Filter Firewall in Oracle Solaris” in <i>Securing the Network in Oracle Solaris 11.3</i>
IP Security Architecture (IPsec)	IPsec provides cryptographic protection for IP datagrams in IPv4 and IPv6 network packets. IPsec includes several components that provide protection for IP packages by authenticating or encrypting the packets.	Chapter 9, “Configuring IPsec” in <i>Securing the Network in Oracle Solaris 11.3</i>
Internet Key Exchange (IKE)	The IKE feature automates key management for IPsec. IKE easily scales to provide a secure channel for a large volume of traffic.	Chapter 11, “Configuring IKEv2” in <i>Securing the Network in Oracle Solaris 11.3</i>

Network Configuration Scenarios

This chapter contains one basic network configuration scenario and three network virtualization scenarios. The basic network configuration scenario describes essential tasks for configuring an Oracle Solaris system on the network. The network virtualization scenarios describe networking strategies that combine multiple network virtualization features for high availability, optimal performance, resource management, and the deployment of workloads in a cloud environment.

For overview information about network administration, see [Chapter 1, “About Network Administration in Oracle Solaris”](#) in *Configuring and Managing Network Components in Oracle Solaris 11.3*.

For more information about administering network virtualization features, see [Managing Network Virtualization and Network Resources in Oracle Solaris 11.3](#).

This chapter contains the following topics:

- [“Basic Network Configuration Scenario”](#) on page 27
- [“Combining Aggregations With VNICs for High Availability”](#) on page 31
- [“Setting Up an EVS Virtual Tenant Network”](#) on page 33
- [“Combining Network Virtualization With Oracle VM Server for SPARC to Create a Cloud Environment”](#) on page 39

Basic Network Configuration Scenario

To perform basic network configuration of an Oracle Solaris system, you must first customize the datalinks on the system. Then, you configure the IP interfaces and IP addresses, as well as add a persistent default route for the system. Additionally, you configure any system-wide network services such as naming and directory services. The following examples assume that you are using the *fixed mode* for network configuration. See [Example 1, “Verifying the Active Network Mode,”](#) on page 28.

Depending on your particular networking needs, you might not need to perform each of the following tasks to configure your network. Or, you might need to perform additional tasks that are not described in this scenario. For a quick reference to commonly used network administration commands, see [Oracle Solaris 11.3 Network Administration Cheatsheet](#).

This section contains the following topics:

- [“Configure Datalinks, IP Interfaces, and IP Addresses” on page 28](#)
- [“Configure Naming Services Through SMF” on page 29](#)
- [“Set a System's Host Name” on page 31](#)

Configure Datalinks, IP Interfaces, and IP Addresses

The following configuration tasks are described:

- Verify the current network configuration mode.
- Determine how network interface names on the system map to physical interfaces.
- Configure a static IP interface and address.
- Add a persistent default route.

EXAMPLE 1 Verifying the Active Network Mode

After installing Oracle Solaris, verify which configuration mode you are using as follows

```
# netadm list
TYPE      PROFILE      STATE
ncp       Automatic    disabled
ncp       DefaultFixed online
loc       Automatic    offline
loc       NoNet        offline
loc       DefaultFixed online
```

The previous output indicates that the system is using the fixed mode, which means you use the `dladm`, `ipadm`, and `route` commands to manage network configuration.

If the system-generated Automatic profile is online, enable the DefaultFixed profile as follows:

```
# netadm enable -p ncp DefaultFixed
```

EXAMPLE 2 Determining How Network Interface Names Map to Physical Interfaces

Prior to configuring an IP interface and static IP address for a system, determine how the network interface names on the system map to the physical interfaces. On a system with multiple physical networks, use the `dladm` command to obtain this information:

```
# dladm show-phys
LINK      MEDIA      STATE    SPEED  DUPLEX  DEVICE
net0      Ethernet   up       1000   full    e1000g0
net1      Ethernet   unknown  0      unknown pcn0
```

EXAMPLE 3 Configuring a Static IP Address

You first create the IP interface and then you configure an IP address for the interface. You can associate multiple IP addresses with a single IP interface. In the following example, `ronj` is used for example purposes only.

```
# ipadm create-ip net0
# ipadm show-if
```

IFNAME	CLASS	STATE	ACTIVE	OVER
lo0	loopback	ok	yes	---
net0	ip	down	no	---

```
# ipadm create-addr -T static -a 10.163.198.20/24 net0/ronj
# ipadm show-if
```

IFNAME	CLASS	STATE	ACTIVE	OVER
lo0	loopback	ok	yes	---
net0	ip	ok	yes	---

```
# ipadm show-addr
```

ADDROBJ	TYPE	STATIC	ADDR
lo0/v4	static	ok	127.0.0.1/8
net0/ronj	static	ok	10.163.198.20/24
lo0/v6	static	ok	:::1/128

If your site implements IPv6 addressing, use the `addrconf` argument with the `-T` option to specify an automatically generated IPv6 address:

```
# ipadm create-ip net0
# ipadm create-addr -T addrconf net0/addr
```

If you need to obtain the IP address from a DHCP server, type the following commands:

```
# ipadm create-ip net0
# ipadm create-addr -T dhcp net0/addr
```

EXAMPLE 4 Adding a Persistent Default Route

After configuring an IP interface and address, add a persistent default route as follows:

```
# route -p add default 10.163.198.1
add net default: gateway 10.163.198.1
add persistent net default: gateway 10.163.198.1
```

For detailed instructions, see [“Creating Persistent \(Static\) Routes” in *Configuring and Managing Network Components in Oracle Solaris 11.3*](#).

Configure Naming Services Through SMF

Because the SMF repository is the primary repository for all naming services configuration in Oracle Solaris 11, the previous method of modifying a configuration file to configure naming

services no longer works. If you make changes to any of these services, for example, `svc:/system/name-service/switch`, `svc:/network/dns/client`, or `svc:/system/name-service/cache`, you must enable and refresh the service for the changes to take effect.

Note - If no network configuration exists, naming services default to `files` only behavior, rather than `nis` files. Note also that the `svc:/system/name-service/cache` SMF service should be enabled at all times.

The following configuration tasks are described:

- Configure DNS.
- Set multiple DNS options.
- Set multiple NIS servers.

EXAMPLE 5 Configuring DNS Through SMF

The following example shows how to configure the Domain Name Service (DNS) by using SMF commands. The DNS configuration on a system provides the ability to look up IP addresses by host name and host names by IP address. As shown in this example, you can set DNS properties from the command line, or you can set the same properties interactively. See [“Configuring a DNS Client” in *Configuring and Managing Network Components in Oracle Solaris 11.3*](#) for an example. After you set the various properties, you must enable and refresh the SMF service for the changes to take effect.

```
# svccfg -s dns/client setprop config/nameserver=net_address: 192.168.1.1
# svccfg -s dns/client setprop config/domain = astring: "examplehost.org"
# svccfg -s name-service/switch setprop config/host = astring: "files dns"
# svcadm refresh name-service/switch
# svcadm refresh dns/client
```

EXAMPLE 6 Configuring Multiple DNS Options Through SMF

One network configuration task that you might need to perform is to set DNS options for a system. The following example shows how to set multiple `/etc/resolv.conf` options simultaneously.

```
# svccg
svc:> select /network/dns/client
svc:/network/dns/client> setprop config/options = "ndots:2 retrans:3 retry:1"
svc:/network/dns/client> listprop config/options
config/options astring      ndots:2 retrans:3 retry:1

# svcadm refresh dns/client
# grep options /etc/resolv.conf
options ndots:2 retrans:3 retry:1
```

EXAMPLE 7 Configuring Multiple NIS Servers Through SMF

The following example shows how to set multiple NIS servers simultaneously.

```
# svccfg -s nis/domain setprop config/ypservers = host: (1.2.3.4 5.6.7.8) (Note that
there is a space between 1.2.3.4 and 5.6.7.8)
```

Set a System's Host Name

Note - The primary interface's TCP/IP host name is a distinct entity from the *system host name* that you set with the `hostname` command. Although not required by Oracle Solaris, the same name is normally used for both. Some network applications depend on this convention.

Permanently set a system's host name as follows:

```
# hostname name-of-host
```

Initially, the `hostname` value is stored in `config/nodename`, but this value is overridden if the system is configured by using DHCP, in which case, DHCP provides the `hostname` value. If the `hostname` command is used, then `hostname` is the value specified in the `config/nodename` file. If you set a system's identity by using the `hostname` command, this setting cannot be overridden by DHCP until you execute the `hostname` command with the `-D` option. The corresponding SMF properties and the associated SMF service are also automatically updated when you use the `hostname` command. See the [hostname\(1\)](#) man page.

Combining Aggregations With VNICs for High Availability

The following scenario describes how to combine a Datalink Multipathing (DLMP) aggregation with VNICs for high availability. [Figure 2, “Combining the Use of Aggregations With VNICs,” on page 18](#) graphically depicts this type of configuration.

The system that is used to create and configure the DLMP aggregation in the following example has a set of 10 Gigabit Ethernet NICs, as shown in the following output:

```
# dladm show-phys
LINK          MEDIA          STATE    SPEED  DUPLEX    DEVICE
net0          Ethernet       up       1000   full      e1000g0
net1          Ethernet       up       1000   full      e1000g1
net2          Ethernet       up       1000   full      e1000g2
```

EXAMPLE 8 Configuring and Virtualizing DLMP Aggregations With VNICs

1. First, create the DLMP aggregation (`aggr0`) with probing enabled for the `net1` and `net2` interfaces, as shown in the following example:

```
# dladm create-aggr -l net1 -l net2 -m dlmp -p probe-ip+= aggr0
```

Setting the `probe-ip` property enables probe-based detection failure with automatic selection of the source and target probing IP addresses. For details, see [“Configuring Probe-Based Failure Detection for DLMP Aggregation”](#) in *Managing Network Datalinks in Oracle Solaris 11.3*.

Then, create an IP interface and address for the aggregation datalink as follows:

```
# ipadm create-ip aggr0
# ipadm create-addr -T dhcp aggr0
```

2. Virtualize the DLMP aggregation.

You can then easily virtualize the aggregation by creating a VNIC over the aggregation datalink. For example, you would create a VNIC over `aggr0` as follows:

```
# dladm create-vnic -l aggr0 vnic0
```

The newly created VNIC (`vnic0`) is now highly available. If one of the aggregated links (`net1` or `net2`) fails, the traffic for that VNIC automatically fails over to the remaining link and the operation is transparent to the VNIC.

View information about the aggregation by using either of the following commands:

```
# dladm show-aggr
```

LINK	MODE	POLICY	ADDRPOLICY	LACPACTIVITY	LACPTIMER
aggr0	dlmp	--	--	--	--

```
# dlstat show-aggr -x
```

LINK	PORT	SPEED	DUPLEX	STATE	ADDRESS	PORTSTATE
aggr0	--	1000Mb	full	up	0:14:4f:fa:ea:42	--
	net1	1000Mb	full	up	0:14:4f:fa:ea:42	attached
	net2	1000Mb	full	up	0:14:4f:f9:c:d	attached

EXAMPLE 9 Specifying an Aggregation Datalink as the Lower Link of a Zone's anet Resource

Alternatively, you can virtualize an aggregation for high availability by specifying the aggregation datalink as the lower link of an Oracle Solaris zone's `anet` resource, as shown in the following example. Or, you can specify an aggregation datalink as the uplink of an EVS node. See [“Setting Up an EVS Virtual Tenant Network”](#) on page 33 for an example of this type of configuration.

The following truncated example shows how you would specify an aggregation datalink as the lower link of a zone's `anet` resource during a `zonecfg` interactive session,.

```
# zonecfg -z zone1
.
.
```



```
.
zonecfg:zone1> add anet
zonecfg:zone1:anet> set lower-link=aggr0
.
.
.
zonecfg:zone1:anet> end
zonecfg:zone1> commit
```

For more information about using the `zonecfg` command interactively, see the [zonecfg\(1M\)](#) man page and [Creating and Using Oracle Solaris Zones](#).

Setting Up an EVS Virtual Tenant Network

A virtual switch is either a software or hardware entity that facilitates inter-virtual machine (VM) communication by looping the inter-VM traffic within a physical machine rather than sending it out on the wire.

EVS enables you to explicitly create virtual switches that span one or more nodes (physical machines), which further virtualizes your network. The virtual switch that you create represents an isolated L2 segment that uses either VLANs or VXLANs to implement the isolation.

For more information about the EVS architecture, see “[EVS Components](#)” in *Managing Network Virtualization and Network Resources in Oracle Solaris 11.3*.

This scenario's overall objective is to set up and deploy an EVS virtual tenant network. The primary objective is to create an elastic virtual switch (vswitch) that connects two compute nodes so that both nodes are part of the same L2 segment and can communicate with each other.

The individual objectives for this scenario are as follows:

- Deploy a virtual tenant network with two zones that are connected to the network through anet VNICs.
- Deploy the VNICs on a private cloud infrastructure that has two compute nodes.
- Use a VLAN L2 infrastructure to instantiate the private virtual tenant network.

Note - Other L2 technologies such as VXLANs are also supported. For more information, see “[Use Case: Configuring an Elastic Virtual Switch for a Tenant](#)” in *Managing Network Virtualization and Network Resources in Oracle Solaris 11.3*.

The following figures represent the virtual and physical components of the elastic virtual switch configuration that is used in this scenario.

FIGURE 4 Virtual Components of an EVS Switch Configuration


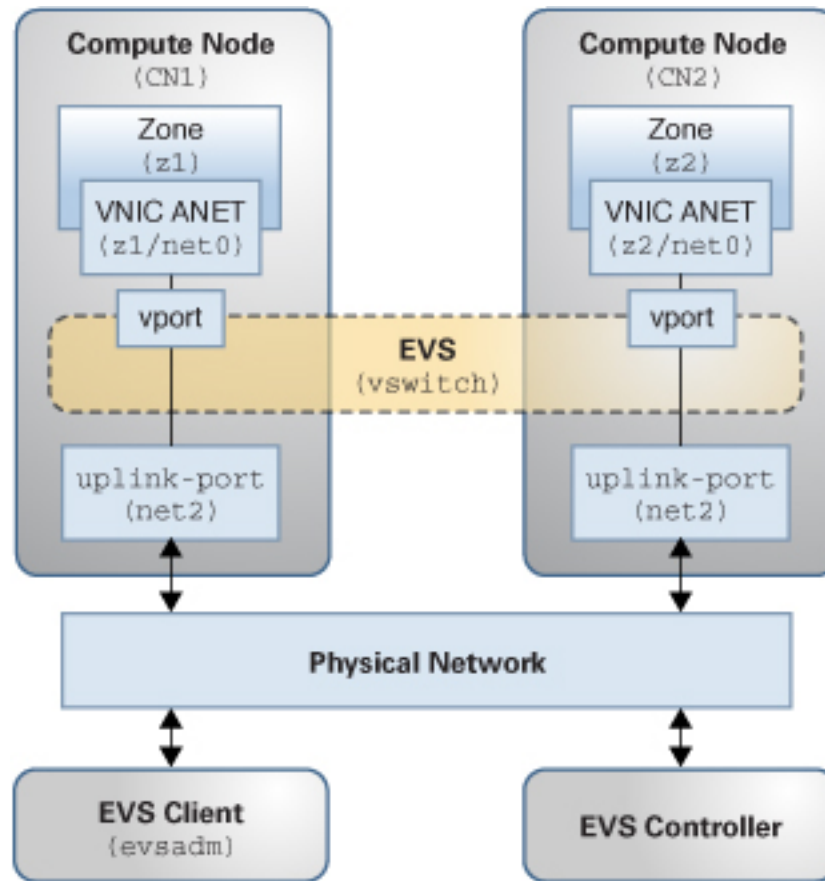


FIGURE 5 Physical Components of an EVS Switch Configuration

The following setup uses four network nodes with the following configuration:

- Two compute nodes (CN1 and CN2).
- Two zones (z1 and z2) that are configured on CN1 and CN2, respectively.
- The two zones (z1 and z2) are configured with a VNIC anet resource on each zone.
- One node that acts as the EVS controller.
- One node that acts as the EVS client.

Note - The EVS controller and the EVS client can be located on the same host.

- Two uplink-ports (net2) that specify the datalink to be used for the VLANs.

Perform Preliminary Tasks Prior to Creating an EVS Virtual Tenant Network

The following one-time setup tasks are described:

- Plan the EVS virtual tenant network deployment by doing the following:
 - Select the two compute nodes.
 - Designate a node to act as the controller.
 - Designate a node to act as the EVS client.

Note - The client and controller nodes can be on the same host.

- Select the VLAN ID range to be used for tenant traffic.
- Decide which datalink to use for tenant traffic on each compute node.
- Install the base EVS package (pkg:/service/network/evs) on every node.
- Install the pkg:/system/management/rad/module/rad-evs-controller package on the controller node.
- Configure each of the nodes so that Remote Administration Daemon (RAD) invocations are enabled.
- On every node, configure EVS to point to the controller.
- From the EVS client node, configure the controller properties.
- From the EVS client node, verify the controller configuration.

EXAMPLE 10 Installing Mandatory EVS Packages

Prior to setting up an EVS switch, you need to install the necessary software packages. You install these packages on each EVS node separately.

Install the base EVS package (pkg:/service/network/evs) on every node (client, controller, and the compute nodes) as follows:

```
# pkg install evs
```

Install the pkg:/system/management/rad/module/rad-evs-controller package on the node that is designated as the EVS controller as follows:

```
# pkg install rad-evs-controller
```

After installing the mandatory EVS packages and prior to configuring and setting properties for the EVS controller, you must configure all of the nodes so that RAD invocations between each

node can take place. For detailed instructions, see [“Security Requirements for Using EVS” in *Managing Network Virtualization and Network Resources in Oracle Solaris 11.3*](#).

EXAMPLE 11 Configuring and Setting Properties for the EVS Controller

The EVS controller provides the resources that are associated with creating and managing elastic virtual switches. You set properties for the controller that specify the necessary information for implementing L2 segments across physical nodes. See [“EVS Controller” in *Managing Network Virtualization and Network Resources in Oracle Solaris 11.3*](#).

Configure each compute node so that it points to the EVS controller. This scenario uses two compute nodes, so you will need to run the following command on each of the compute nodes:

```
# evsadm set-prop -p controller=CONTROLLER
```

From the client node, configure the EVS controller properties.

1. Set the L2 topology.

```
# evsadm set-controlprop -p l2-type=vlan
```

2. Set the VLAN range.

```
# evsadm set-controlprop -p vlan-range=200-300
```

3. Specify the uplink-port (datalinks) that are used for the VLAN.

```
# evsadm set-controlprop -p uplink-port=net2
```

4. Verify the controller configuration on the EVS client.

```
# evsadm show-controlprop -p l2-type,vlan-range,uplink-port
```

NAME	VALUE	DEFAULT	HOST
l2-type	vlan	vlan	--
vlan-range	200-300	--	--
uplink-port	net2	--	--

Create an EVS Virtual Tenant Network (vswitch)

The following examples show how you would set up and configure an EVS virtual tenant network named vswitch. Pay particular attention to where you perform each task.

The following configuration tasks are described:

- From the client node, set up a virtual switch.
- On each compute node, create a zone and then connect the zone to the virtual switch.
- From the client node, display the EVS configuration.

For an overview of the EVS feature, see [Chapter 5, “About Elastic Virtual Switches” in *Managing Network Virtualization and Network Resources in Oracle Solaris 11.3*](#).

EXAMPLE 12 Setting Up an EVS Switch

The following example shows how you would set up an EVS virtual tenant network. You perform this task from the client node.

First, create the EVS switch, named `vswitch` in this example, as follows:

```
# evsadm create-evs vswitch
```

Add the IPnet information to the EVS switch and verify the configuration.

```
# evsadm add-ipnet -p subnet=192.168.70.0/24 vswitch/ipnet
# evsadm show-ipnet
```

NAME	TENANT	SUBNET	DEFROUTER	AVAILRANGE
vswitch/ipnet	sys-global	192.168.70.0/24	192.168.70.1	192.168.70.2-192.168.70.254

Verify that the EVS switch was successfully created.

```
# evsadm
NAME          TENANT      STATUS  VNIC    IP          HOST
vswitch       sys-global  --      --      vswitch_ipnet  --
```

Check the VLAN ID that is associated with the virtual switch.

```
# evsadm show-evs -L
EVS    TENANT  VID    VNI
vswitch  sys-global  200    --
```

EXAMPLE 13 Creating and Connecting a Zone to an EVS Switch

The following example describes how you would create a zone on each tenant, and then connect the zone to the virtual switch.

On each tenant, configure a zone with an `anet` VNIC resource as follows:

```
# zonecfg -z z1
zonecfg:z1> create
.
.
.
zonecfg:z1> add anet
zonecfg:z1:anet> set evs=vswitch
zonecfg:z1:anet> end
zonecfg:z1> commit
zonecfg:z1> exit
```

For more information, see [“Creating a VNIC `anet` Resource for an Elastic Virtual Switch” in *Managing Network Virtualization and Network Resources in Oracle Solaris 11.3*](#).

For information about setting anet resource properties that pertain to an EVS switch, see [“Resource Types and Properties” in Oracle Solaris Zones Configuration Resources](#).

Boot the zone.

```
# zoneadm -z z1 boot
```

Verify that the VNIC was created and is connected to the virtual switch.

```
# dladm show-vnic -c
```

LINK	TENANT	EVS	VPORT	OVER	MACADDRESS	VIDS
z1/net0	sys-global	vswitch	sys-vport0	net2	2:8:20:1a:c1:e4	200

From within the zone, verify that the IP address has been assigned.

```
# zlogin z1 ipadm
```

NAME	CLASS/TYPE	STATE	UNDER	ADDR
lo0	loopback	ok	--	--
lo0/v4	static	ok	--	127.0.0.1/8
lo0/v6	static	ok	--	::1/128
net0	ip	ok	--	--
net0/v4	inherited	ok	--	192.168.84.3/24

From the client node, display the EVS configuration.

```
# evsadm
```

NAME	TENANT	STATUS	VNIC	IP	HOST
vswitch	sys-global	-- --	vswitch_ipnet		

EVS provides a rich set of functionality that is not fully described in this scenario. For additional tasks and use cases, see [Chapter 6, “Administering Elastic Virtual Switches” in Managing Network Virtualization and Network Resources in Oracle Solaris 11.3](#).

Combining Network Virtualization With Oracle VM Server for SPARC to Create a Cloud Environment

The following scenario combines network virtualization features with Oracle VM Server for SPARC to create a multilevel virtual network that parallels a cloud environment. This deployment method provides highly efficient, enterprise-class virtualization capabilities for Oracle's SPARC T-Series servers and supported M-Series servers.

This scenario assumes that you are running an Oracle VM Server for SPARC version that supports Oracle Solaris 11.2. For more information about Oracle VM Server for SPARC, refer to the documentation library at <http://www.oracle.com/technetwork/documentation/vm-sparc-194287.html>.

At a high level, the objective of this scenario is to carve up a SPARC based system into multiple Oracle Solaris VM Server guest domains, where each domain corresponds to a node within a cloud environment. You can deploy per-tenant workloads as zones within these Oracle VM Server for SPARC guest domains.

Configuring network virtualization features in this way enables you to build an entire cloud within a single SPARC based system. Or, you can use this type of configuration to integrate a SPARC based system into a larger cloud environment, where the system appears as a set of nodes within that environment.

Combining network virtualization features with Oracle VM Server for SPARC parallels a traditional cloud in the following ways:

- Compute nodes are implemented as Oracle VM Server for SPARC guest domains.
- Compute nodes communicate with each other through the virtual network infrastructure that is provided by Oracle VM Server for SPARC and Oracle Solaris 11 running on the service domain.
- The vnet driver instances that are within each guest domain correspond to a physical NIC within a physical compute node.

The benefits of this type of configuration include the following:

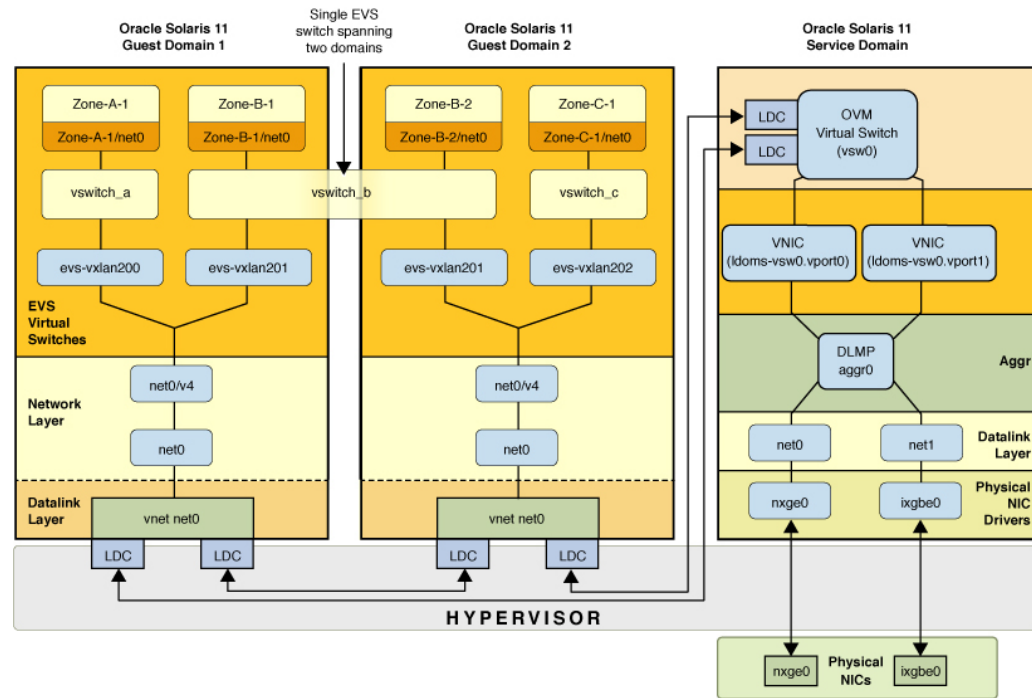
- Enables more flexibility by enabling you to run smaller domains that you can upgrade individually without affecting other workloads that are running on the system.
- Takes advantage of SPARC Reliability, Availability, and Serviceability (RAS) features.
- Uses a faster virtual network for communication between nodes instead of relying on a physical infrastructure.

Objectives for Creating and Deploying a Cloud Environment

The deployment objectives for this scenario are as follows:

- Configure a virtual network on an Oracle VM Server for SPARC service domain.
- Configure two Oracle VM Server for SPARC guest domains to be used as containers for multiple zones that are configured within each guest.
- Have each guest domain then correspond to a specific compute node within the cloud that will run the various workloads.
- Configure the elastic virtual switches that will be used to connect the zones running in the guest domains.
- Carve up the guest domains into multiple zones that will run the various workloads.

The following figure illustrates the two distinct levels of network virtualization that you create with this configuration.

FIGURE 6 Combining Network Virtualization Features With Oracle VM Server for SPARC

On the first level, you configure network virtualization features that are supported by Oracle VM Server for SPARC. This part of the network virtualization combines Oracle VM Server for SPARC configuration with the Oracle Solaris 11 OS that is running on the service domain. The vnet configuration takes place at this first level of virtualization. Because the configuration only relies upon IP connectivity from the guest domains, no additional support from Oracle VM Server for SPARC is required for the configuration on the second network virtualization level to work.

On the second level, EVS is used to create elastic virtual switches across the guest domains. EVS is configured to use the vnet interfaces as uplinks. VXLAN datalinks are automatically created by EVS from each guest domain and then used to encapsulate the traffic of the individual elastic virtual switches.

The figure represents the following configuration:

- Two physical NICs, `nxge0` and `ixgbe0`, which are directly assigned to the service domain where they are represented by datalinks `net0` and `net1`.
- To provide high availability in case of failure of the physical NICs, `net0` and `net1` in the service domain are grouped into the DLMP aggregation (`aggr0`).

- The aggregation, `aggr0`, is then connected to an Oracle VM Server for SPARC virtual switch in the service domain named `vsw0`.
Two VNICs, `ldoms-vsw.vport0` and `ldoms-vsw.vport1`, are automatically created by `vsw0`, with each VNIC then corresponding to the Oracle VM for SPARC `vnet` instances within the guest domains.
- The `vsw0` and the `vnet` instances communicate with each other through the hypervisor by using Logical Domain Channels (LDCs).
- Each guest uses its instance of the `vnet0` driver, which appears in the guest domain as a datalink (`net0`) for the purpose of communicating with other guest domains and the physical network.
- In each guest domain, the `vnet` datalinks (`net0`) are configured with the IP interface `net0/v4`.
- Each guest domain is an EVS compute node, with three EVS switches, `vswitch_a`, `vswitch_b`, and `vswitch_c`, that are configured from the EVS controller (not shown in this figure).
- EVS is configured to use a VXLAN as its underlying protocol. For each guest domain that uses an elastic virtual switch, EVS automatically configures a VXLAN datalink. These VXLAN datalinks are named `evs-vxlanid`, where `id` is the VXLAN ID that is assigned to the virtual switch.
- In the guest domains, Oracle Solaris zones are configured to run the tenants' workload. Each zone is connected through a VNIC and a virtual port (not shown in this figure) to one of the EVS switches.
- Zone-B1 and Zone-B2 belong to the same user and are running on two different guest domains. The EVS switch, `vswitch_b`, is instantiated on both guest domains. To the two zones, it appears as if each zone is connected to a single Ethernet segment that is represented by `vswitch_b` and isolated from the other virtual switches.
- EVS automatically creates the VXLAN datalinks that are needed by the various elastic virtual switches. For example, for `vswitch_b`, EVS automatically created a VXLAN datalink named `evs-vxlan201` on each of the guest domains.

Configure a Virtual Network on Oracle VM Server for SPARC Service and Guest Domains

Perform the following configuration tasks:

- On the service domain, create and configure a DLMP aggregation.
- On the service domain, configure the Oracle VM Server for SPARC virtual switch.
- On the service domain, configure the Oracle VM Server for SPARC virtual network devices for use on the guest domains
- On each of the guest domains, configure an IP address for each `vnet`.

The following examples assume that you have already configured your Oracle VM Server for SPARC (formerly called Sun Logical Domains, or LDoms) infrastructure with a control domain and a service domain and that you have created two guest domains to be used as cloud nodes.

For step-by-step instructions on setting up your Oracle VM Server for SPARC infrastructure, see the white papers at <http://www.oracle.com/technetwork/server-storage/vm/overview/index.html>.

Note - The examples that describe this scenario are presented in the order in which you would perform each of the individual tasks.

EXAMPLE 14 Creating and Configuring a DLMP Aggregation

The following example describes the first configuration task in this scenario, which is to create a DLMP aggregation on the Oracle VM Server for SPARC service domain. In this example, you create the DLMP aggregation (aggr0) with probing enabled for the net1 and net2 interfaces.

```
servicedomain# dladm create-aggr -l net1 -l net2 -m dlmp -p probe-ip+= aggr0
```

For additional information, see [Example 8, “Configuring and Virtualizing DLMP Aggregations With VNICs,” on page 31](#).

EXAMPLE 15 Creating an Oracle VM Server for SPARC Virtual Switch

One of the basic components that Oracle VM Server for SPARC uses for virtual networking is a virtual switch (vsw). A virtual switch is similar to an Ethernet switch, in that it runs in an I/O or service domain and switches Ethernet packets over Logical Domain Channels (LDCs), as well using the Oracle Solaris 11 built-in virtual switch.

The following example shows how to create a virtual switch on the DLMP link portion of the configuration. You perform this task on the service domain.

```
servicedomain# ldm add-vsw net-dev=aggr0 primary-vsw0 primary
```

EXAMPLE 16 Creating Virtual Network Devices for Oracle VM Server for SPARC Guest Domains

The second basic component that Oracle VM Server for SPARC uses for virtual networking is a virtual network device (vnet). A virtual network device is plumbed into the guest domain.

The following example shows the next configuration task in this scenario, which is to create a virtual network device for each guest domains. You also perform this task on the service domain.

```
servicedomain# ldm add-vnet
```

You create one virtual network device, per guest domain. For every device that you create, a vnet instance is also created in the corresponding guest domain.

Next, configure an IP address for each vnet on each guest domain as follows:

```
guestdomain1# ipadm create-ip net0
guestdomain# ipadm create-addr -t -a 192.168.70.1 net0
```

```
guestdomain2# ipadm create-ip net0
guestdomain# ipadm create-addr -t -a 192.168.70.2 net0
```

For more information about configuring virtual switches and creating virtual network devices for Oracle VM Server for SPARC, see the relevant documentation at <http://www.oracle.com/technetwork/documentation/vm-sparc-194287.html>.

Create an EVS Switch to Deploy the Cloud Workload

The next group of tasks involves creating an EVS switch that you use to deploy the cloud workload. Some configuration tasks are performed on the Oracle VM Server for SPARC service domain, while others are performed on the guest domains.

The following EVS setup is used:

- Two compute nodes that correspond to two guest domains. Each guest domain has a `net0` interface for its vnet datalink, which are then used as uplink-ports by the elastic virtual switch.
- One node that acts as the EVS controller.
- One node that acts as the EVS client.

Note - The EVS controller and the EVS client can be located on the same host.

- Set of four zones: Zone-A1 and Zone-B1, which are configured on the first guest domain, and Zone-B2 and Zone-C2, which are configured on second guest domain.
- The four zones are configured with a VNIC (anet) resource on each zone, which are then connected to the EVS switches.

▼ How to Configure an EVS Virtual Switch to Deploy a Cloud Workload

Before You Begin Perform all of the necessary planning and prerequisite tasks, which include installing the EVS packages and configuring the appropriate authorizations.

For planning instructions, see [“Perform Preliminary Tasks Prior to Creating an EVS Virtual Tenant Network” on page 36](#).

For security requirements, see [“Security Requirements for Using EVS” in *Managing Network Virtualization and Network Resources in Oracle Solaris 11.3*](#).

1. Configure each compute node to point to the EVS controller.

```
# evsadm set-prop -p controller=CONTROLLER
```

You can deploy the EVS controller on any node, as long as it can reach the Oracle VM Server for SPARC guest domains through the vnet interfaces.

For example, you can deploy the EVS controller in any of the following ways:

- In the service domain's global zone
- In a non-global zone of the service domain
- In its own guest domain
- On a separate physical machine

For more information about configuring an EVS controller, see [“Configuring an EVS Controller” in *Managing Network Virtualization and Network Resources in Oracle Solaris 11.3*](#).

2. On the controller, configure the required EVS properties for the compute nodes.

a. Set the L2 topology.

```
# evsadm set-controlprop -p l2-type=vxlan
```

b. Set the VXLAN range and IP address.

```
# evsadm set-controlprop -p vxlan-range=200-300
# evsadm set-controlprop -p vxlan-addr=192.168.70.0/24
```

Note that you determine the VXLAN range during the planning phase, prior to setting up the EVS switch. For more information about configuring EVS controller properties, see [Example 11, “Configuring and Setting Properties for the EVS Controller,” on page 37](#).

c. Specify the uplink-port (datalinks) that are used for the VXLAN.

```
# evsadm set-controlprop -p uplink-port=net0
```

d. Verify the configuration.

```
# evsadm show-controlprop -p l2-type,vxlan-range,vxlan-addr
```

NAME	VALUE	DEFAULT	HOST
l2-type	vxlan	vxlan	--
vxlan-range	200-300	--	--
vxlan-addr	192.168.70.0/24	0.0.0.0	--
uplink-port	net0	--	--

The controller must be assigned an IP address that is reachable from each of the guest domains. For this example, that IP address is 192.168.70.10.

3. Create and verify the EVS virtual switch, which is named `vswitch_a` in this example.

a. Create the EVS switch.

```
# evsadm create-evs vswitch_a
```

Repeat this step to create the two other EVS switches that are used in the configuration (`vswitch_b` and `vswitch_c`).

b. Add the IPnet information to the EVS switch and verify the configuration.

```
# evsadm add-ipnet -p subnet=192.168.80.0/24 vswitch_a/ipnet
# evsadm show-ipnet
```

NAME	TENANT	SUBNET	DEFROUTER	AVAILRANGE
vswitch_a/ipnet	sys-global	192.168.80.0/24	192.168.80.1	192.168.80.2-192.168.80.254

Repeat this step for the two other EVS switches that are used in the configuration (`vswitch_b` and `vswitch_c`).

c. Verify that the virtual switch was successfully created.

```
# evsadm
NAME          TENANT      STATUS  VNIC      IP          HOST
vswitch_a     sys-global  --      --        vswitch_a/ipnet  --
```

d. Check the VLAN ID that is associated with the virtual switch.

```
# evsadm show-evs -L
EVS      TENANT      VID      VNI
vswitch_a  sys-global  --       200
vswitch_b  sys-global  --       201
vswitch_c  sys-global  --       202
```

Create Oracle Solaris Zones on the Oracle VM Server for SPARC Guest Domains

The following example shows how to create a zone within an Oracle VM Server for SPARC guest domain for the purpose of deploying the cloud workload. The following commands create

a zone on the guest domain with an anet that uses a VXLAN as the underlying link of the Oracle VM Server for SPARC virtual switch.

```
# zonecfg -z B-1
zonecfg:B-1> create
.
.
.
zonecfg:B-1> add anet
zonecfg:B-1:anet> set evs=vswitch_b
zonecfg:B-1:anet> end
zonecfg:B-1> commit
zonecfg:B-1> exit
```

For more information about configuring zones, see [Creating and Using Oracle Solaris Zones](#).

Index

A

- active network mode
 - verifying, 28
- aggregation
 - DLMP
 - example, 31
 - virtualizing DLMP
 - example, 32
- aggregation description, 11
- aggregations
 - combining with VNICs
 - use case, 31
- anet resource properties, 39

B

- basic network configuration
 - summary, 10
- basic network configuration scenarios
 - example, 27
- bridging description, 11
- building blocks for network virtualization, 19

C

- cloud
 - configuring an EVS switch, 44
 - creating by using EVS, 39
- cloud architecture, 21
 - utility computing model, 21
- cloud environment
 - combining network virtualization with Oracle Solaris VM Server for SPARC, 40
 - configuring an Oracle Solaris VM Server for SPARC virtual device, 43

- configuring an Oracle Solaris VM Server for SPARC virtual switch, 43
- configuring Oracle Solaris VM Server for SPARC, 43
- configuring zones, 46
- cloud environments
 - using for high availability, 22
- cloud network
 - description of, 21
- cloud networking
 - network virtualization strategies, 21
- combining aggregations with VNICs
 - illustration, 19
 - use case, 31
- combining multiple networking features
 - figure of, 23
- combining networking features
 - networking strategies, 18
- commands for managing network resources
 - dladm, 24
- configuring DNS
 - example, 30
- configuring EVS controller properties, 37
- configuring IP interfaces and addresses
 - example, 28
- configuring naming services
 - through SMF, 29
- configuring NIS
 - example, 31
- configuring persistent default route
 - example, 29
- configuring static IP addresses
 - example, 29
- configuring zones
 - Oracle VM Server for SPARC, 46
- controller

- creating elastic virtual switches, 20
- setting EVS properties, 37
- creating a cloud environment
 - with EVS, 39
- creating an EVS switch, 37
 - connecting to a zone, 38
 - displaying configuration, 39
- creating DLMP aggregation
 - use case, 31
- creating elastic virtual switches
 - EVS controller, 20

D

- Data Center Bridging (DCB), 12
- datalink configuration
 - example, 28
- default route
 - configuring persistently
 - example, 29
- deploying a cloud workload
 - using an EVS switch, 44
- deploying workloads efficiently
 - cloud networking, 21
- description of a flow, 12
- description of a load balancer
 - ILB, 12
- description of a trunk aggregation, 11
- description of a VLAN, 13
- description of a VNIC, 13, 13
- description of a VXLAN, 13
- description of an aggregation, 11
- description of an DLMP, 11
- description of an etherstub, 12
- description of bridging, 11
- description of DCB, 12
- description of EVB, 12
- description of EVS, 12
- description of IP tunnels, 13
- description of IPMP
 - ILB, 12
- description of LLDP, 13
- description of VRRP, 13
- determining mappings between network interface names and physical interfaces
 - example, 28

- displaying configuration information
 - EVS switch, 39
- dladm
 - commands for managing network resources, 24
- DLMP aggregation
 - virtualizing, 32
- DLMP description, 11
- DNS
 - configuring through SMF, 30

E

- Edge Virtual Bridging (EVB), 12
- elastic virtual switch
 - installing packages, 36
 - setting up a virtual tenant network, 33
- Elastic Virtual Switch (EVS), 12
- Elastic Virtual Switch feature
 - network virtualization building blocks, 19
- etherstub, 12
- EVS
 - anet resource properties, 39
 - vports, 24
- EVS configuration
 - description, 24
- EVS controller
 - configuring, 37
 - creating elastic virtual switches, 20
- EVS packages
 - virtual tenant network, 36
- EVS switch
 - creating, 37
 - creating and connecting to a zone, 38
 - deploying a cloud workload, 44
 - displaying configuration, 39
 - using to create a cloud environment, 39
- evsadm command
 - example, 39
- example
 - adding a default persistent route, 29
 - configuring DNS, 30
 - configuring naming services, 29
 - configuring NIS, 31
 - DLMP aggregation, 31
 - EVS virtual tenant network, 33

- mapping network interface names to physical interfaces, 28
 - network virtualization use case, 31
 - setting a system's host name, 31
 - static IP address configuration, 29
 - verifying active network mode, 28
 - virtualizing DLMP aggregation, 32
 - examples of basic network configuration
 - datalink and IP interfaces, 28
 - use cases, 27
 - examples of datalink and IP interface configuration, 28
 - examples of network configuration
 - use cases, 27
- F**
- feature descriptions
 - network administration, 11
 - features for managing network security, 25
 - flowadm
 - managing network resources, 24
 - flows, 12, 24
 - functional areas
 - networking features, 16
- H**
- high availability
 - combining aggregations with VNICs, 31
 - networking features that support, 16
 - using cloud environments, 22
 - high availability examples
 - combining aggregations with VNICs, 19
 - host name
 - how to set, 31
 - hostname command
 - example, 31
- I**
- Integrated Load Balancer (ILB), 12
 - interface names
 - mapping physical interface to network names, 28
 - IP configuration
 - configuring static addresses, 29
 - IP interface and IP address configuration
 - example, 28
 - IP network multipathing (IPMP), 12
 - IP tunnel, 13
- K**
- key network administration features, 11
- L**
- Link Layer Discovery Protocol (LLDP), 13
- M**
- managing network resources, 24
 - features, 24
 - flowadm, 24
 - flows, 24
 - managing virtual switches
 - features in Oracle Solaris, 20
 - mandatory package installation
 - setting up an EVS virtual tenant network, 36
 - mode used for configuration
 - verifying, 28
- N**
- naming service configuration
 - SMF commands
 - example, 29
 - naming services
 - configuring DNS
 - example, 30
 - configuring NIS
 - example, 31
 - network administration
 - by functional area, 16
 - network administration by functional area
 - high availability, 18
 - network virtualization, 18
 - performance, 18
 - resource management, 18
 - security, 18
 - network administration feature descriptions, 11
 - network administration features, 9

- network configuration
 - combining multiple networking features
 - figure of, 23
- network configuration in Oracle Solaris
 - summary, 10
- network configuration scenarios, 27
- network consolidation
 - creating virtual private networks, 21
- network resources
 - managing, 24
- network security features, 25
- network stack
 - description of, 14
- network stack layers
 - network administration, 15
- network storage
 - networking features that support, 16
- network virtualization
 - combining Oracle VM Server for SPARC with zones and EVS, 40
 - networking features that support, 16
- network virtualization building blocks, 19
- network virtualization in Oracle Solaris
 - description, 19
- network virtualization strategies
 - cloud networking, 21
 - virtual private networks, 21
 - workload consolidation, 21
- networking examples
 - using aggregations with VNICs
 - high availability features, 19
- networking feature descriptions
 - aggregation, 11
 - bridging, 11
 - DCB, 12
 - DLMP, 11
 - etherstub, 12
 - EVB, 12
 - EVS, 12
 - flows, 12
 - I/O virtualization (SR-IOV), 13
 - ILB
 - load balancer, 12
 - IP tunnels, 13
 - IPMP, 12
 - LLDP, 13

- trunk aggregation, 11
- virtual switch, 14
- VLAN, 13
- VNIC, 13
- VRRP, 13
- VXLAN, 13
- networking features by network protocol stack layer, 15
- NIS
 - configuring through SMF, 31

O

- Oracle Solaris network protocol stack
 - description of, 14
- Oracle VM Server for SPARC
 - combining with network virtualization, 39
 - configuring a virtual device, 43
 - configuring a virtual switch, 43
 - configuring service and guest domains, 43
 - configuring zones on guest domain, 46
 - network virtualization use case, 40

P

- package installation
 - EVS, 36
- performance
 - networking features that support, 16
- physical interface names
 - mapping to network interface names
 - example, 28
- private virtual networks
 - security, 21
 - using for testing and simulation, 21
- properties
 - configuring for the EVS controller, 37

R

- resource management
 - networking features that support, 16
- resource properties
 - anetEVS switches, 39
- root I/O virtualization
 - (SR-IOV), 13

- route
 - configuring persistently, 29
- route command
 - example, 29

S

- scenario
 - setting up an EVS virtual tenant network, 33
- scenarios
 - basic network configuration, 27
 - combining aggregations with VNICs for high availability, 31
 - configuring datalinks, 28
 - network configuration, 27
 - scenario
 - zones and Oracle VM Server, 40
- security
 - creating private virtual networks, 21
 - networking features that support, 16
- security features
 - networking, 25
- setting host name
 - example, 31
- SMF commands
 - configuring naming services, 29
- stack
 - description of the network protocol stack, 14
- stack layers
 - feature descriptions, 15
- static IP addresses
 - configuring
 - example, 29
- strategies for network administration
 - combining features, 18
- strategies for network virtualization, 21
- summary of basic network configuration, 10
- summary of network virtualization in Oracle Solaris, 19
- supported network administration features, 9
- switches
 - virtual, 14

T

- testing and simulation

- creating private virtual networks, 21
- trunk aggregation description, 11
- tunnels
 - description of, 13

U

- use case
 - combining aggregations with VNICs, 31
 - creating DLMP aggregation, 31
 - setting up an EVS virtual tenant network, 33
 - virtualizing DLMP aggregation, 32
- use cases
 - basic network configuration, 27
 - configuring datalinks and IP interfaces, 28
 - network configuration, 27
- using multiple networking features, 18
- utility computing model
 - cloud networking, 21

V

- verifying active network mode
 - example, 28
- virtual device
 - configuring Oracle VM Server for SPARC, 43
- virtual eXtensible area network (VXLAN), 13
- virtual local area network (VLAN), 13
- virtual network
 - Oracle VM Server for SPARC configuration, 43
- virtual network interface card (VNIC)
 - virtual NIC, 13
- virtual network stack
 - cloud environments
 - high availability, 22
- virtual private networks
 - network consolidation, 21
 - strategies for network virtualization, 21
- Virtual Router Redundancy Protocol (VRRP), 13
- virtual switch
 - configuring Oracle VM Server for SPARC, 43
 - connecting to a zone
 - EVS, 38
 - creating a virtual tenant network
 - EVS, 37

- network virtualization building blocks, 14
- virtual switches
 - managing
 - EVS, 20
 - network virtualization building blocks, 19
- virtual tenant network
 - configuring the EVS controller
 - example, 37
 - creating an EVS switch
 - connecting to a zone example, 38
 - example, 37
 - EVS use case
 - example, 33
 - installing packages
 - example, 36
- VNICs
 - combining with aggregations, 31
 - network virtualization building blocks, 19
- vports
 - EVS, 24
- VXLANs
 - using with EVS configuration, 24

W

- workload consolidation
 - network virtualization strategies, 21

Z

- zones
 - deploying in a cloud environment, 46